

Introduction to Machine Learning (IML)

# LECTURE #2: CLASSIFICATION - PRELIMINARIES

---

236756 – 2024 WINTER – TECHNION

LECTURER: YONATAN BELINKOV



# Classification

## Today:

- Fundamentals of classification via simple example
- But don't confuse *simple* with *easy*
- Sets ground for the first  $\sim 2/3$  of the entire course
- We'll revisit many of the issues we'll discuss today
- So make sure you understand these fundamentals now

# Classification

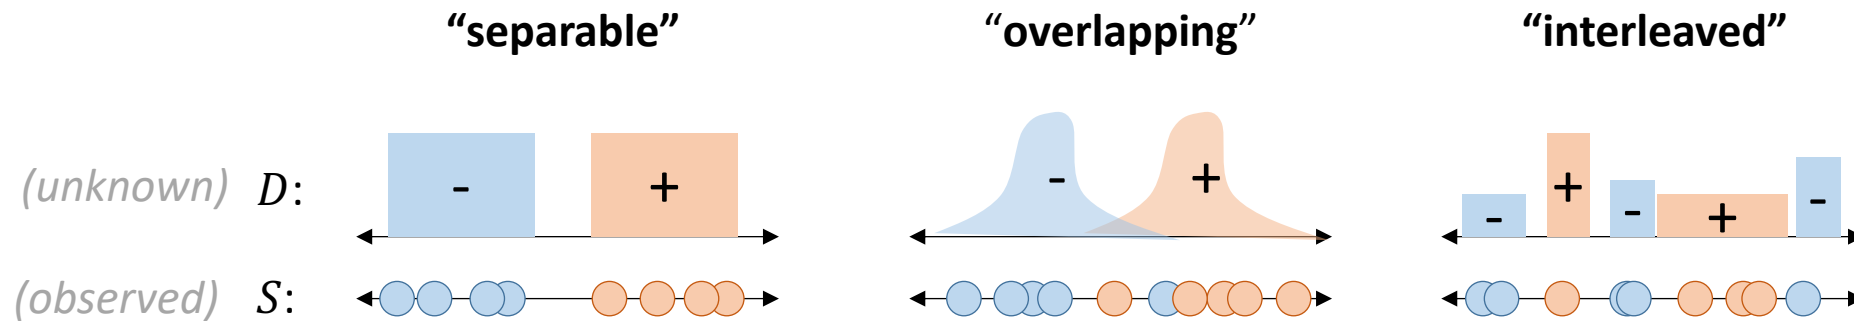
## Recall:

- **Data generating process:**  $(x, y) \stackrel{iid}{\sim} D, y \in \{\pm 1\}$  (classification)
- **Sample set:**  $S = \{(x_i, y_i)\}_{i=1}^m \sim D^m$  (data)
- **Model class:**  $H = \{h : h: \mathcal{X} \rightarrow \mathcal{Y}\}$
- **Expected error:**  $L_D(h) = \mathbb{P}_D[y \neq h(x)] = \mathbb{E}_D[\mathbb{1}\{y \neq h(x)\}]$
- **Empirical error:**  $L_S(h) = \mathbb{P}_S[y \neq h(x)] = \frac{1}{m} \sum_{i=1}^m \mathbb{1}\{y_i \neq h(x_i)\}$
- **Goal:** learn  $h \in H$  with low  $L_D(h)$
- **Means:** return  $\hat{h} = \operatorname{argmin}_{h \in H} L_S(h)$  (ERM)

# Classification

## Today:

- Learning on the real line,  $\mathcal{X} = \mathbb{R}$
- Consider three “types” of distributions of increasing difficulty:



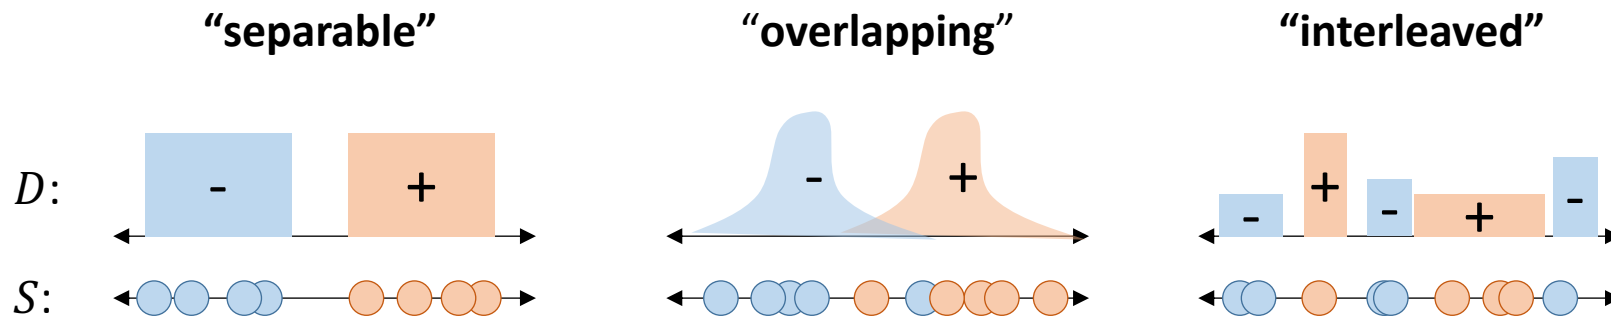
- For each, we will discuss the **three aspects of learning** we care about:

Modeling, Optimization, Statistics

# Classification

## Today:

- As we progress, we will make less and less assumptions
- Holy grail: having no assumptions on  $D$
- Our goal for today: develop intuition for *distribution-independent* learning algorithms
- **Head's up:** we will ask many questions (but will answer relatively few)



# Case I: Separable Data

# Separable data

- **Definition (interim)\*:**

$D$  (over  $\mathbb{R}$ ) is *separable* if exists  $\theta^* \in \mathbb{R}$  such that  $y = 1$  iff  $x > \theta^*$

- Note this means labels are deterministic.


\* This definition is not precise, and is in fact class-dependent. This will later be refined.

# Separable data

- **Definition (interim)\*:**

$D$  (over  $\mathbb{R}$ ) is *separable* if exists  $\theta^* \in \mathbb{R}$  such that  $y = 1$  iff  $x > \theta^*$

- Note this means labels are deterministic.
- Example distribution:



\* This definition is not precise, and will later be refined.




# Separable data

- **Definition (interim)\*:**

$D$  (over  $\mathbb{R}$ ) is *separable* if exists  $\theta^* \in \mathbb{R}$  such that  $y = 1$  iff  $x > \theta^*$

- Note this means labels are deterministic.
- Example distribution:



\* This definition is not precise, and will later be refined.

# Modeling

- Time to choose a model class!
- **Think** – what are good criteria for choosing?
- Let's try *threshold functions*:

$$H_{\text{thresh}} = \{h_{\theta}(x) = \text{sign}(x \geq \theta) : \theta \in \mathbb{R}\}$$

- **Q:** Why is this a good class?
- (**Remember:** we're assuming separability, but not any specific distribution)
- **A:** Because it's “just right” – intuitively:
  1. For any separable  $D$  there is some  $h \in H$  with  $L_S(h) = 0$
  2. And, it's not too large – no “redundant”  $h$ -s
- **Note:**  $|H| = \infty$ , suggesting that “large” is not the right measure (we'll return to this)

# Optimization

- Recall the ERM template
- To actually learn, need to implement for  $H_{\text{thresh}}$
- Plug  $h_{\theta}(x) = \text{sign}(x \geq \theta)$  into formulas

## ERM: (Empirical Risk Minimization)

- **Input:**

- Sample set  $S = \{(x_i, y_i)\}_{i=1}^m$
- Model class  $H$  (*candidate classifiers*)

- **Objective:**

classifier with lowest **expected** error

$$h^* = \operatorname{argmin}_{h \in H} \mathbb{E}_{(x,y) \sim D} [\mathbb{1}\{y \neq h(x)\}]$$

- **Return:**

classifier with lowest **empirical** error

$$\hat{h} = \operatorname{argmin}_{h \in H} \frac{1}{m} \sum_{i=1}^m \mathbb{1}\{y_i \neq h(x_i)\}$$

# Optimization

- Recall the ERM template
- To actually learn, need to implement for  $H_{\text{thresh}}$
- Plug  $h_{\theta}(x) = \text{sign}(x \geq \theta)$  into formulas
- We will call  $\theta$  the **parameter(s)** of  $h_{\theta}$
- Optimizing over  $h \equiv$  optimizing over  $\theta$
- This is an example of model class *structure*, which we will see is helpful (and necessary)
- Can we solve ERM now?

## ERM: (Empirical Risk Minimization)

- **Input:**

- Sample set  $S = \{(x_i, y_i)\}_{i=1}^m$
- Model class  $H$  (*candidate classifiers*)

- **Objective:**

classifier with lowest expected error  
 $\theta^* = \underset{\theta \in \mathbb{R}}{\operatorname{argmin}} \mathbb{E}_{(x,y) \sim D} [\mathbb{1}\{y \neq \text{sign}(x \geq \theta)\}]$

- **Return:**

classifier with lowest empirical error

$$\hat{\theta} = \underset{\theta \in \mathbb{R}}{\operatorname{argmin}} \frac{1}{m} \sum_{i=1}^m \mathbb{1}\{y_i \neq \text{sign}(x_i \geq \theta)\}$$

# Optimization

## Simple algorithm:

1. Sort  $x_i$ -s
- 1b. Validate separability
2. Find largest  $x_i$  with  $y_i = -1$ ; call it  $x_{i^*}$
3. Return  $\hat{\theta} = x_{i^*}$

- **Congrats!** We just implemented our first ERM
- What do we need to prove?
- (Remember the algorithm is designed for *threshold classifiers on separable data*)

## ERM: (Empirical Risk Minimization)

- **Input:**

- Sample set  $S = \{(x_i, y_i)\}_{i=1}^m$
- Model class  $H$  (*candidate classifiers*)

- **Objective:**

classifier with lowest expected error  
$$\theta^* = \operatorname{argmin}_{\theta \in \mathbb{R}} \mathbb{E}_{(x,y) \sim D} [\mathbb{1}\{y \neq \text{sign}(x \geq \theta)\}]$$

- **Return:**

classifier with lowest empirical error

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \mathbb{R}} \frac{1}{m} \sum_{i=1}^m \mathbb{1}\{y_i \neq \text{sign}(x_i \geq \theta)\}$$

# Optimization

## Simple algorithm:

1. Sort  $x_i$ -s  $O(m \log m)$
- 1b. Validate separability  $O(m)$
2. Find largest  $x_i$  with  $y_i = -1$ ; call it  $x_{i^*}$   $O(m)$
3. Return  $\hat{\theta} = x_{i^*}$

- **Congrats!** We just implemented our first ERM
- What do we need to prove?
- (Remember the algorithm is designed for *threshold classifiers on separable data*)

## 1. Runtime: $O(m \log m)$ – great!

- Recall  $m = |S|$
- But  $H$  is also “input”, and  $|H| = \infty!$   
*How can this be?*

## 2. Correctness: $L_S(h_{\hat{\theta}}) = 0$

- Easy to show.
- **Not only solution** (what are others?)

## 3. Generalization: $L_D(h_{\hat{\theta}}) = ?$

- How can we relate  $L_S$  and  $L_D$ ?
- Enter *statistics*.

# Statistics I

- Algorithm returns  $\hat{\theta}$  with zero *empirical* error,  $L_S(h_{\hat{\theta}}) = 0$
- But we actually care about its *expected* error,  $L_D(h_{\hat{\theta}})$
- **Q:** What does  $L_D(h_{\hat{\theta}})$  depend on? What can cause it to vary?
- **A:**  $m$
- We expect  $L_D$  to go down with  $m$  – but at what rate?
- **Let's simulate!**
- Empirically, we can see that  $L_D \approx \frac{1}{m}$
- Can we run this simulation on real problems? (or – where did we cheat?)
- This is where theory comes in: bound rate *analytically*

# Statistics II

- Our algorithm takes in  $S$  and returns  $\hat{\theta}$
- But different  $S \Rightarrow$  different  $\hat{\theta} \Rightarrow$  different  $L_D(\hat{h})!$  (even for fixed  $m$ )
- Output of ERM is a **random variable** (because it's a function of the random variable  $S$ )
- How much variance in performance should we expect?
- **Let's simulate!**
- **Q:** So how can we tell if our algorithm is “good”?
- **A:** Use **probabilistic criteria** (w.r.t.  $S \sim D^m$ )
- **Corollary:**


Performance on a single dataset has less meaning than we'd like to believe



# Statistics III

- Let's see this in action.
- **Claim:**  $P_{S \sim D^m}(L_D(\hat{h}) > \epsilon) \leq 2e^{-\epsilon m}$
- **Proof:** [on board; UML book Sec. 6.1, p67]
- **Corollary:**  $\epsilon \geq \frac{1}{m} \log\left(\frac{2}{\delta}\right)$

# Statistics III



Let  $a^*$  be a threshold such that the hypothesis  $h^*(x) = \text{sign}(x \geq a^*)$  achieves  $L_D(h^*) = 0$ . Let  $a_0 < a^* < a_1$  be two points as in the figure. Formally,  $P_{D^m}(x \in (a_0, a^*)) = P_{D^m}(x \in (a^*, a_1)) = \epsilon$ .

Given a set  $S$ , define  $b_0 = \max\{x_i : y_i = -1\}$  and  $b_1 = \min\{x_i : y_i = 1\}$

If  $b_S$  is a threshold for ERM hypothesis  $h_S$ , then  $b_S \in (b_0, b_1)$

A sufficient condition for  $L_D(h_S) \leq \epsilon$  is:  $b_0 \geq a_0$  and  $b_1 \leq a_1$ . Or:

$P_{S \sim D^m}(L_D(h_S) > \epsilon) \leq P_{S \sim D^m}(b_0 < a_0 \text{ or } b_1 > a_1) \leq P_{S \sim D^m}(b_0 < a_0) + P_{S \sim D^m}(b_1 > a_1)$ , where the last inequality is by union bound.

Now,  $b_0 < a_0$  iff all samples in  $S$  are not in  $(a_0, a^*)$ , whose prob mass is  $\epsilon$ , so:

$P_{S \sim D^m}(b_0 < a_0) = P_{S \sim D^m}(\forall (x, y) \text{ in } S, x \notin (a_0, a^*)) = (1 - \epsilon)^m \leq e^{-\epsilon m}$ .

In the same way,  $P_{S \sim D^m}(b_1 > a_1) \leq e^{-\epsilon m}$ . Adding both, we get the desired bound.

# Discussion

- **Recall:** our algorithm sets  $\hat{\theta}$  to be the largest negative example
- **Observation:** just as good – any point between *largest negative* and *smallest positive*
- **Conclusion:** there can be multiple possible ERM rules!


*largest negative, smallest positive, mean, median, convex combination, random in interval, ...*

- **Ask yourself:**
  - Are some ERM rules better than others for some  $D$ ?
  - Are some ERM rules better than others for all  $D$ ?
  - Do your answers implicitly rely on additional assumptions?
  - Are your claims statistical (*observed*), distributional (*latent*), or both?
  - How would you validate your claims?
  - Actually – what does “better” even mean? Be precise!

## Case II: Overlapping Classes


# Overlapping classes

- **Hand-wavy:** continuous “middle” region where labels are uncertain
- Some example distributions:



# Overlapping classes

- **Hand-wavy:** continuous “middle” region where labels are uncertain
- Some example distributions:



- Separable is special case (=no overlap)
- (Our new algorithm should be “backward-compatible”)
- But in general, no  $\theta^*$ -like object
- (**Remember:** we only observe samples with binary  $y$ )


# Modeling

- What would be a good model class now?
- How about... thresholds?
- **Recall:**

$$H_{\text{thresh}} = \{h_{\theta}(x) = \text{sign}(x \geq \theta) : \theta \in \mathbb{R}\}$$

- **(Sneak-peak):** will work out well
- But with *one big difference*:

**minimal error can be strictly positive**



# Optimization


Recall our previous (separable) ERM algorithm:

## Simple algorithm:

1. Sort  $x_i$ -s
- 1b. Validate separability
2. Find largest  $x_i$  with  $y_i = -1$ ; call it  $x_{i^*}$
3. Return  $\hat{\theta} = x_{i^*}$

- **Q:** Will it still work?
- **A:** no – it breaks! (consider mixture of Gaussians)

➤ **Take away:** be careful in what you assume





# Optimization

Recall our previous (separable) ERM algorithm:


## Simple algorithm:

1. Sort  $x_i$ -s
- 1b. Validate separability
2. Find largest  $x_i$  with  $y_i = -1$ ; call it  $x_{i^*}$
3. Return  $\hat{\theta} = x_{i^*}$

- **Q:** Will it still work?
- **A:** no – it breaks! (consider mixture of Gaussians)

➤ **Take away:** be careful in what you assume

- Nonetheless, is there an underlying principal we can salvage?
- Recall:  $\hat{\theta} = x_{i^*}$
- Claim:  $x_{i^*}$  was useful because it was an **informative summary** of the data
- Let's apply this principle to our current problem



# Optimization

- Let's try again:

## Revised algorithm:

1. Sort  $x_i$ -s
2. For each  $x_i$ , compute empirical error  
“as if”  $x_i$  was the threshold
3. Find example with lowest error; call it  $x_{i^*}$
4. Return  $\hat{\theta} = x_{i^*}$

- **Runtime:**  $O(m \log m)$  (backward and forward counting passes)
- **Correctness:**  $\hat{h} = h_{\hat{\theta}} = \operatorname{argmin}_{h \in H} L_S(h)$  (easy to show)

- Nonetheless, is there an underlying principle we can salvage?
- Recall:  $\hat{\theta} = x_{i^*}$
- Claim:  $x_{i^*}$  was useful because it was an **informative summary** of the data
- Let's apply this principle to our current problem

# Optimization

- Let's try again:

## Revised algorithm:

1. Sort  $x_i$ -s
2. For each  $x_i$ , compute empirical error  
“as if”  $x_i$  was the threshold
3. Find example with lowest error; call it  $x_{i^*}$
4. Return  $\hat{\theta} = x_{i^*}$


- **Runtime:**  $O(m \log m)$
- **Correctness:**  $\hat{h} = h_{\hat{\theta}} = \operatorname{argmin}_{h \in H} L_S(h)$

- **Congrats!** We just implemented a more general ERM algorithm.
- Note that the learned  $\hat{h}$  is a very simple function of the data.
- **Hand-wavy:** this in effect “removes” the runtime dependence on  $H$ .
- This principle appears in many learning algorithms.
- Challenge is often in finding (and appropriately combining) the “important” elements of the data.

# Statistics

- **Observation:** minimal empirical error can be *strictly positive* (e.g., when data is not separable)
- Imagine we ran our ERM and got an (empirical) error of  $L_S(\hat{h}) = 0.1$ .
- **Q:** Should we be happy?
- **A:** It's complicated.
- Our principle concern is **generalization**.
- But we can't compute it...


**generalization:**


$$L_D(\hat{h})$$

# Statistics

- **Observation:** minimal empirical error can be *strictly positive* (e.g., when data is not separable)
  - Imagine we ran our ERM and got an (empirical) error of  $L_S(\hat{h}) = 0.1$ .
  - **Q:** Should we be happy?
  - **A:** It's complicated.
- 
- Our principle concern is **generalization**.
  - But we can't compute it...
  - **Observations:**
    1. We can't just assume empirical error  $\approx$  expected error

**generalization:**


$$L_D(\hat{h}) \stackrel{?}{\gtrless} L_S(\hat{h}) \geq 0$$

# Statistics

- **Observation:** minimal empirical error can be *strictly positive* (e.g., when data is not separable)
- Imagine we ran our ERM and got an (empirical) error of  $L_S(\hat{h}) = 0.1$ .
- **Q:** Should we be happy?
- **A:** It's complicated.


- Our principle concern is **generalization**.
- But we can't compute it... nor a bound.

- **Observations:**

1. We can't just assume empirical error  $\approx$  expected error
2. "Good" is *relative* (to optimal classifier  $h^* = \operatorname{argmin}_{h \in H} L_D(h)$ )

## generalization:

separable:

$$0 = L_D(h^*) \leq L_D(\hat{h}) \stackrel{?}{\gtrless} L_S(\hat{h}) \geq 0$$


# Statistics

- Trivial (but helpful!) *error decomposition*:


$$\begin{array}{ccccc} L_D(\hat{h}) & = & L_D(h^*) & + & L_D(\hat{h}) - L_D(h^*) \\ \text{generalization} & & \text{approximation} & & \text{estimation} \\ & & \text{error} & & \text{error} \end{array}$$

- **Approximation** = “how good  $H$  is for  $D$ ” (independent of  $S$ )
- **Estimation** = “given  $H$ , how good is  $S$  for  $D$ ” (empirical error as estimate of expected)
- **Conclusion**: model class  $H$  is useful if:
  - We believe it can approximate the data well [*modeling; approximation*]
  - We can associate empirical and expected errors [*statistics; estimation*]
  - We know we can compute ERM [*optimization; estimation*]
- **Remember**: The only thing we can compute (exactly) is  $L_S(\hat{h})$ ! **But where is it?**

# Statistics

$$\underbrace{L_D(\hat{h})}_{\text{generalization}} = \underbrace{L_D(h^*)}_{\text{approximation error}} + \underbrace{L_D(\hat{h}) - L_D(h^*)}_{\text{estimation error}}$$

- **Intuition:** as we increase the “size” of  $H$  (keeping  $m$  fixed):
  - larger  $H$  -> less bias (“assumptions”) -> lower approx. error -> better generalization






# Statistics

$$\underbrace{L_D(\hat{h})}_{\text{generalization}} = \underbrace{L_D(h^*)}_{\text{approximation error}} + \underbrace{L_D(\hat{h}) - L_D(h^*)}_{\text{estimation error}}$$


- **Intuition:** as we increase the “size” of  $H$  (keeping  $m$  fixed):
  - larger  $H$  -> less bias (“assumptions”) -> lower approx. error -> better generalization
  - larger  $H$  -> worse estimation -> higher error -> worse generalization



# Statistics

$$\underbrace{L_D(\hat{h})}_{\text{generalization}} = \underbrace{L_D(h^*)}_{\text{approximation error}} + \underbrace{L_D(\hat{h}) - L_D(h^*)}_{\text{estimation error}}$$

- **Intuition:** as we increase the “size” of  $H$  (keeping  $m$  fixed):
  - larger  $H$   $\rightarrow$  less bias (“assumptions”)  $\rightarrow$  lower approx. error  $\rightarrow$  better generalization
  - larger  $H$   $\rightarrow$  worse estimation  $\rightarrow$  higher error  $\rightarrow$  worse generalization



- This is called the **bias-complexity tradeoff**, we will later see why.


# Discussion

- **Recall:** we assume data is *jointly* sampled  $(x, y) \sim D_{XY}$
- Classification focuses on  $p(y|x)$  (where  $p(x, y) = p(x)p(y|x)$ )
- But sometimes it helps to consider  $p(x|y)$ , such as when we:
  - Have assumptions on data distribution
  - Directly model the distribution
- (Think of the distribution “types” we discussed today)
- Above quantities are related through *Bayes rule*:
$$p(y|x) = p(x|y)p(y)/p(x)$$
- Learning  $p(x|y)$  is called **generative learning** – we will return to this later in the course.

# Case III: Interleaved Classes

# Interleaved classes

- **Hand-wavy:** more than one  $D_{X|Y=y}$  for each class  $y$
- Example distributions:



- **Note:** effectively, unbounded "chunks" + non-separable = no assumptions
- This is where we're aiming!

# Modeling


$$L_D(\hat{h}) = L_D(h^*) + L_D(\hat{h}) - L_D(h^*)$$

- **Q:** Are thresholds still a good choice?
- **A:** **Yes!** Because they are statistically well-behaved (low *estimation error*)
- **A:** **No!** Because may not be best alternative (high *approximation error*)
- Ideas for better class?
- **Union of (half) intervals:**

$$h_I(x) = \mathbb{1}\{x \in \cup_j (a_j, b_j)\}, \quad I = \{(a_j, b_j)\}_{j=1}^d$$


- **On separable data:** approximation error = 0 (easy to show)
- But...

$D$



$|S| = 7$

$\hat{h}$



$L_S(\hat{h})$


$L_D(\hat{h})$

0

$\approx L_D(h^*)$

$|S| = 14$

$\hat{h}$




0

worse

$|S| = 24$

$\hat{h}$



0


even worse

# Statistics

- **Conclusion:** more data is... *worse*? What's going on here?
- **Observation:** with enough intervals ( $d$ ) – can always get zero empirical error!
- Models with larger  $d$  are more *expressive*
- We will think of *expressivity* as controlled by model class *complexity*
- For intervals, complexity = number of intervals
- **Notice:** as we added more data, we also increased model complexity
- Our example shows: increasing complexity can hurt expected error
- This is a fundamental concept in learning called ***overfitting***
- (Insufficient complexity is called *underfitting*)



$D$



$|S| = 7$

$\hat{h}$



$L_S(\hat{h})$


$L_D(\hat{h})$

0

worse

$|S| = 14$

$\hat{h}$




0

$\approx L_D(h^*)$

$|S| = 24$

$\hat{h}$




0

worse

# Statistics

- Let's generalize this intuition. For fixed data size  $m$ :




- Recall that we would like to minimize *expected* error.
- Plot suggests that we should **control complexity**. But how? Think:
  - What can we compute?
  - What *can't* we compute?

# Statistics

- In our example, model complexity increased when *more examples* were observed.
- What role does data size  $m$  play in over/under-fitting?

# Statistics

- In our example, model complexity increased when *more examples* were observed.
- What role does data size  $m$  play in over/under-fitting?




- Data and complexity are **complementary** in their effect on performance.
- **Q:** What happens if we restrict model complexity?

# Statistics

- Assume we restrict model complexity (e.g., use at most  $d = 5$  intervals)
- How does additional data effect performance now?

# Statistics

- Assume we restrict model complexity (e.g., use at most  $d = 5$  intervals)
- How does additional data effect performance now?






- **Q:** Will they always converge?
- **A:** Technically – yes (as  $m \rightarrow \infty$ ), but helpful only if this happens **polynomially-fast**.
- **Important:** not all model classes are *learnable* ( $\approx$  polynomial convergence)

# Statistics

- This is all in **theory** (and on average!) – what happens in **practice**?
- Typically, we have a finite data set (fixed  $m$ ), and can choose the model class (complexity).
- Three possible scenarios:

# Statistics

- This is all in **theory** (and on average!) – what happens in **practice**?
- Typically, we have a finite data set (fixed  $m$ ), and can choose the model class (complexity).
- Three possible scenarios:



- **Remember:** expect smooth behavior only when averaging – single datasets are *noisy*



# Optimization

- **Separable**: easy (won't go into details)
- **Non-separable**: hard
- **Intuition** –ERM as discrete optimization: (fix  $d$ )
  - Our previous approaches constructed  $\hat{\theta}$  from data points
  - Can think of  $d$  intervals as partition  $\{x_i\}_{i=1}^m$  into  $d$  (consecutive) subsets
  - This means we now need to find best *subset*
  - This (plus a discrete objective) results in a *hard combinatorial problem*
- This is typical of many learning problems
- We will therefore mostly deal with *approximate*\* learning algorithms
- **Remember this** when you reason about the performance of  $\hat{h}$

\* approximate in the algorithmic sense – don't confuse this with approximation error

# Modeling, revisited


- **Recall:** ERM implementations are typically class-specific
- Models in the class need to have “structure” that can be utilized (think parameterization)
- Unions of intervals are not ideal in that sense
- **Q:** Can we perhaps tweak thresholds to “behave” like intervals?
- **A:** Yes - using... **polynomials!** (?)

# Modeling, revisited

- **Recall:** ERM implementations are typically class-specific
- Models in the class need to have “structure” that can be utilized (think parameterization)
- Unions of intervals are not ideal in that sense
- **Q:** Can we perhaps tweak thresholds to “behave” like intervals?
- **A:** Yes - using... **polynomials!** (?)


**Thresh. polynomial classifier:**

$$h_{p,\theta}(x) = \text{sign}(p(x) > \theta)$$



# Modeling, revisited


- **Parameterization:**  $p(x) = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots w_dx^d$
- **Complexity:**  $d$  = number of intervals  $\leftrightarrow$  degree of polynomial



- **Polynomial thresholds are great!**
  - Generalize 1D thresholds – *intuitions and conclusions carry over*
  - **Algorithmically:** can solve (tractable) continuous optimization problem
  - **Statistically:** easy to analyze (well-studied objects)




# Statistics, revisited

- **Intuition:** *overfitting* happens when overly-complex models fit *noise*
- Easy to visualize for *polynomial regression* ( $\mathcal{Y} = \mathbb{R}$ ; e.g., *least squares*)



# Statistics, revisited

- **Intuition:** *overfitting* happens when overly-complex models fit *noise*
- Easy to visualize for *polynomial regression* ( $\mathcal{Y} = \mathbb{R}$ ; e.g., *least squares*)



- (For binary classification – imagine “bit-flip” noise)

# Discussion

# Recap

- ERM aims to minimize *empirical error*, but we care about *expected error*
- Our key quantity of interest is **generalization**
- Seems straightforward, but in practice – many, many subtleties to consider:
  - Approximation error and estimation error
  - Overfitting and underfitting
  - The relation between  $H$  and  $m$
  - The computational hardness of learning
  - What we can compute and what we cannot
  - How to define “good”
  - The randomness in... everything
  - ...
- **Beware of potential pitfalls!**



# Discussion

- Today we discussed several “types” of distributions
- In practice, we won’t know our type
- Can make assumptions, but this should be done with care
- (Remember: we *always* make *some* assumptions)
- **Alternative:** learning algorithms that are **distribution-free**
- **Pro:** guaranteed to work well for any  $D$
- **Con:** guarantees are always relative (e.g., vs.  $h^*$ )
- (In effect, choosing  $H$  + sufficing with relative measures = making an assumption)

# Next: beyond 1D data

- Consider a degree- $d$  polynomial over reals ( $z \in \mathbb{R}$ ):

$$p(z) = w_0 + w_1 z + w_2 z^2 + w_3 z^3 + \cdots w_d z^d$$

- Now define a *vector representation* of powers of  $z$ :

$$\mathbf{x} = (x_1, x_2, \dots, x_d) = (z, z^2, \dots, z^d)$$

- Notice that  $p(z)$  is linear in  $\mathbf{x}$ :

$$p(z) = b + \mathbf{w}^\top \mathbf{x}, \quad b = w_0, \mathbf{w} = (w_1, \dots, w_d)$$

- Transformed problem:** univariate inputs  
complex model class  $\Rightarrow$  multivariate inputs  
simple model class
- Next step:** generalize to arbitrary high-dimensional vector inputs,  $\mathbf{x} \in \mathbb{R}^d$
- This is called ***linear classification*** – a problem we will devote much of our time to.

# Next week

- **Classification** – three methods:

1. Similarity-based
2. Rule-based
3. Linear classifiers

