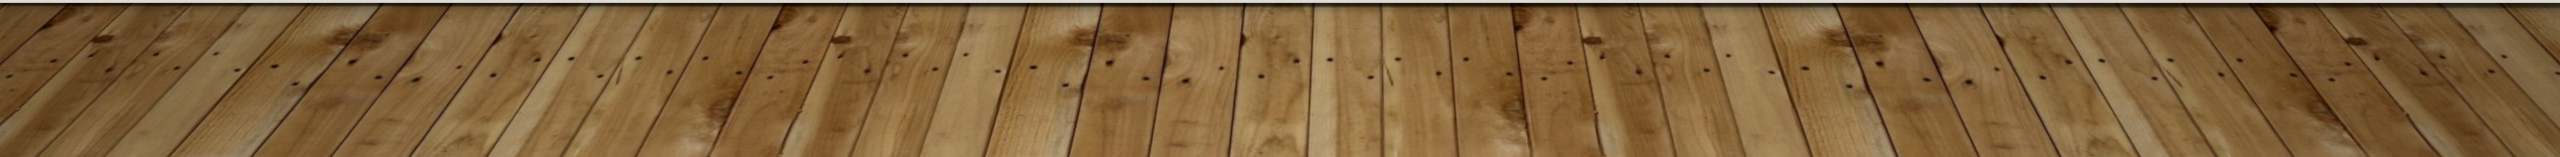


Introduction to Machine Learning (IML)

LECTURE #4: CLASSIFICATION – SVM IN DEPTH

236756 – 2024 SPRING – TECHNION

LECTURER: NIR ROSENFELD



Today

- Last lecture of first part of course (!)
- More SVM (!!)
- (Mostly optimization; some modeling)
- Hard SVM (separable)
 - short recap
 - finish up
- Soft SVM (non-separable)
- Dual SVM and kernels (non-linearity via linearity)

Recap

- Def: margin of hyperplane w :

$$\text{margin}(w; S) = \min_{i \in [m]} \frac{|w^\top x_i|}{\|w\|} := \gamma(w; S)$$

- **SVM looks for max margin classifier**



- Hard SVM works under **linear separability**:

$$\exists w \text{ s.t. } y_i \cdot w^\top x_i \geq 0 \quad \forall i \in [m]$$

(homogeneous case)

- Hyperplanes are scale-invariant, and so are margins:

$$\gamma(w; S) = \gamma(\alpha w; S) \quad \forall \alpha \in \mathbb{R}$$



Discussion


- **Original max-margin motivation:** $\operatorname{argmax}_w \gamma(w; S)$ s.t. $y_i \cdot w^\top x_i \geq 0 \quad \forall i \in [m]$
- **Final Hard SVM objective:** $w_{\text{H-SVM}} = \operatorname{argmin}_w \|w\|_2^2$ s.t. $y_i \cdot w^\top x_i \geq 1 \quad \forall i \in [m]$
- **Main insight:** increasing margin \equiv reducing norm (for fixed margin of size=1)

$$\begin{array}{ccc} \text{what we want} & \text{constrain as =1} & \\ \gamma(w; S) = \min_{i \in [m]} \frac{|w^\top x_i|}{\|w\|_2} & \equiv & \frac{1}{\|w\|_2} \\ & & \text{what we do} \end{array}$$

- Results in simple, convex objective with linear constraints (easy to optimize + unique solution)

Discussion

- **Final Hard SVM objective:** $w_{\text{H-SVM}} = \operatorname{argmin}_w \|w\|_2^2 \quad \text{s.t.} \quad y_i \cdot w^\top x_i \geq 1 \quad \forall i \in [m]$
- **Claim:** at least one example (but possibly more) “touches” margin (=constraint is tight)
- **Margin-touching examples are called “support vectors”:**
 - (hence the name - support vector *machines*)
 - removing “support” examples changes learned model
 - removing other examples does not
- These will pop up again later



Soft SVM

Hard SVM derivation

- Last week we derived a sequence of three optimization problems:

$$(1) \ w_1 = \operatorname{argmax}_w \frac{1}{\|w\|_2} \min_{i \in [m]} |w^\top x_i| \quad \text{s.t.} \quad y_i \cdot w^\top x_i \geq 0 \quad \forall i \in [m]$$


$$(2) \ w_2 = \operatorname{argmax}_w \frac{1}{\|w\|_2} \quad \text{s.t.} \quad \min_{i \in [m]} |w^\top x_i| = 1, \quad y_i \cdot w^\top x_i \geq 0 \quad \forall i \in [m]$$

$$(3) \ w_3 = \operatorname{argmin}_w \|w\|_2^2 \quad \text{s.t.} \quad y_i \cdot w^\top x_i \geq 1 \quad \forall i \in [m]$$

- Now it's time for #4 (and then #5!)


Removing the separability assumption

- **Q:** What happens if we use Hard SVM on non-separable data?



Removing the separability assumption


- **Q:** What happens if we use Hard SVM on non-separable data?
- **A:** Optimization problem has no feasible solutions.
(meaning constraints cannot be satisfied)
- **Solution** – use “soft” margin constraints:
penalize w by how much constraints are violated
- **Soft SVM:** penalize violations linearly (on average)



Removing the separability assumption

- **Q:** What happens if we use Hard SVM on non-separable data?
- **A:** Optimization problem has no feasible solutions.
(meaning constraints cannot be satisfied)
- **Solution** – use “soft” margin constraints:
penalize w by how much constraints are violated
- **Soft SVM:** penalize violations linearly (on average)


$$\begin{aligned} \operatorname{argmin}_{w \in \mathbb{R}^d, \xi \in \mathbb{R}^m} \quad & \lambda \|w\|_2^2 + \frac{1}{m} \sum_{i=1}^m \xi_i \\ \text{s.t.} \quad & y_i \cdot w^\top x_i \geq 1 - \xi_i \quad \forall i \in [m] \\ & \xi_i \geq 0 \end{aligned}$$



Soft SVM

- Penalties $\xi_i \geq 0$ are also called *slack variables*
- Separable data \Rightarrow optimal solution has $\xi_i = 0$
- Penalization allows points to be inside margin, or even on “wrong” side!
- λ – will get back to this

$$\begin{aligned} \operatorname{argmin}_{w \in \mathbb{R}^d, \xi \in \mathbb{R}^m} \quad & \lambda \|w\|_2^2 + \frac{1}{m} \sum_{i=1}^m \xi_i \\ \text{s.t.} \quad & y_i \cdot w^\top x_i \geq 1 - \xi_i \quad \forall i \in [m] \\ & \xi_i \geq 0 \end{aligned}$$




Soft SVM

$$\begin{aligned} \underset{w \in \mathbb{R}^d, \xi \in \mathbb{R}^m}{\operatorname{argmin}} \quad & \lambda \|w\|_2^2 + \frac{1}{m} \sum_{i=1}^m \xi_i \\ \text{s.t.} \quad & y_i \cdot w^\top x_i \geq 1 - \xi_i \quad \forall i \in [m] \\ & \xi_i \geq 0 \end{aligned}$$

- Constraints are ugly! (and not fun to optimize)
- Because we minimize over $\xi_i \geq 0$:
$$\xi_i = \begin{cases} 0 & \text{if } y_i \cdot w^\top x_i \geq 1 \\ 1 - y_i \cdot w^\top x_i & \text{if } y_i \cdot w^\top x_i < 1 \end{cases}$$
- Rewrite: $\xi_i = \max\{0, 1 - y_i \cdot w^\top x_i\}$

Soft SVM

$$\begin{aligned} \operatorname{argmin}_{w \in \mathbb{R}^d, \xi \in \mathbb{R}^m} \quad & \lambda \|w\|_2^2 + \frac{1}{m} \sum_{i=1}^m \xi_i \\ \text{s.t.} \quad & y_i \cdot w^\top x_i \geq 1 - \xi_i \quad \forall i \in [m] \\ & \xi_i \geq 0 \end{aligned}$$



- Constraints are ugly! (and not fun to optimize)
- Because we minimize over $\xi_i \geq 0$:
$$\xi_i = \begin{cases} 0 & \text{if } y_i \cdot w^\top x_i \geq 1 \\ 1 - y_i \cdot w^\top x_i & \text{if } y_i \cdot w^\top x_i < 1 \end{cases}$$
- Rewrite: $\xi_i = \max\{0, 1 - y_i \cdot w^\top x_i\} = z$


Soft SVM

- Plug in to get final **Soft SVM** formulation:

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

no more ξ !

$$\begin{aligned} \operatorname{argmin}_{w \in \mathbb{R}^d, \xi \in \mathbb{R}^m} \quad & \lambda \|w\|_2^2 + \frac{1}{m} \sum_{i=1}^m \xi_i \\ \text{s.t.} \quad & y_i \cdot w^\top x_i \geq 1 - \xi_i \quad \forall i \in [m] \\ & \xi_i \geq 0 \end{aligned}$$



- Constraints are ugly! (and not fun to optimize)
- Because we minimize over $\xi_i \geq 0$:
$$\xi_i = \begin{cases} 0 & \text{if } y_i \cdot w^\top x_i \geq 1 \\ 1 - y_i \cdot w^\top x_i & \text{if } y_i \cdot w^\top x_i < 1 \end{cases}$$
- Rewrite: $\xi_i = \max\{0, 1 - y_i \cdot w^\top x_i\} = z$

Hinge formulation

- We now see a template emerging:

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \boxed{\max\{0, 1 - y_i \cdot w^\top x_i\}} + \boxed{\lambda \|w\|_2^2}$$

loss **regularization**

Hinge formulation


- We now see a template emerging:

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \boxed{\max\{0, 1 - y_i \cdot w^\top x_i\}} + \lambda \|w\|_2^2$$

loss regularization

Loss:

- Penalizes model for being wrong
- SVM loss called *hinge loss* (= מפרק, ציר)



Hinge formulation

- We now see a template emerging:


$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \boxed{\max\{0, 1 - y_i \cdot w^\top x_i\}} + \lambda \|w\|_2^2$$

loss regularization






Loss:

- Penalizes model for being wrong
- SVM loss called *hinge loss* (= מפרק, צירק)
- Illustration is helpful, but can be misleading!
 - loss is plotted for one example – whereas real loss is average over many examples
 - loss appears to vary in a single "dimension" – but w can change in any direction in \mathbb{R}^d


single example:




multiple multidimensional examples: hinge loss as function of w




$$x_1 = (0.7, 0.7)$$



$$x_2 = (0.24, 0.97)$$



$$x_1, x_2$$




$$x_1, x_2, x_3 = (-0.83, 0.55)$$

Hinge formulation

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \boxed{\max\{0, 1 - y_i \cdot w^\top x_i\}} + \lambda \|w\|_2^2$$

loss regularization

- Recall our real goal was to minimize the 0/1 loss, which is difficult to optimize




Hinge formulation

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \boxed{\max\{0, 1 - y_i \cdot w^\top x_i\}} + \lambda \|w\|_2^2$$

loss **regularization**

- Recall our real goal was to minimize the 0/1 loss, which is difficult to optimize
- Hinge loss is a **tractable** alternative: it is continuous and **convex**, and upper bounds 0/1 loss
note: sum of convex = convex
- Replaces optimization over (discrete) **classifiers** h with optimization over (continuous) **score functions** f
- But tractability has its price:**
 - may suffer loss even if correct (i.e., when correct but by less than the margin)
 - if wrong, suffer linear loss (can be unbounded! vs. fixed loss of 1, as in 0-1 loss)
 - well-known weakness:** **outliers** (think – why?)



Retelling our story

- Once upon a time we wanted to minimize expected 0/1 loss:

$$\min_{h \in H} \mathbb{E}_{(x,y) \sim D} [\mathbb{1}\{y \neq h(x)\}]$$

- But this is statistically hard (no access to D),
so instead we tried to minimize the empirical 0/1 loss (ERM):

$$\operatorname{argmin}_{h \in H} \frac{1}{m} \sum_{i=1}^m \mathbb{1}\{y_i \neq h(x_i)\}$$

Retelling our story

- Once upon a time we wanted to minimize expected 0/1 loss:

$$\min_{h \in H} \mathbb{E}_{(x,y) \sim D} [\mathbb{1}\{y \neq h(x)\}]$$

- But this is statistically hard (no access to D),
so instead we tried to minimize the empirical 0/1 loss (ERM):

$$\operatorname{argmin}_{h \in H} \frac{1}{m} \sum_{i=1}^m \mathbb{1}\{y_i \neq h(x_i)\}$$

- But this is computationally hard (NP-hard discrete optimization),
so instead we ended up minimizing a proxy loss – the hinge loss:

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \boxed{\max\{0, 1 - y_i \cdot w^\top x_i\}} + \lambda \|w\|_2^2$$

loss

Retelling our story

- Once upon a time we wanted to minimize expected 0/1 loss:

$$\min_{h \in H} \mathbb{E}_{(x,y) \sim D} [\mathbb{1}\{y \neq h(x)\}]$$

- But this is statistically hard (no access to D),
so instead we tried to minimize the empirical 0/1 loss (ERM):

$$\operatorname{argmin}_{h \in H} \frac{1}{m} \sum_{i=1}^m \mathbb{1}\{y_i \neq h(x_i)\}$$

- But this is computationally hard (NP-hard discrete optimization),
so instead we ended up minimizing a proxy loss – the hinge loss:

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2 \quad \leftarrow \text{not ERM}$$

- **But oh no!** Turns out we're not minimizing the empirical error! **What are we minimizing?**

Regularization


$$\operatorname{argmin}_{w \in \mathbb{R}^d} \underbrace{\frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\}}_{\text{loss}} + \underbrace{\lambda \|w\|_2^2}_{\text{regularization}}$$

- The additive term $+\lambda \|w\|_2^2$ is called regularization – it *regularizes* solutions to have small norms
- Can be thought of as penalty on w -s with a large norm
- This is a way to inject prior knowledge, expressing a believe that low-norm w -s are “better”
- **Recall:** got $\|w\|$ by wanting max margin
- So in what sense are low-norm solutions better?

Hinge formulation

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \underbrace{\max\{0, 1 - y_i \cdot w^\top x_i\}}_{\text{loss}} + \underbrace{\lambda \|w\|_2^2}_{\text{regularization}}$$


- **Claim:** low-norm w -s provide better generalization
- (formal proof next week; for now – some intuition)
- λ controls trade-off between loss and regularization:
 - Small $\lambda \Rightarrow$ large $\|w\|$ (and lower loss)
 - Large $\lambda \Rightarrow$ small $\|w\|$ (and possibly worse loss)
- Can think of each λ as inducing hard constraint $\|w\| \leq B_\lambda$ for some B_λ (which we don't know!)
- A series $\lambda_1 > \lambda_2 > \dots > \lambda_N > \dots$ induces a nested hierarchy of models
- Can think of these as having **increased model complexity**, so choosing λ = choosing model class!



Hinge formulation

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \underbrace{\frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\}}_{\text{loss}} + \underbrace{\lambda \|w\|_2^2}_{\text{regularization}}$$

- **Conclusion:** tuning λ can control overfitting



- **Q1:** in a non-homogeneous model ($b \neq 0$), should we also regularize b ?
- (A1: no – b just shifts the data, and has no effect on overfitting)
- **Q2:** we saw (and used!) the fact that H is scale invariant; don't norms just change the scale?
- (A2: no – even though for a given w “changing the scale” is similar to “changing the norm”, this is not the same as considering a set of w -s with (at most) given norm)

Hinge formulation

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \underbrace{\max\{0, 1 - y_i \cdot w^\top x_i\}}_{\text{loss}} + \underbrace{\lambda \|w\|_2^2}_{\text{regularization}}$$



Regularization coefficient (λ):

- Soft SVM penalizes violations of margin constraints
- But now, must decide – what is more important:
large margin (min norm) –or– *small penalty* (constraint violations)?
- Determine **tradeoff** using λ
- As $\lambda \rightarrow 0$, constraints become “hard” constraints (recall ξ formulation)
- How to choose λ ? Later in course!

Beyond SVM

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \underbrace{\max\{0, 1 - y_i \cdot w^\top x_i\}}_{\text{loss}} + \underbrace{\lambda \|w\|_2^2}_{\text{regularization}}$$

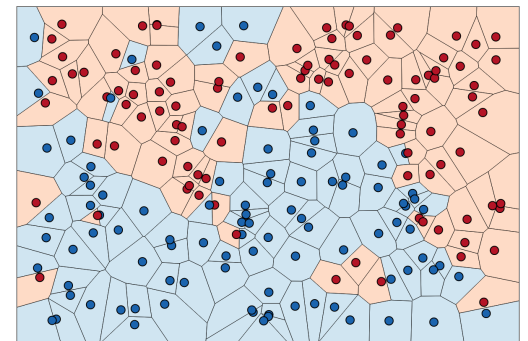

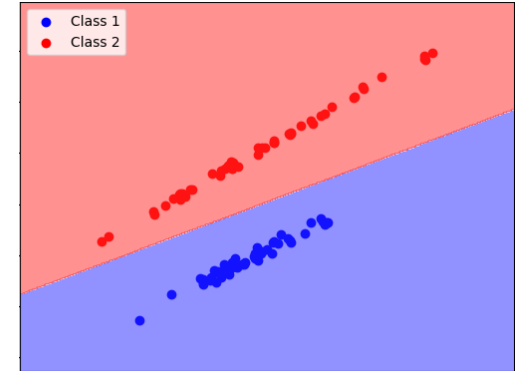
- Think of current objective as “template”
- Can use:
 - Other losses (logistic, exponential, ramp, ...)
 - Other regularization (L_1 , L_p , mixed, structured, ...)
- Each has its pros, cons, use cases, and derivations
- We will see some later in course



Kernels

Linear non-linearity


- SVMs are great!
- But a clear limitation is expressivity:
work only for **linear thresholds**
- Compare this to k-NN and decisions trees
- What can we do to make SVMs **more expressive**?



Linear non-linearity

- Recall 1D example from lecture #2:
- Can't get perfect accuracy with a linear threshold classifier!

$$h(x) = \text{sign}(f(x)), \quad f(x) = w_0 + w_1 x$$



Linear non-linearity


- Recall 1D example from lecture #2:
- Can't get perfect accuracy with a linear threshold classifier!
- But, can do this with polynomial threshold: (here, of degree=2)

$$h(x) = \text{sign}(q(x)), \quad q(x) = w_0 + w_1x + w_2x^2$$

- But a non-linear polynomial is also... linear! (?!
- Can rewrite as *linear* function of polynomial *representation*:




$$q(x) = f_w(\phi(x)) = w^\top \phi(x), \quad \phi(x) = (1, x, x^2)$$

- $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^{d'}$ is called a **feature mapping**
(these will pup-up again and again...)



Linear non-linearity

- Also works for higher dimensional *inputs*
- E.g., 2D example: (cannot be linearly separated in x)
- Possible useful feature mappings:
 1. **Polynomial:** $\phi(x) = (1, x_1, x_2, x_1x_2, x_1^2, x_2^2)$
 2. **Radial:** $\phi(x) = (r, \theta)$ (if data is centered)
- Sometimes we'll need higher-dimensional *outputs* (=polynomials)



Linear non-linearity


- **General Feature mappings:** $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^{d'}$
- Can make linear SVMs be non-linear!
- (any multivariate mapping ϕ works, not restricted to polynomials – more on this later!)
- **Recipe:**
 - choose ϕ (e.g., as some (multi-variate) polynomial)
 - apply $x'_i = \phi(x_i)$ to all x_i in the dataset
 - use the SVM algorithm to learn on modified data (x'_i, y_i)
- **Result:** a linear classifier $w^\top \phi(x)$ that behaves like a non-linear classifier on x !

Linear non-linearity

- Using a non-linear representation ϕ is a *modeling choice*
- Will it work well? How can we tell?
- As always, consult our **three pillars**:
 - **Modeling**: choose ϕ if you believe it (linearly) separates the data well (approx.)
 - **Statistics**: (will see next week)
 - **Optimization**: up next!

Linear non-linearity


- Consider ϕ to be a polynomial of degree k and dimension d
- **Q:** What is the dimension d' of the induced linear classifier?
- **A:** $d' = |\phi(x)| = \text{number of monomials} = \binom{k + d - 1}{k}$
- **Examples:**
 - $d = 3, \quad k = 3 \quad \rightarrow \quad d' = 10$
 - $d = 30, \quad k = 3 \quad \rightarrow \quad d' = 4960$
 - $d = 30, \quad k = 30 \rightarrow d' \approx 6 \times 10^{17}$
- **Problem:** d' grows *fast*, making SVM computationally demanding
- (even just storing $\phi(x)$ can be challenging!)
- **Goal for this section:** make SVM work for arbitrary ϕ – **even huge!**
- **Solution:** transform into *another* optimization problem (yes, again)



Quadratic constrained optimization

- **Gentle start:** Hard-SVM with one example (so single constraint):

$$\min_{w \in \mathbb{R}} w^2 \text{ s.t. } xyw \geq 1$$



- **Quadratic + linear constraints:** global optimum \iff gradient = 0 (or at boundary)
- **Hand-wavy:** local improvement \Rightarrow global optimum \Rightarrow exact and efficient optimization
- (more on this – week 7)

Duality – simple case

- 1D Hard-SVM: $\min_{w \in \mathbb{R}} w^2$ s.t. $w \geq \frac{1}{xy} = b$
- **Lagrange method:** express constraint within objective

$$\text{Lagrangian: } L(w, \alpha) = w^2 - \alpha(w - b), \alpha \in \mathbb{R}_+$$

Duality – simple case

- 1D Hard-SVM: $\min_{w \in \mathbb{R}} w^2$ s.t. $w \geq \frac{1}{xy} = b$
- **Lagrange method:** express constraint within objective

$$\text{Lagrangian: } L(w, \alpha) = w^2 - \alpha(w - b), \alpha \in \mathbb{R}_+$$

- “New” objective: solve $\min_{w \in \mathbb{R}} \max_{\alpha \in \mathbb{R}_+} L(w, \alpha)$
- **Claim:** problems are equivalent
- **Intuition:**
 - min and max “compete”
 - $\alpha \geq 0$ ensures w satisfies constraints
 - “Pushes” violating w towards edge of constraints

- **Why equivalence?**

- Case I: $w < b \rightarrow w - b < 0$
 $\rightarrow \max_{\alpha} -\alpha(w - b) = \infty$
 \min_w will never choose such w
- Case II: $w > b \rightarrow w - b > 0$
 $\rightarrow \max_{\alpha} -\alpha(w - b) = 0$
min doesn't mind,
objective recovered: $L(w, \alpha) = w^2$
- Case III: $w = b \rightarrow$ any α works
objective recovered: $L(w, \alpha) = w^2$

Duality – general case

- Hard-SVM: $\operatorname{argmin}_{w \in \mathbb{R}^d} \|w\|_2^2 \quad \text{s.t.} \quad y_i w^\top x_i \geq 1 \quad \forall i \in [m]$
- **Lagrangian:** $L(w, \alpha) = \|w\|_2^2 - \sum_{i=1}^m \alpha_i (y_i w^\top x_i - 1), \alpha \in \mathbb{R}_+^m$ (multiple constraints => sum)

Duality – general case

- Hard-SVM: $\operatorname{argmin}_{w \in \mathbb{R}^d} \|w\|_2^2 \quad \text{s.t.} \quad y_i w^\top x_i \geq 1 \quad \forall i \in [m]$
- **Lagrangian:** $L(w, \alpha) = \|w\|_2^2 - \sum_{i=1}^m \alpha_i (y_i w^\top x_i - 1), \alpha \in \mathbb{R}_+^m$ (multiple constraints => sum)
- Primal objective: $\min_{w \in \mathbb{R}^d} \max_{\alpha \in \mathbb{R}_+^m} L(w, \alpha) \quad ? \quad \max_{\alpha \in \mathbb{R}_+^m} \min_{w \in \mathbb{R}^d} L(w, \alpha) \quad (\text{dual objective})$
- **Dubious move:** swap min \leftrightarrow max

Duality – general case

- Hard-SVM: $\operatorname{argmin}_{w \in \mathbb{R}^d} \|w\|_2^2 \quad \text{s.t.} \quad y_i w^\top x_i \geq 1 \quad \forall i \in [m]$
- **Lagrangian:** $L(w, \alpha) = \|w\|_2^2 - \sum_{i=1}^m \alpha_i (y_i w^\top x_i - 1), \alpha \in \mathbb{R}_+^m$ (multiple constraints => sum)
- Primal objective:
$$\min_{w \in \mathbb{R}^d} \max_{\alpha \in \mathbb{R}_+^m} L(w, \alpha) \geq \max_{\alpha \in \mathbb{R}_+^m} \min_{w \in \mathbb{R}^d} L(w, \alpha) \quad (\text{dual objective})$$
- **Dubious move:** swap min \leftrightarrow max
- In general, $\min \max \geq \max \min$ (“max min inequality”; see wiki)

Duality – general case

- Hard-SVM: $\operatorname{argmin}_{w \in \mathbb{R}^d} \|w\|_2^2 \quad \text{s.t.} \quad y_i w^\top x_i \geq 1 \quad \forall i \in [m]$
- **Lagrangian:** $L(w, \alpha) = \|w\|_2^2 - \sum_{i=1}^m \alpha_i (y_i w^\top x_i - 1), \alpha \in \mathbb{R}_+^m$ (multiple constraints => sum)
- Primal objective:
$$\min_{w \in \mathbb{R}^d} \max_{\alpha \in \mathbb{R}_+^m} L(w, \alpha) = \max_{\alpha \in \mathbb{R}_+^m} \min_{w \in \mathbb{R}^d} L(w, \alpha) \quad (\text{dual objective})$$
- **Dubious move:** swap $\min \leftrightarrow \max$
- In general, $\min \max \geq \max \min$ (“max min inequality”; see wiki)
- **But:** **convex in w** (for fixed α) + **concave in α** (for fixed w) \Rightarrow **equality!**
- (aka minimax theorem; won’t prove)
- **Bonus:** lies at core of game theory (zero-sum games); adversarial learning, GANs.

Dual SVM

- **Dual objective:** $\max_{\alpha \in \mathbb{R}_+^m} \min_{w \in \mathbb{R}^d} \frac{1}{2} \|w\|_2^2 - \sum_{i=1}^m \alpha_i (y_i w^\top x_i - 1), \alpha \in \mathbb{R}_+^m$
- **Next step:** solve for optimal w (as function of α) by setting derivative = 0

$$\frac{\partial L}{\partial w} = w - \sum_i \alpha_i y_i x_i = 0 \rightarrow w = \sum_i \alpha_i y_i x_i \quad (w, x \in \mathbb{R}^d)$$

- Plugging into objective (and simplifying) gives:

$$\text{Dual (Hard) SVM: } \operatorname{argmax}_{\alpha \in \mathbb{R}_+^m} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j x_i^\top x_j$$

- Time to reap the fruits!

Dual SVM

- **Dual (Hard) SVM:**

$$\operatorname{argmax}_{\alpha \in \mathbb{R}_+^m} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j x_i^\top x_j$$

- **Q:** Where did w go?

Dual SVM

- **Dual (Hard) SVM:**

$$\operatorname{argmax}_{\alpha \in \mathbb{R}_+^m} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j x_i^\top x_j$$

- **Q:** Where did w go?
- **A:** Derivation reveals optimality constraint:

$$w = \sum_{i=1}^m \alpha_i y_i x_i \Rightarrow f_w(x) = w^\top x = \sum_{i=1}^m \alpha_i y_i x_i^\top x = f_\alpha(x)$$

- Classifier as *linear combination of data points* (special case of “representer theorem”)
- Dual SVM optimizes *directly over α*
- In practice, α is almost always sparse
(We know this! Think: only support vectors have $\alpha_i > 0$)

The kernel trick

- **Recall:** we want to work with a *feature mapping* $x \rightarrow \phi(x)$

learning objective: (*train time*)

$$\max_{\alpha \in \mathbb{R}_+^m} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j x_i^\top x_j$$

classifier: (*test time*)

$$f_\alpha(x) = \sum_{i=1}^m \alpha_i y_i x_i^\top x$$

- **Observation:** in dual, features appear only as *inner products* $x^\top x'$
- **Kernel trick:** given representation ϕ , replace inner product: $x^\top x' \rightarrow \phi(x)^\top \phi(x')$
- New *kernelized* objective and classifier:

$$\max_{\alpha \in \mathbb{R}_+^m} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j \phi(x_i)^\top \phi(x_j)$$

$$f_\alpha(x) = \sum_{i=1}^m \alpha_i y_i \phi(x_i)^\top \phi(x)$$

The kernel trick

- **Definition:** *kernel function* $K(x, x') = \phi(x)^\top \phi(x')$

- Can rewrite:

learning objective: (*train time*)

$$\max_{\alpha \in \mathbb{R}_+^m} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j K(x_i, x_j)$$

classifier: (*test time*)

$$f_\alpha(x) = \sum_{i=1}^m \alpha_i y_i K(x_i, x)$$

- To optimize (=train), need only *entries* of *kernel matrix* $K_{ij} = K(x_i, x_j) = \phi(x_i)^\top \phi(x_j)$
(kernel matrix can be *precomputed* from original features before training)
- To classify, need only kernel *queries* $K(x, x_i) = \phi(x)^\top \phi(x_i)$
(must be computed on the fly for all x_i for which $\alpha_i > 0$)
- **But:** kernels are useful if $\phi(x)^\top \phi(x')$ can be **computed efficiently** – let's see why!

Polynomial kernel

- Recall that a polynomial ϕ has d' that is exponential in k, d
- This means that computing $\phi(x)^\top \phi(x')$ directly is intractable

- Luckily – a polynomial kernel admits a simple closed-form solution:

$$K_{\text{poly}}(x, x'; k) = \phi(x)^\top \phi(x') = (x^\top x')^k$$

- (To see why, recall that $(a + b)^2 = a^2 + ab + b^2$; now expand to power of k)
- **Compute time:** $O(d)$

Kernels - examples

- There are many, many types of kernels
- Examples of some known kernels with compact formulation:

- Polynomials of degree = k : $K(x, x'; k) = (x^\top x')^k$
- Polynomials of degree $\leq k$: $K(x, x'; k) = (x^\top x' + 1)^k$ *(more in tirgul)*
- RBF/Gaussian (scale sigma): $K(x, x'; \sigma) = \exp\left\{-\frac{1}{2\sigma^2} \|x - x'\|_2^2\right\}$ *(next up)*
- Sigmoid: $K(x, x'; \eta, \nu) = \tanh(\eta x^\top x' + \nu)$

- Compute time for each of the above is $O(d)$
- Notice kernels have *parameters*

Handwritten notes in blue ink:

$\sigma \rightarrow \infty: e^0 = 1$
 $\sigma \rightarrow 0: e^{-\infty} = 0$
↑
"כ" / הפסד על ההיחזקים מהנך' <=

Well-defined kernels

- How do we know a given kernel is “valid”?
- For example, is this a kernel?

RBF kernel: $K(x, x'; \sigma) = \exp \left\{ -\frac{1}{\sigma^2} \|x - x'\|_2^2 \right\}$

- Two options to validate:
 1. Figure out what the underlying ϕ is; show $K(x, x') = \phi(x)^\top \phi(x')$
 2. Use **kernel algebra**
 3. Show kernel matrix is (semi)-positive-definite (PSD) – in week 7!

Well-defined kernels

- **Kernel composition rules:**

any of the following operations give valid kernels

1. $K(x, x') = x^\top x'$
2. $K(x, x') = x^\top A x', A \succeq 0$ (PSD)
3. $K(x, x') = cK_1(x, x'), c \geq 0$
4. $K(x, x') = K_1(x, x') + K_2(x, x')$
5. $K(x, x') = K_1(x, x') \cdot K_2(x, x')$
6. $K(x, x') = g(K_1(x, x'))$, g polynomial (with ≥ 0 coeffs.)
7. $K(x, x') = f(x)K(x, x')f(x')$, f is any function
8. $K(x, x') = \exp\{K_1(x, x')\}$

- Can use above composition rules to check validity
- And to create new kernels!

Claim: RBF kernel

$$K(x, x'; \sigma) = \exp\left\{-\frac{1}{\sigma^2} \|x - x'\|_2^2\right\}$$

is a valid kernel.

Proof: using composition rules

1. $K_1(x, x') = x^\top x' \quad (1)$

2. $K_2(x, x') = \frac{2}{\sigma^2} K_1(x, x') \quad (2)$

3. $K_3(x, x') = e^{K_2(x, x')} = e^{\frac{2x^\top x'}{\sigma^2}} \quad (8)$

4. $K_4(x, x') = e^{\frac{-x^\top x}{\sigma^2}} K_3(x, x') e^{\frac{-x'^\top x'}{\sigma^2}} \quad (7)$

$$\begin{aligned} &= e^{\frac{-x^\top x}{\sigma^2}} e^{\frac{2x^\top x'}{\sigma^2}} e^{\frac{-x'^\top x'}{\sigma^2}} \\ &= e^{\frac{-x^\top x + 2x^\top x' - x'^\top x'}{\sigma^2}} \\ &= e^{\frac{-\|x - x'\|_2^2}{\sigma^2}} = K_{\text{RBF}}(x, x') \end{aligned}$$

More kernels!

- Kernels for *structured inputs*:
 - graphs
 - trees
 - sets
 - strings
 - text
 - molecules
 - ...
- See lists, surveys, packages, ...

Contents

1. [Kernel Methods](#)
 1. [The Kernel Trick](#)
 2. [Kernel Properties](#)
 3. [Choosing the Right Kernel](#)
2. [Kernel Functions](#)
 1. [Linear Kernel](#)
 2. [Polynomial Kernel](#)
 3. [Gaussian Kernel](#)
 4. [Exponential Kernel](#)
 5. [Laplacian Kernel](#)
 6. [ANOVA Kernel](#)
 7. [Hyperbolic Tangent \(Sigmoid\)](#)
 8. [Rational Quadratic Kernel](#)
 9. [Multiquadric Kernel](#)
 10. [Inverse Multiquadric Kernel](#)
 11. [Circular Kernel](#)
 12. [Spherical Kernel](#)
 13. [Wave Kernel](#)
 14. [Power Kernel](#)
 15. [Log Kernel](#)
 16. [Spline Kernel](#)
 17. [B-Spline Kernel](#)
 18. [Bessel Kernel](#)
 19. [Cauchy Kernel](#)
 20. [Chi-Square Kernel](#)
 21. [Histogram Intersection Kernel](#)
 22. [Generalized Histogram Intersection](#)
 23. [Generalized T-Student Kernel](#)
 24. [Bayesian Kernel](#)
 25. [Wavelet Kernel](#)
3. Source code

http://crsouza.com/2010/03/17/kernel-functions-for-machine-learning-applications/#kernel_functions

RBF primal representation

- **Q:** What is the feature representation $\phi(x)$ underlying RBF?

RBF primal representation

- **Q:** What is the feature representation $\phi(x)$ underlying RBF?

- **A:**
$$\begin{aligned} \exp\left(-\frac{1}{2}\|\mathbf{x} - \mathbf{x}'\|^2\right) &= \exp\left(\frac{2}{2}\mathbf{x}^\top \mathbf{x}' - \frac{1}{2}\|\mathbf{x}\|^2 - \frac{1}{2}\|\mathbf{x}'\|^2\right) \\ &= \exp(\mathbf{x}^\top \mathbf{x}') \exp\left(-\frac{1}{2}\|\mathbf{x}\|^2\right) \exp\left(-\frac{1}{2}\|\mathbf{x}'\|^2\right) \\ &= \sum_{j=0}^{\infty} \frac{(\mathbf{x}^\top \mathbf{x}')^j}{j!} \exp\left(-\frac{1}{2}\|\mathbf{x}\|^2\right) \exp\left(-\frac{1}{2}\|\mathbf{x}'\|^2\right) \\ &= \underbrace{\left(\sum_{j=0}^{\infty} \sum_{\sum n_i=j} \exp\left(-\frac{1}{2}\|\mathbf{x}\|^2\right) \frac{x_1^{n_1} \cdots x_k^{n_k}}{\sqrt{n_1! \cdots n_k!}}\right)}_{\text{single entry in } \phi(x)} \underbrace{\exp\left(-\frac{1}{2}\|\mathbf{x}'\|^2\right) \frac{x_1'^{n_1} \cdots x_k'^{n_k}}{\sqrt{n_1! \cdots n_k!}}}_{\text{single entry in } \phi(x')} \end{aligned}$$
 (source: wikipedia)

- Basis expansion of *infinite* support (that is, $d' = |\phi(x)| = \infty$)
- Nonetheless, can still be learned using Dual SVM (aka *Kernel SVM*)
- Unrestricted complexity - can fit *any function* (*universal kernel*)
- Computationally, no worries – but **beware overfitting!**

Modeling

- Kernels define a *similarity measure* via inner products $\phi(x)^\top \phi(x')$
- Compare:


linear kernel:

$$K(x, x') = x^\top x'$$

RBF kernel:

$$K(x, x') = \exp \left\{ -\frac{1}{\sigma^2} \|x - x'\|_2^2 \right\}$$

- **Dual – similarity**; **Primal – distance** (from hyperplane)
- (well separated (*distance*) \approx clustered together (*similarity*))
- This provides flexible and expressive modeling power:
instead of thinking in terms of distances, can think in terms of similarities
- (kNN anyone?)
- Relation between **distances** and **similarities** is fundamental in learning!



Discussion

- SVM as major milestone in ML history
- Had all you could ask for:
 1. Computationally tractable
(fast convergence to optimum)
 2. Highly expressive (with kernels)
 3. Statistical guarantees (next week!)
- Prime example of “complete” learning approach, great to learn from.
- At time, state of the art performance
(today, still competitive baseline)
- “Just need the right kernel”
(vs: “just need the right neural architecture”)

“Give me a place to stand, and a lever long enough, and I shall move the earth”

- *Archimedes*

Next week(s)

- **Finished part I:** *supervised binary classification*
- **Next up – part II:** *the different aspects of learning*
 1. Statistics: generalization and PAC theory
 2. Modeling: model selection and evaluation
 3. Optimization: convexity, gradient descent
 4. Practical aspects and potential pitfalls
- (will mostly use SVM as use case)

