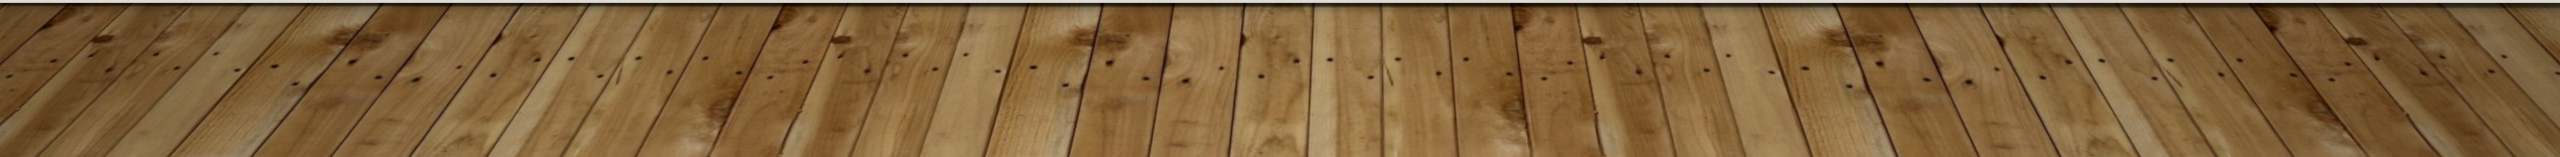


Introduction to Machine Learning (IML)

LECTURE #6: MODEL SELECTION

236756 – 2024 SPRING – TECHNION

LECTURER: NIR ROSENFELD



Today

- **Part II:** *the different aspects of learning*
 1. Statistics: generalization and PAC theory (finish up)
 2. Modeling: (today)
 - Model selection
 - Regularization
 - Evaluation and validation
 3. Optimization: convexity, gradient descent
 4. Practical aspects and potential pitfalls

PAC and VC – wrap up

Recall

- **Remember:** everything in learning is a random variable (sample set, learned model, performance, ...)

- **Definition:**

H is **PAC-learnable** if $\exists A, \exists m_H(\epsilon, \delta) \in \text{poly}\left(\frac{1}{\epsilon}, \frac{1}{\delta}\right)$

such that $\forall D$ (**realizable**) and $\forall \epsilon, \delta \in [0,1]$, if $m \geq m_H(\epsilon, \delta)$, then:

$$P_{S \sim D^m}(L_D(h_S) \geq \epsilon) \leq \delta$$

- For **finite model classes**, following bound holds:

$$P_{S \sim D^m}(L_D(h_S) \geq \epsilon) \leq |H|e^{-m\epsilon} \leq \delta$$

- **Interpretations:**

$$1. m \geq \frac{\log|H| + \log\frac{1}{\delta}}{\epsilon} \qquad 2. \epsilon \geq \frac{\log|H| + \log\frac{1}{\delta}}{m} \approx \frac{1}{m}$$

- **What about non-realizable?**

Recall

- **Remember:** everything in learning is a random variable (sample set, learned model, performance, ...)

- **Definition:**

H is **Agnostic-PAC-learnable** if $\exists A, \exists m_H(\epsilon, \delta) \in \text{poly}\left(\frac{1}{\epsilon}, \frac{1}{\delta}\right)$

such that $\forall D$ and $\forall \epsilon, \delta \in [0,1]$, if $m \geq m_H(\epsilon, \delta)$, then:

$$P_{S \sim D^m}(L_D(h_S) - L_D(h^*) \geq \epsilon) \leq \delta$$

- For **finite model classes**, got following bound:

$$P_{S \sim D^m}(L_D(h_S) - L_D(h^*) \geq \epsilon) \leq 2|H|e^{-2m\epsilon^2} \leq \delta$$

- **Interpretations:**

$$1. m \geq \frac{\log 2|H| + \log \frac{1}{\delta}}{2\epsilon^2} \quad 2. \epsilon \geq \sqrt{\frac{\log 2|H| + \log \frac{1}{\delta}}{2m}} \approx \frac{1}{\sqrt{m}} \quad (\text{vs. } \frac{1}{m} \text{ in realizable case})$$

- **What about infinite classes?**

Beyond finite classes

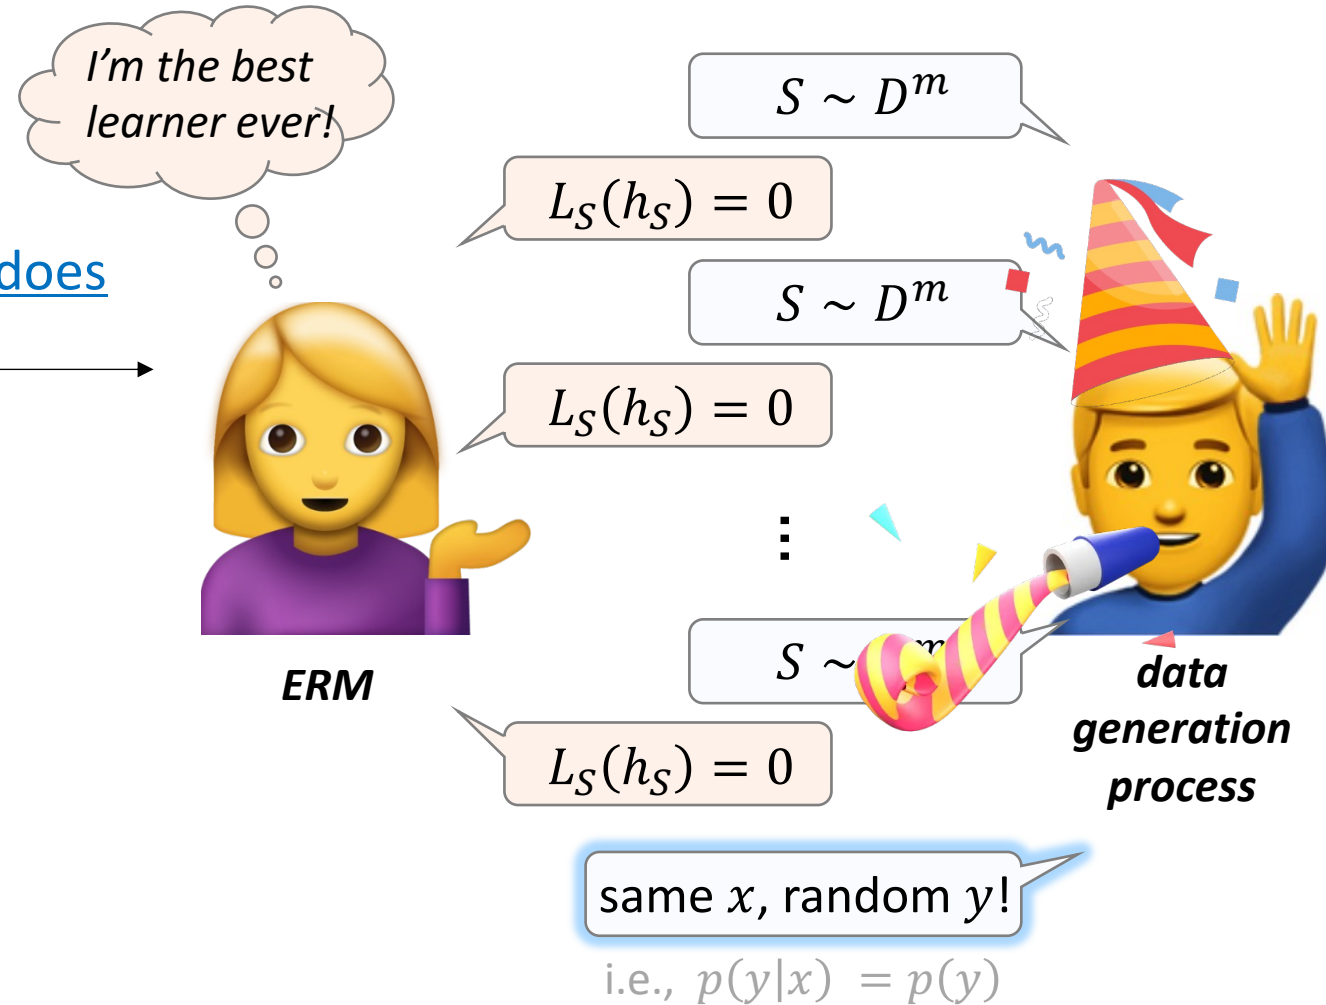
- The previous bound showed finite H are learnable (with dependence on $\log|H|$)
- Is this bound useful for...
 - decision trees? (think!)
 - linear halfspaces? (think!)
 - RBF kernels? (think!)
 - 1D thresholds? (think!)
- **Q:** does $|H| = \infty$ mean we can't learn?
- **A:** Not necessarily! (we proved finite $H \Rightarrow$ learnable, not the negation)
- **Recall:** for 1D thresholds (infinite class!), we showed $\epsilon \approx O\left(\frac{1}{m}\right)$ (under realizability)
- **Conclusion:** $|H|$ is probably not the “correct” measure
- **Note:** there is no single “correct” measure, only *useful* measures; we will see one next

VC dimension

- **Idea:**
consider not what each h is, but what it does

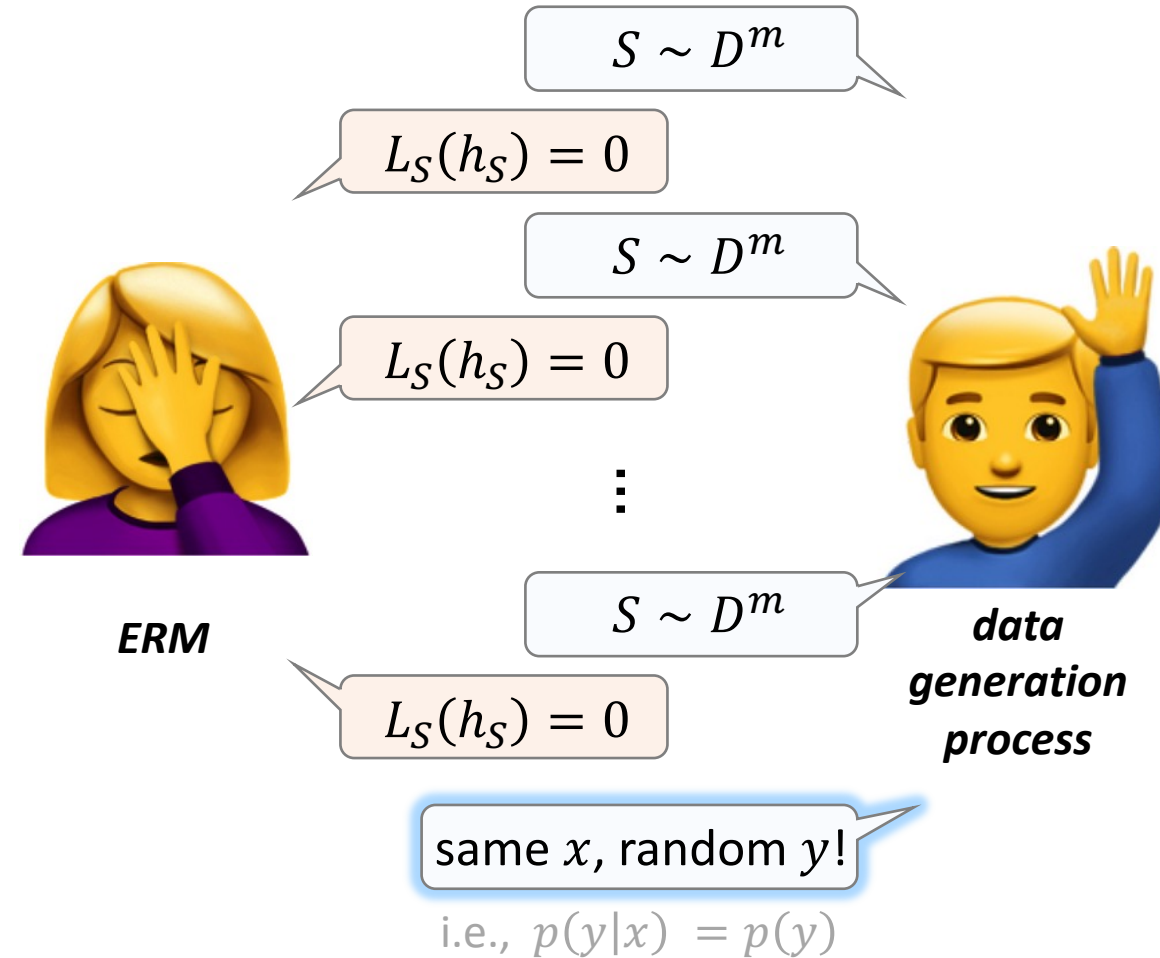
VC dimension

- Idea:
consider not what each h is, but what it does
- Intuition – when learning fails



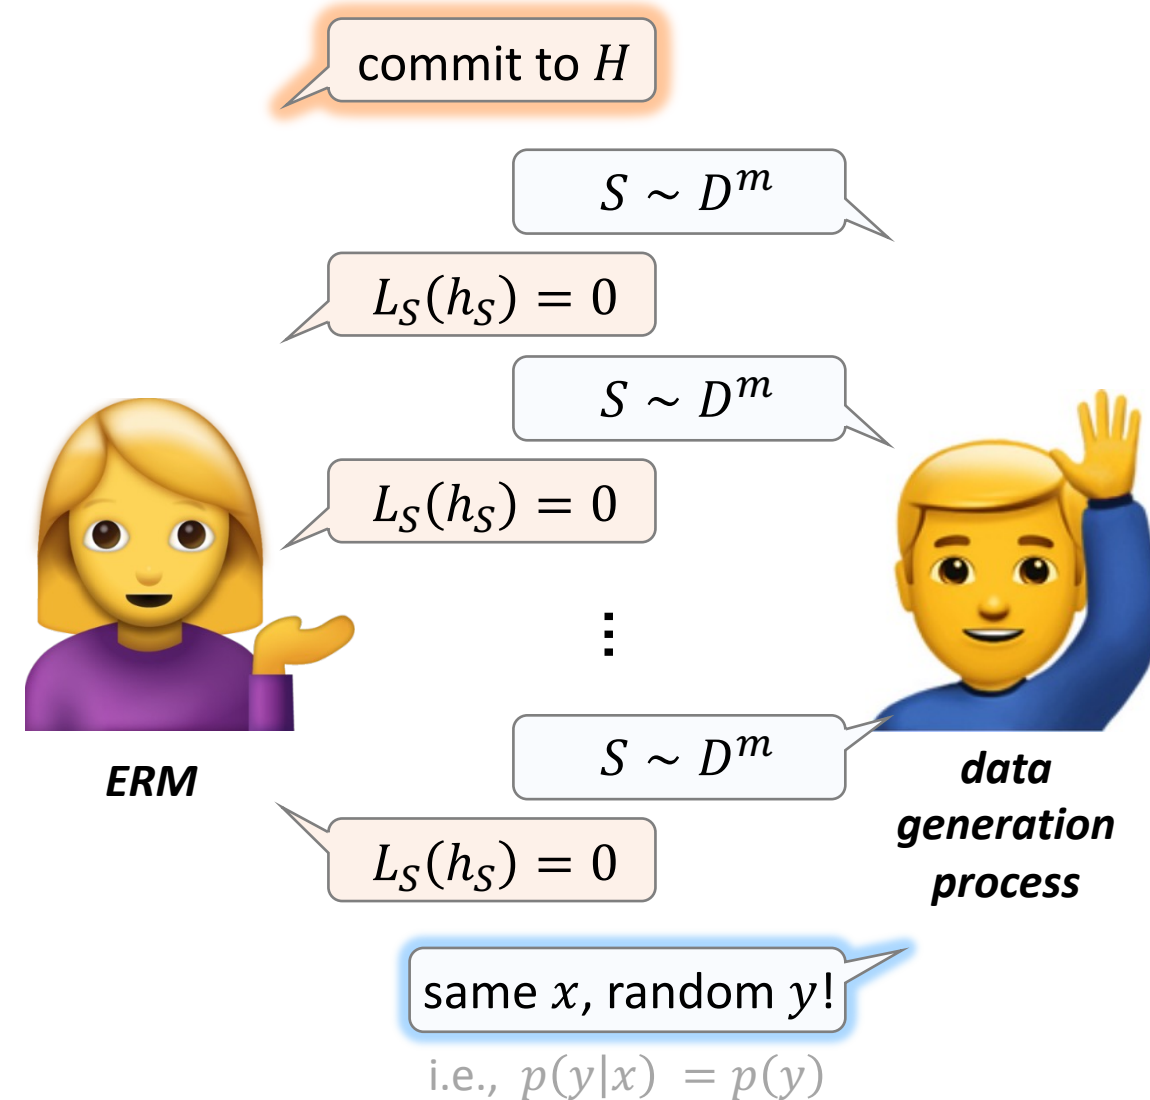
VC dimension

- **Idea:**
consider not what each h is, but what it does
- **Intuition – when learning **fails**** →
- **Take away:**
“explaining everything \equiv explaining nothing”
- **VC theory** quantifies this idea (Vapnik–Chervonenkis)



VC dimension

- **Idea:**
consider not what each h is, but what it does
- **Intuition – when learning fails** →
- **Take away:**
“explaining everything \equiv explaining nothing”
- **VC theory** quantifies this idea (Vapnik–Chervonenkis)
- **Remember:** learnability is property of H !
- The **VC dimension** of H is the **largest set** on which $L_S = 0$ is possible for any labeling
- **Main result:** learning breaks once H can perfectly fit arbitrary label assignments (= noise! remember overfitting?)



VC dimension

- The notion of “explaining everything” is defined using the idea of *shattering* a set of examples:

- **Definition:** Let $C = \{x_i\} \in \mathcal{X}^m$, then H *shatters* C if:

$$\forall \{y_i\} \in \{\pm 1\}^m \quad \exists h \in H \text{ s.t. } h(x_i) = y_i \quad \forall i \in [m]$$

i.e., for any labeling of C , applying ERM to $S(C) = \{(x_i, y_i)\}_{i=1}^m$ gives $L_{S(C)}(h_{S(C)}) = 0$.


- **Definition:** The *VC-dimension* of H is the *size* of the largest set that H *shatters*, denoted $VC(H)$

i.e., $VC(H) = m$ if exists C of size m that H shatters, but H does not shatter all larger sets

Finding VC

- Given H , how do we prove $VC(H)$?
- **Rules of the game:**
 1. guess some m
 2. show exists C of size m that H shatters $\Rightarrow VC \leq m$
 3. show H does **not shatter** all sets C of size $m + 1 \Rightarrow VC \geq m$

- **Examples:**
 1. 1D thresholds [on board]
 2. 1D intervals [on board]
 3. Linear halfspaces? (tirgul!)
 4. RBF kernel? (think!)

- 
1. $m: \exists \{x\} \forall y \exists$ perfect h (shatter)
 2. $m + 1: \forall \{x\} \exists y \forall h$ errs (can't shatter)

VC dimension

- The VC dimension of H tells us how many samples are needed for learning
- This is called the *sample complexity* of H , denoted $m_H(\epsilon, \delta)$ (look familiar?)

- **Fundamental theorem of learning:** (partial; won't prove)

If $VC(H) \leq \infty$, then H is:

1. **PAC-learnable** with $\text{vs. } \log|H|$

what about $|H| = \infty$?

$$m_H(\epsilon, \delta) = \Theta\left(\frac{VC(H) \log 1/\epsilon + \log 1/\delta}{\epsilon}\right)$$

exact characterization

2. **Agnostic PAC-learnable** with

$$m_H(\epsilon, \delta) = \Theta\left(\frac{VC(H) + \log 1/\delta}{\epsilon^2}\right)$$

(almost) same ϵ, δ rates as in finite H

- **Think:** what is the VC of a finite H ?
- **Bonus:** If $VC(H) = \infty$, then H is essentially not learnable (in the PAC sense)
- But there are other notions of learning! (e.g., SVM-RBF has $VC = \infty$, but is still great for learning)

Uses and limitations

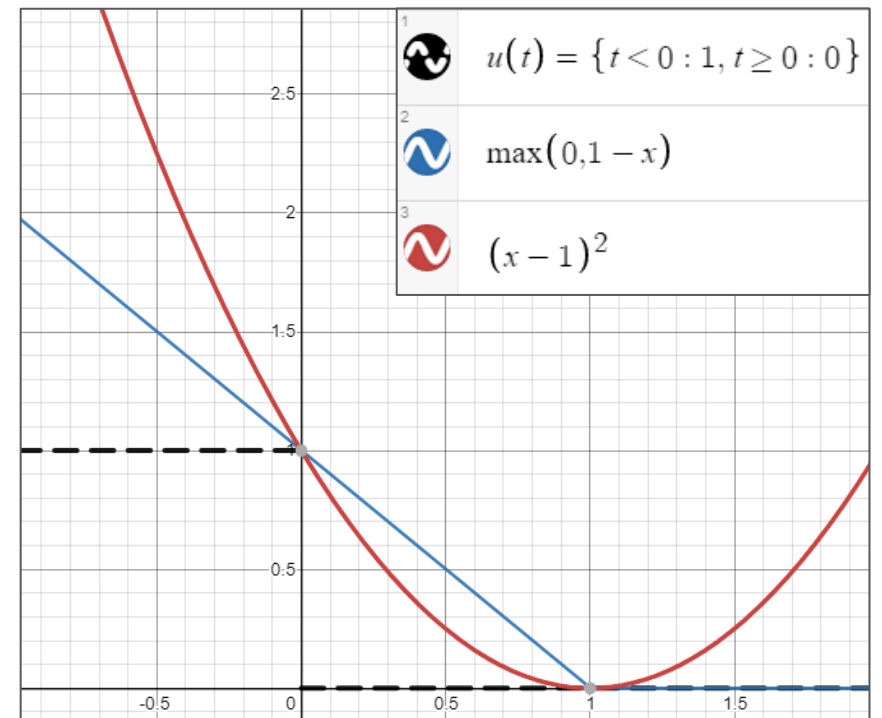
- Say you want to learn with SVM (and assume you know the VC of halfspaces*)
- **Theory is your friend:**
 - Theory asks: tell me your desired ϵ and δ (this is unavoidable!)
 - Theory says: you need (order of) m examples!
- **Great, but need to remember:**
 1. VC and PAC are worst-case (are you really doomed if you only have $< m$ samples?)
 2. Even agnostic PAC relies on distributional assumptions (the elephant in the room: i.i.d.)
 3. ERM is (computationally) hard! SVM minimizes *hinge loss*, not 0/1 loss (UC assumes exact ERM)
 4. Guarantees are probabilistic but (in most cases) **you only see one sample set** → **up next!**

Model Selection

A tale of bias and variance

- **Recall:** PAC asks: $P_{S \sim D^m}(L_D(h_S) - L_D(h^*) \geq \epsilon) \leq \delta$
- **Alternative:** expected error $\mathbb{E}_{S \sim D^m}[L_D(h_S)]$
- **Think:** why are we not considering $L_D(h^*)$ here?
- To simplify the math, we'll analyze **squared loss**:
$$L_D^{sqr}(h) = \mathbb{E}_{(x,y) \sim D}[(h(x) - y)^2]$$

(instead of 0/1 or hinge; will revisit when we talk about regression)
- **Let's analyze!** [on board]
- **Definitions:**
 - Expected label: $\bar{y}(x) = \mathbb{E}_{y \sim D_{Y|X=x}}[y] \in [0,1]$
(for squared error, this is the optimal classifier; won't prove)
 - Expected loss (given h): $\mathbb{E}_{(x,y) \sim D}[(h(x) - y)^2]$
 - Expected "classifier": $\bar{h} = \mathbb{E}_{S \sim D^m}[h_S]$



Error decomposition

$$\mathbb{E}_{S \sim D^m} [L_D^{sqr}(h_S)] = \mathbb{E}_{S \sim D^m} \mathbb{E}_{(x,y) \sim D} [(h_S(x) - y)^2]$$

...

$$= \mathbb{E}_{x,y} [(\bar{y}(x) - y)^2] + \mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2] + \mathbb{E}_{S,x} [(h_S(x) - \bar{h}(x))^2]$$

Full derivation here:

<https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote12.html>

Bias-Variance decomposition

- Got elegant decomposition – three interpretable sources of error:

$$\mathbb{E}_{S \sim D^m} [L_D^{sq}(h_S)] = \underbrace{\mathbb{E}_{x,y} [(\bar{y}(x) - y)^2]}_{\text{expected error}} + \underbrace{\mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2]}_{\text{noise}} + \underbrace{\mathbb{E}_{S,x} \left[\left(h_S(x) - \bar{h}(x) \right)^2 \right]}_{\text{bias}^2}$$

given that we chose hypothesis h_S

variance

- **Noise:**

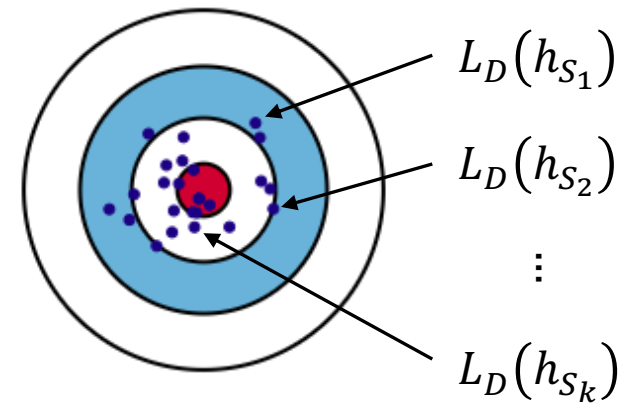
- For squared loss, \bar{y} is an optimal predictor (a.k.a. Bayes-optimal; won't show)
- Property of data distribution (i.e., the statistical relation between x and y)
- Does not depend on choice of model (hence called “irreducible” error)

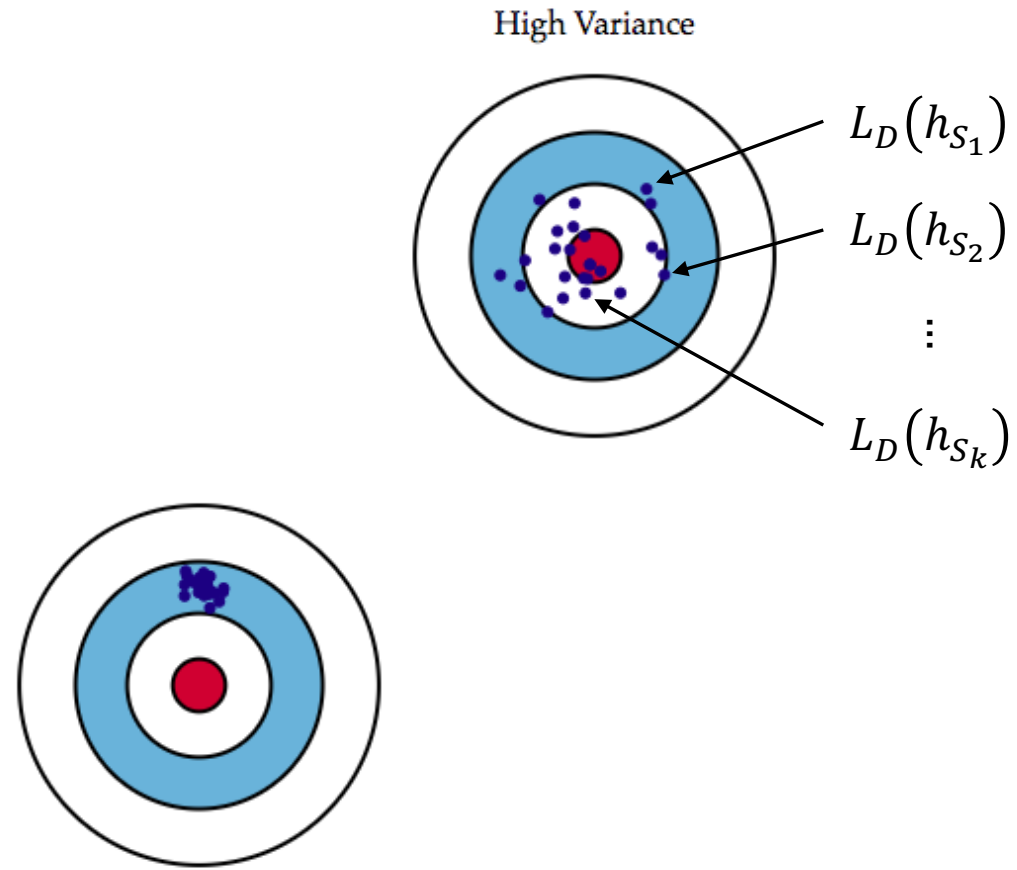
Bias-Variance decomposition

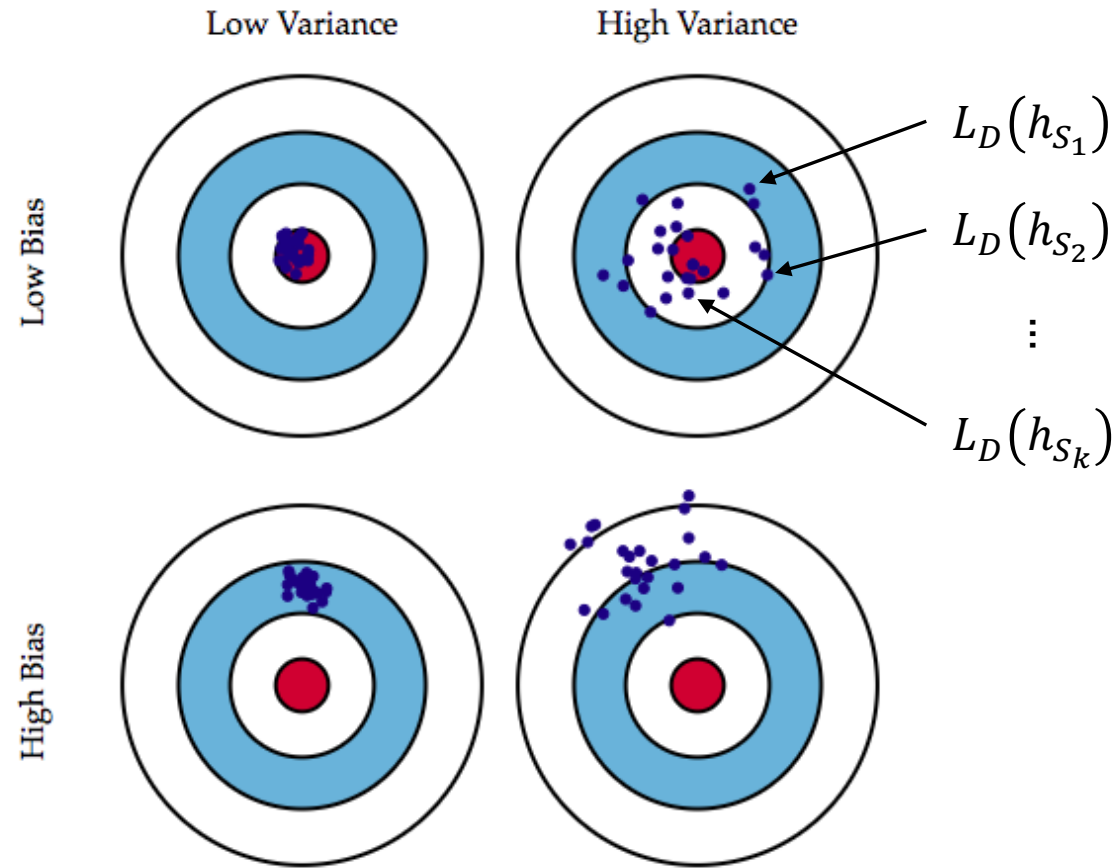
- Got elegant decomposition – three interpretable sources of error:

$$\mathbb{E}_{S \sim D^m} [L_D^{sq}(h_S)] = \underbrace{\mathbb{E}_{x,y} [(\bar{y}(x) - y)^2]}_{\text{expected error}} + \underbrace{\mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2]}_{\text{noise}} + \underbrace{\mathbb{E}_{S,x} [(h_S(x) - \bar{h}(x))^2]}_{\text{bias}^2} + \underbrace{\mathbb{E}_{S,x} [(h_S(x) - \bar{h}(x))^2]}_{\text{variance}}$$

- **Bias:**
 - Quantifies how well our chosen model class fits the data (on average)
 - Does not depend on the data sampled (but does depend on data size) nor on learned model
- **Variance:** (of algorithm; w.r.t. S)
 - Measures how learned models h_S vary
(how “sensitive” the learning algorithm is to changes in its input S)
 - Average model \bar{h} as reference point
(asks: relative to \bar{h} , how “specialized” is h_S to S ?)
 - Does not consider predictive error directly (no dependence on y)

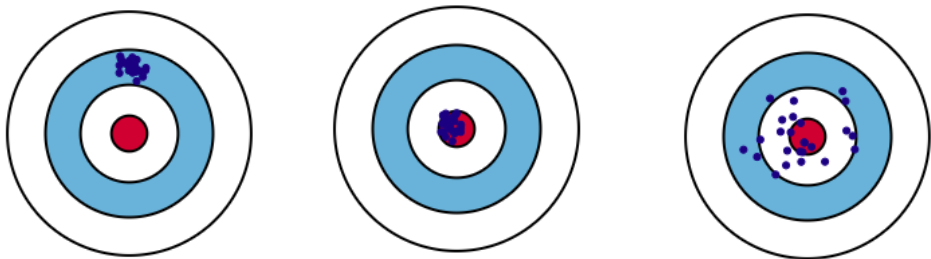
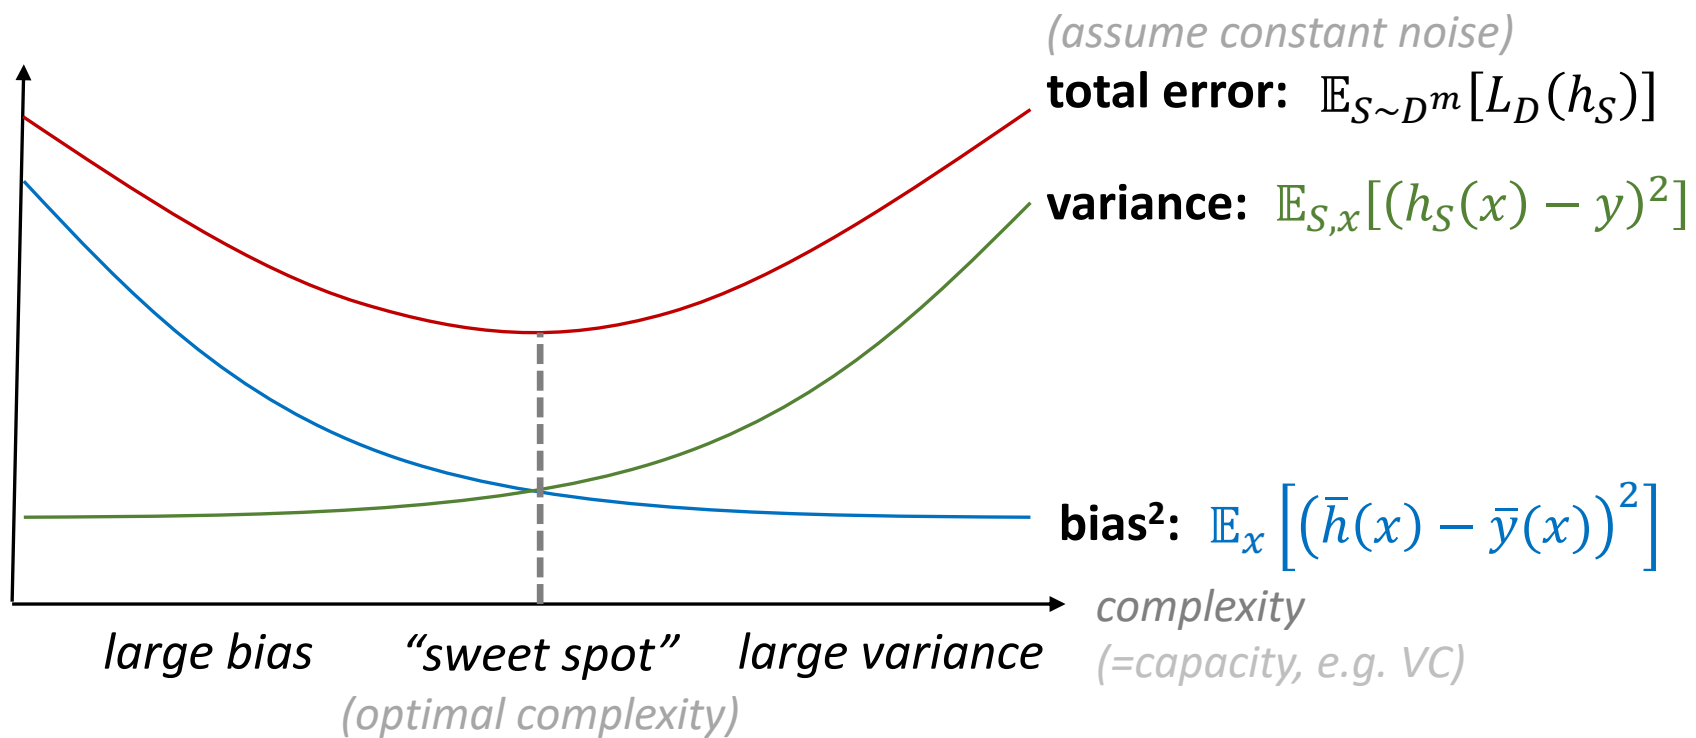


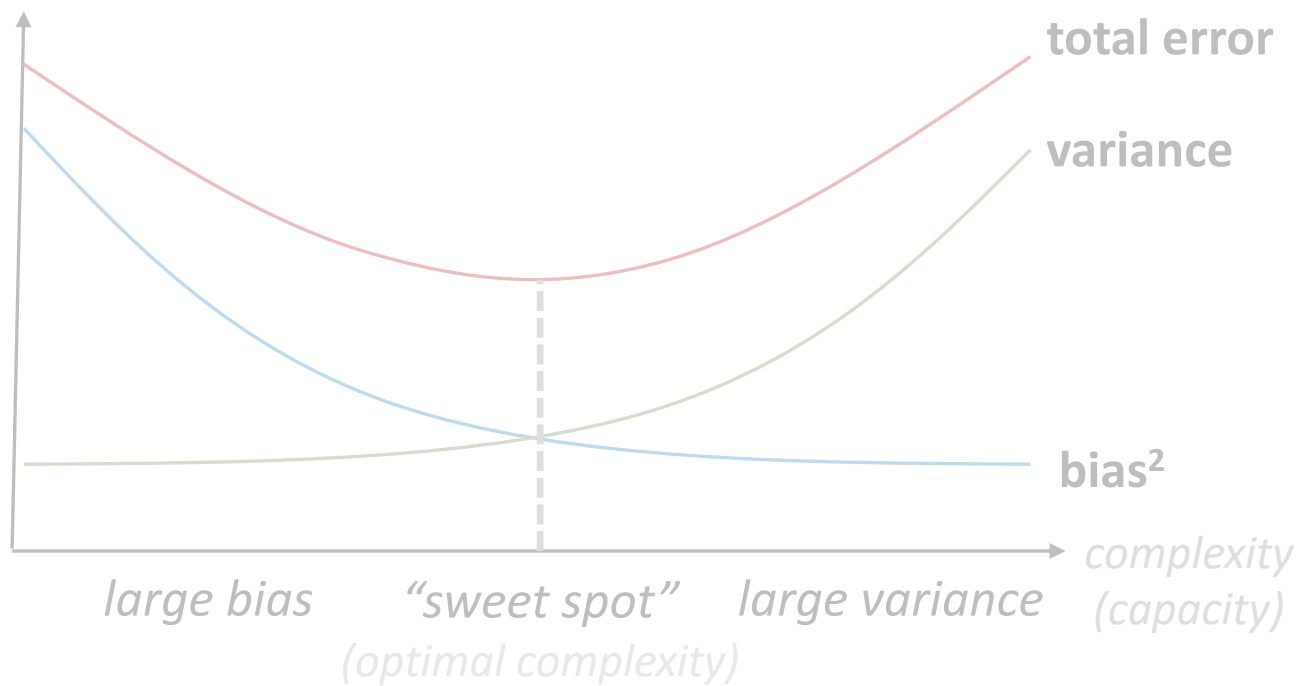




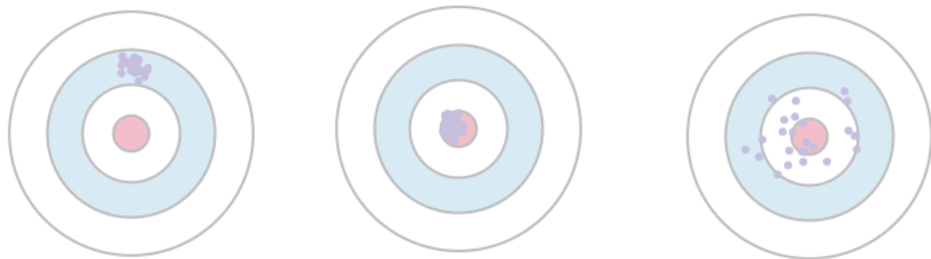
Q: Can we choose what regime we're in?

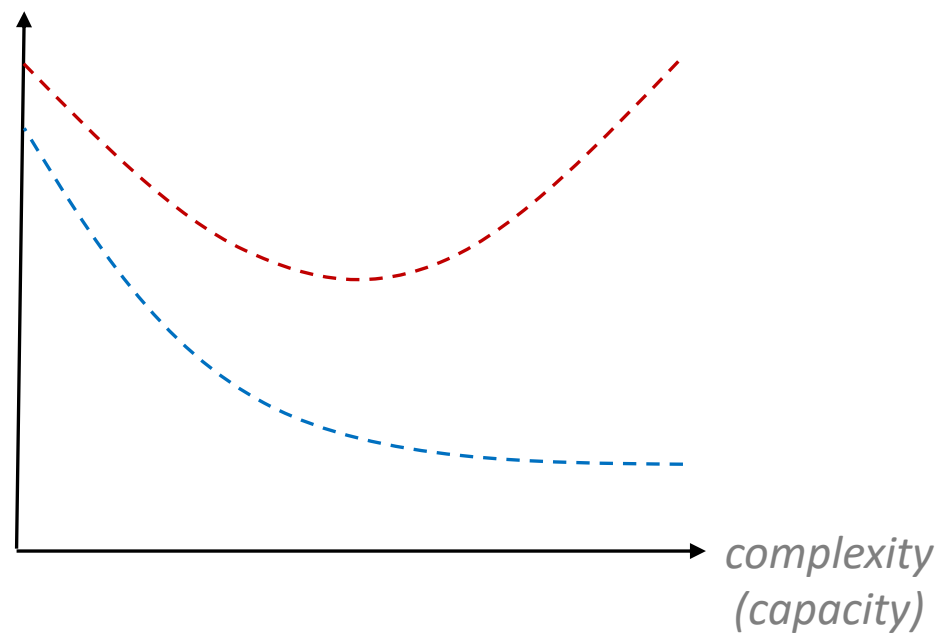
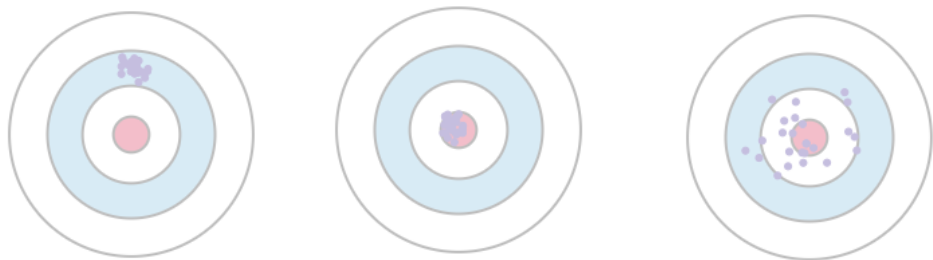
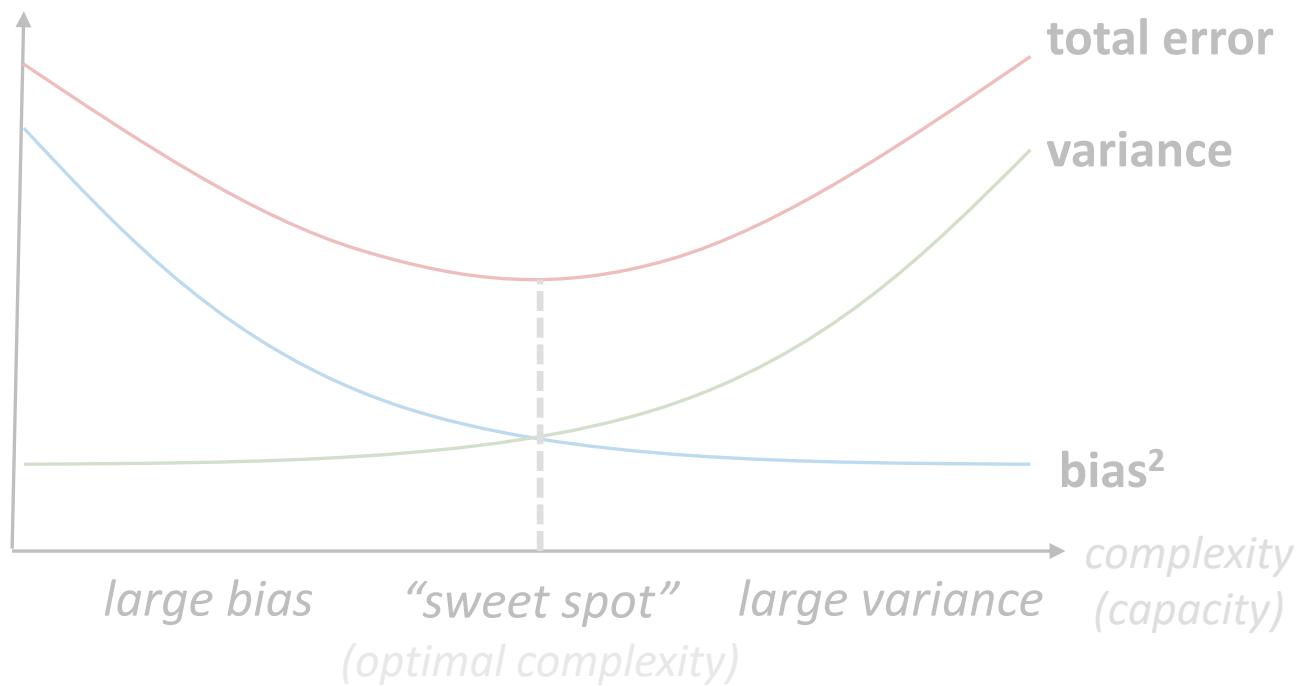
A: Yes, to some extent – by controlling model complexity.

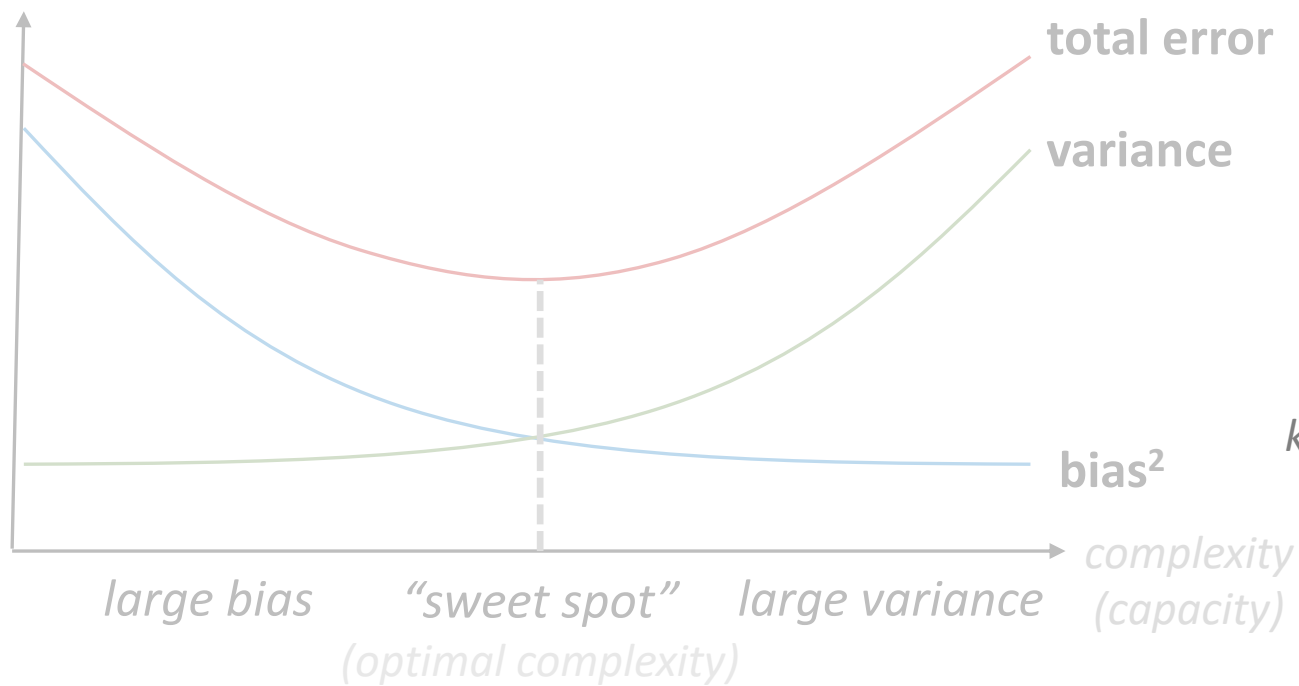




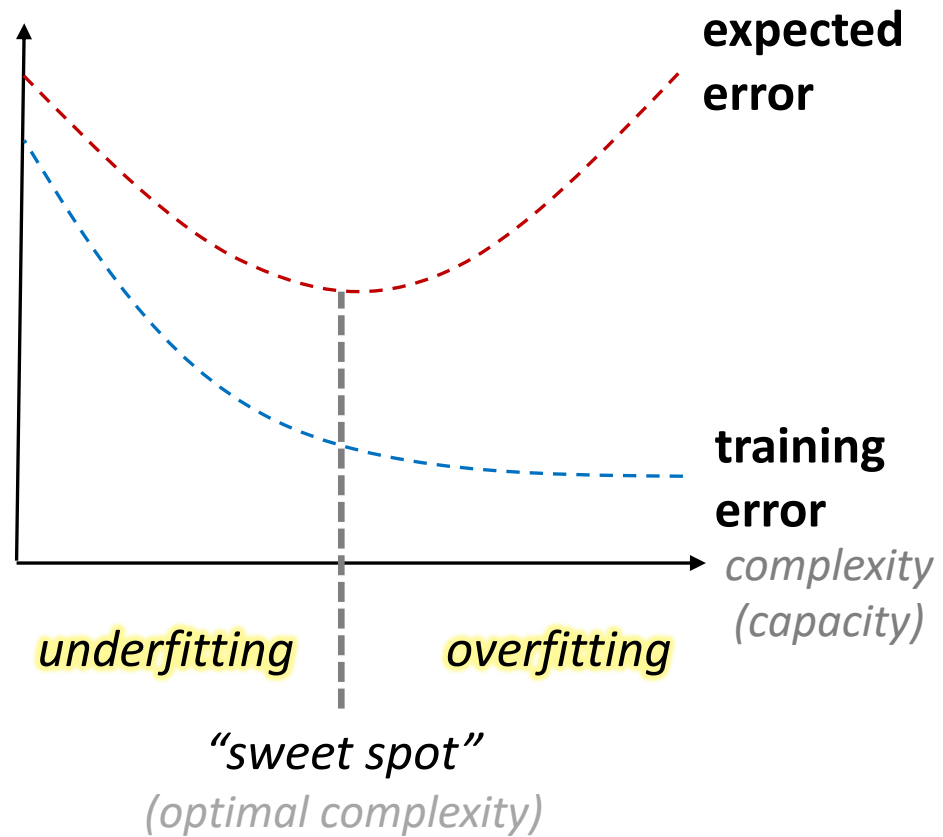
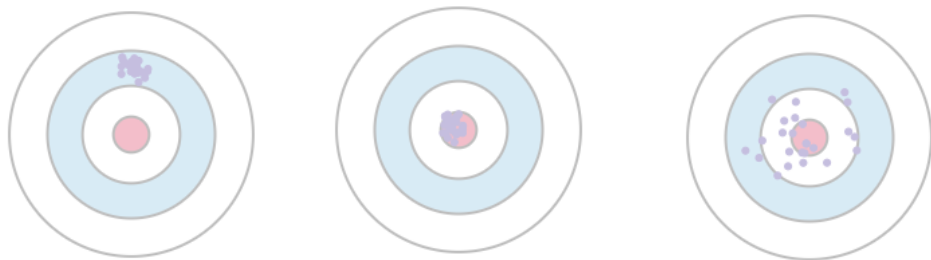
we've seen this sort of plot before...







now we know why!



Q: What is same? What is different?

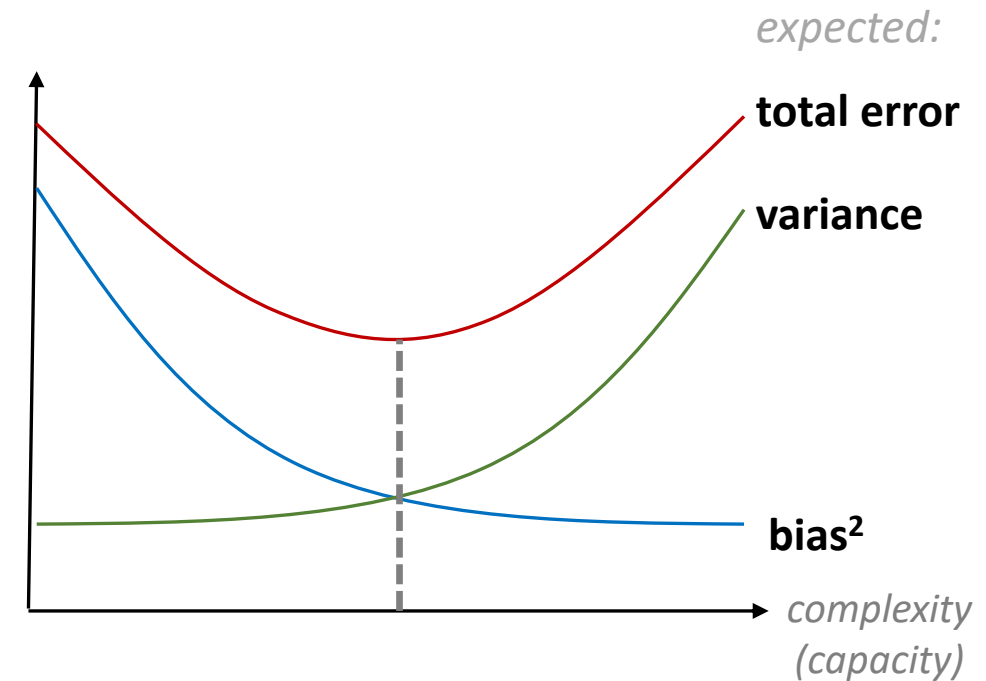
From theory to practice

- **Recall:** goal in learning is to reduce expected error

- **Decomposition says:**

$$\text{error} = \text{noise} + \text{bias}^2 + \text{variance}$$

- **Take-away I:**
Inherent tradeoff between bias and variance
- And we can control operating point (to some extent)
 - Large bias? Increase complexity!
 - Large variance? decrease complexity!
- Many learning problems have U-shaped errors
- We'd like to find the “sweet spot” (will see how soon)

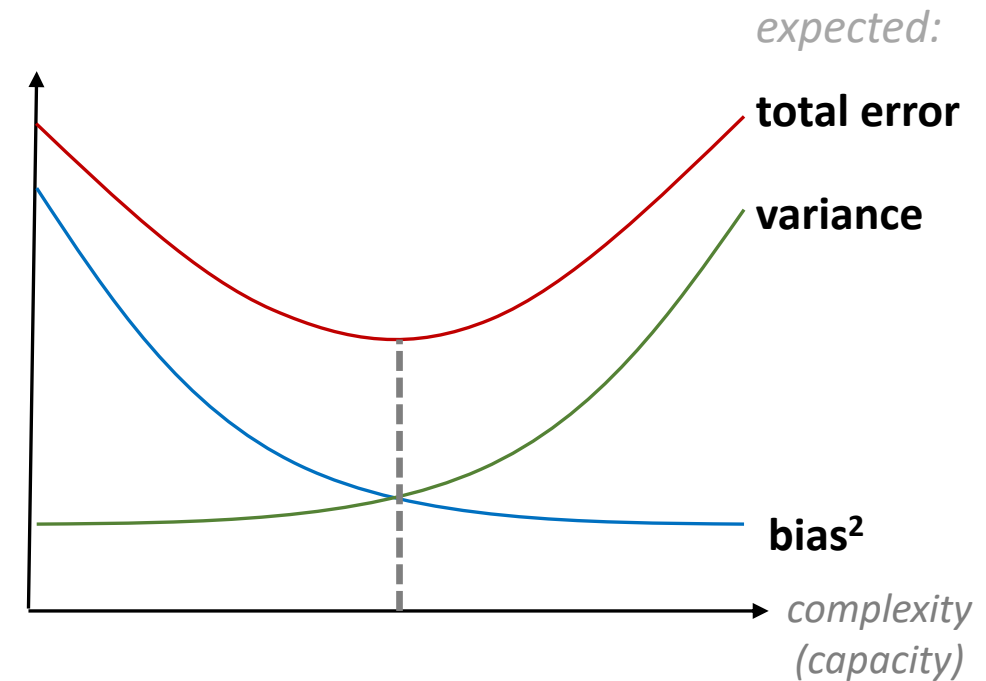


From theory to practice

- **Recall:** goal in learning is to reduce expected error
- **Decomposition says:**

$$\text{error} = \text{noise} + \text{bias}^2 + \text{variance}$$

- **Take-away II:**
Effective modeling = “choose your battles”
- Easier to individually target each source of error
- **One solution:** hard-code into learning objective
- (Hint: we’ve actually already seen this in action!)



Regularization

SVM revisited

- **Recall:**

- **Want:** low expected error $L_D(h)$
- **Have:** low empirical error $L_S(h)$
- **Care:** generalization $L_D(h_S)$ for $h_S = A(S)$

- **Soft SVM objective:**
$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

- **Q:** We know SVM is not ERM... so what is it?
- **Loss:** replaced 0/1 with hinge proxy for optimization reasons
- **Norm:** pushes solution \hat{w} away from 0/1-optimal w^*
- Justification for adding $\lambda \|w\|$:
 - 1. **Modeling:** max margin
 - 2. **Optimization:** strong convexity
 - 3. **Stats:** bias-variance! (up next)
- Approach called **Regularized Loss Minimization** (RLM)

SVM as RLM

- **Soft SVM objective:**

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

● total error ● variance ● bias²

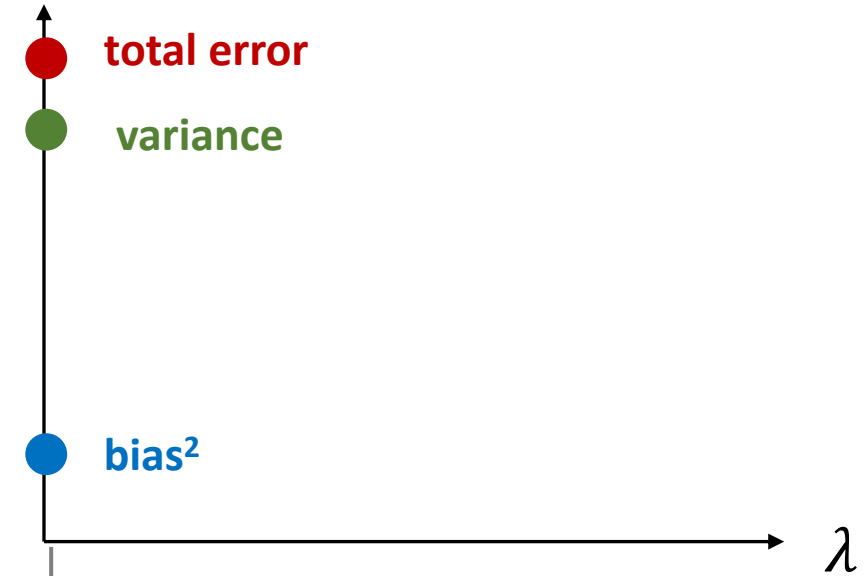


SVM as RLM

- **Soft SVM objective:**

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

● total error ● variance ● bias²



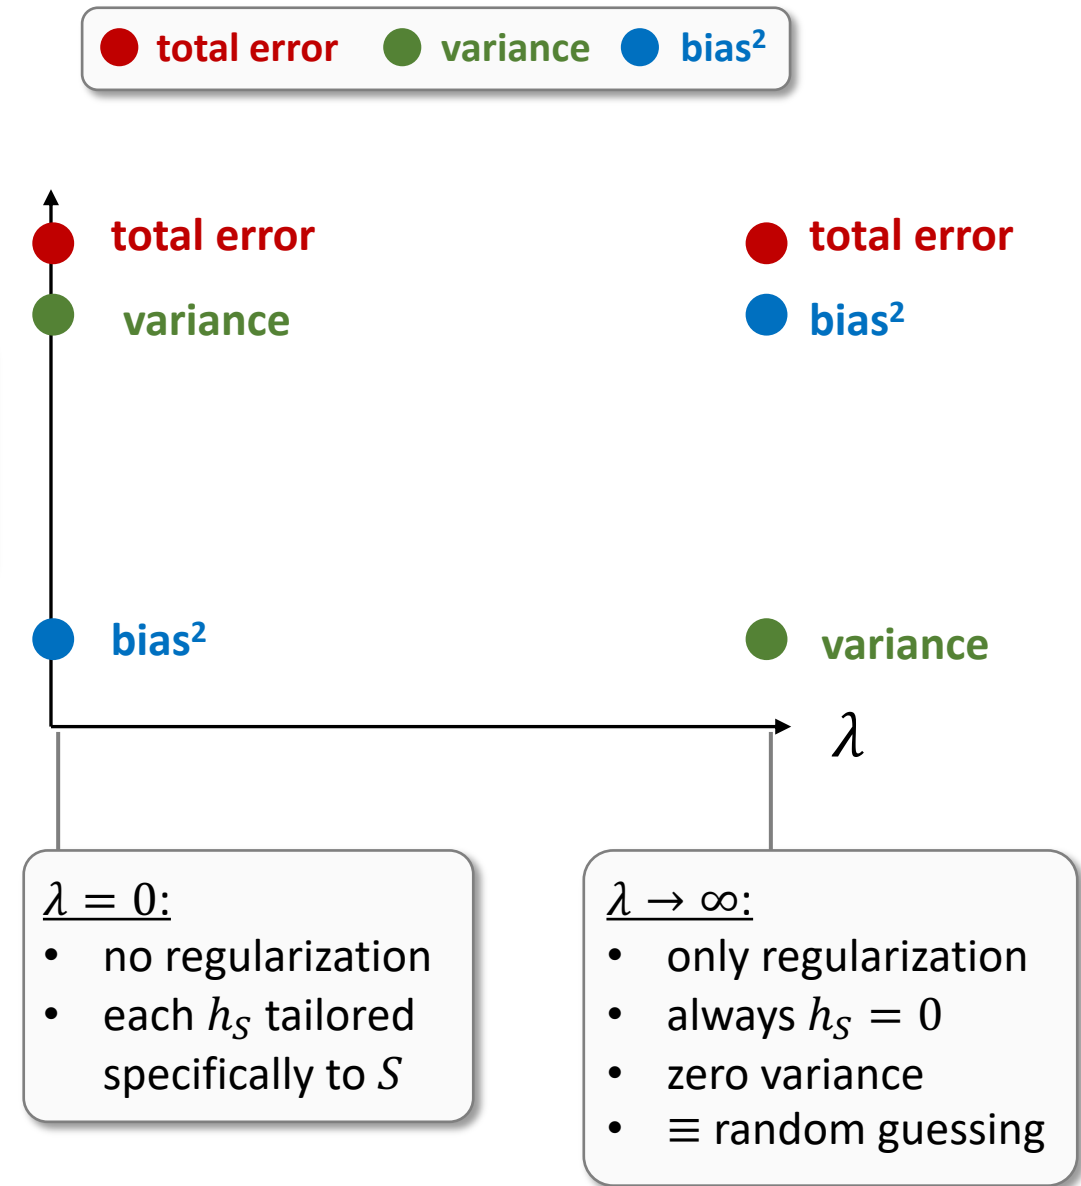
$\lambda = 0$:

- no regularization
- each h_S tailored specifically to S

SVM as RLM

- **Soft SVM objective:**

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$



SVM as RLM

- **Soft SVM objective:**

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

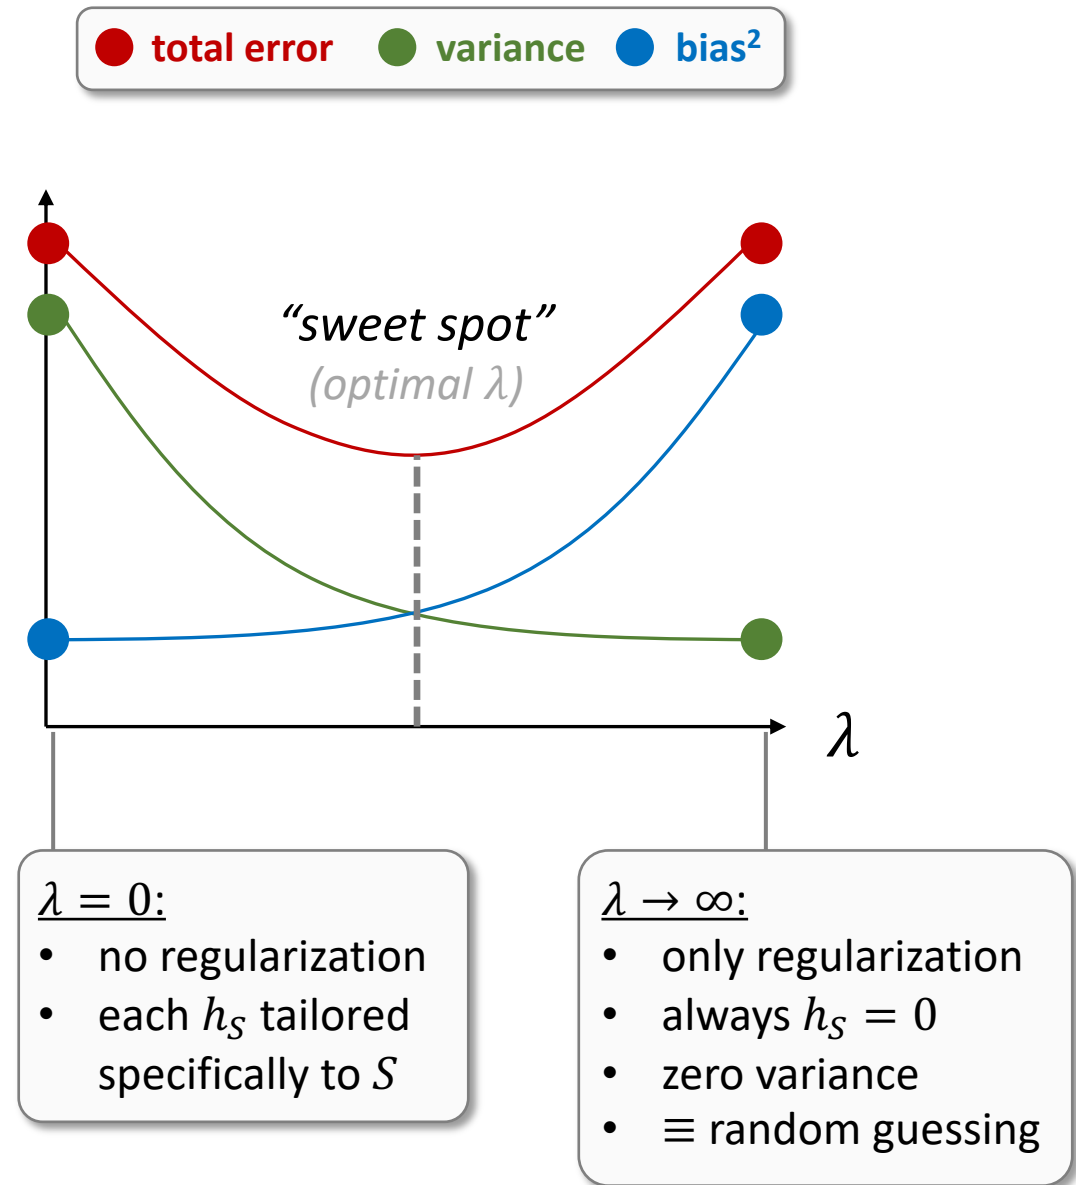
- **This looks like a bias-variance plot!**

(even though it's hinge and not squared loss)
(and not by coincidence)

- **Interpretation – decoupled learning objective:**

- **loss** – directly reduces **bias**
(and ensures tractable optimization!)
- **regularization** – indirectly reduces **variance**
(while knowingly incurring of some bias)

- In effect, regularization controls (“regulates”) complexity. **Let’s see how.**



Generalization of SVM

(use $w \in H$ and $h \in H$ interchangeably)

- **Define:** bounded-norm linear models $H_B = \{w \in \mathbb{R}^d : \|w\|_2 \leq B\}$

- **Theorem:** (won't prove) for any m , Soft SVM with $\lambda = \sqrt{2/(B^2 m)}$ satisfies:

$$\mathbb{E}_{S \sim D^m} \left[L_D^{\text{hinge}}(w_S) \right] \leq \min_{w \in H_B} L_D^{\text{hinge}}(w) + B \sqrt{\frac{8}{m}}$$

- **Conclusion:** low-norm models $w \in H_B$ generalize better!

- Can use above to get PAC-style sample complexity: $m_H(\epsilon, \delta) \geq O\left(\frac{B^2}{\epsilon^2 \delta}\right)$
- Compare to VC bound: $m_H(\epsilon, \delta) \geq \tilde{O}\left(\frac{VC + \log 1/\delta}{\epsilon^2}\right)$ (for 0/1 loss)
- **Conclusion:** restricting norm as means to control complexity

A broader modeling perspective

- **Soft SVM:**
$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

- **General template:**
$$\operatorname{argmin}_{w \in \mathbb{R}^d} \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(y_i, w^\top x_i)}_{\text{loss}} + \underbrace{\lambda R(w)}_{\text{regularization}}$$

- **Regularized Loss Minimization (RLM):**
 - **Interpretation #1:** knowingly add bias to reduce variance
 - **Interpretation #2:** "of all models with equally low loss, choose the one that has [...]"
 - **Interpretation #3:** means to structurally (and softly) plug in prior knowledge

A broader modeling perspective

$$\underset{w \in \mathbb{R}^d}{\operatorname{argmin}} \frac{1}{m} \sum_{i=1}^m \ell(y_i, w^\top x_i) + \lambda R(w)$$

loss regularization

- **Q:** What type of prior knowledge does “small L2-norm” express?
(don’t confuse small scale with small norm! small norm implies “small” in very certain way)
- **A:** for Soft SVM – large margin (but also useful in general!)
- **There are many, many ways to regularize.** For example, **other norms**.
- **[DESMOS]**
- **Q:** What type of prior knowledge does the L1-norm express?
- **A: Sparsity** – encourages learned w_s to have only few non-zero entries (see tirgul)

Model selection

Model selection

- Assume you decide to learn using Soft SVM:

$$\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

- **Q:** How should you choose λ ?
- **Attempt #1:** optimize loss over train set:

$$\operatorname{argmin}_{w \in \mathbb{R}^d, \lambda \in \mathbb{R}_+} \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \cdot w^\top x_i\} + \lambda \|w\|_2^2$$

- **Wrong!** Will always give $\lambda = 0$ (think – why?)
- **Attempt #2:** use data to (empirically) estimate $L_D(h_S; \lambda)$

Error estimation

- Assume you have access to an additional sample set $V \sim D^{m_v}$ of size m_v
- Define $L_V(h)$ as the empirical error over V
- **Idea:** use V to estimate expected error of classifier

- **Theorem:** Let $h \in H$. Then for any $\delta \in [0,1]$, with probability $\geq 1 - \delta$, it holds that

$$|L_V(h) - L_D(h)| \leq \sqrt{\frac{\log 2/\delta}{2m_v}}$$

- **Proof:** Apply Hoeffding bound
- **Interpretation:** $L_V(h)$ is very good estimator of $L_D(h)$
- **Notice:** bound is independent of the choice of $H \ni h$
(**read:** estimation works equally well for any regardless the complexity of h)

Error estimation

- **Q:** Can L_V help us choose λ when learning? (rather than for a fixed h and given λ)

- **“Tuning”:**

1. Set range of λ , e.g., $\lambda \in \Lambda = \{2^{-8}, 2^{-7}, \dots, 2^8\}$
2. For each λ , use S to learn best h (using that λ)
3. Choose best λ using L_V (as estimate of L_D)

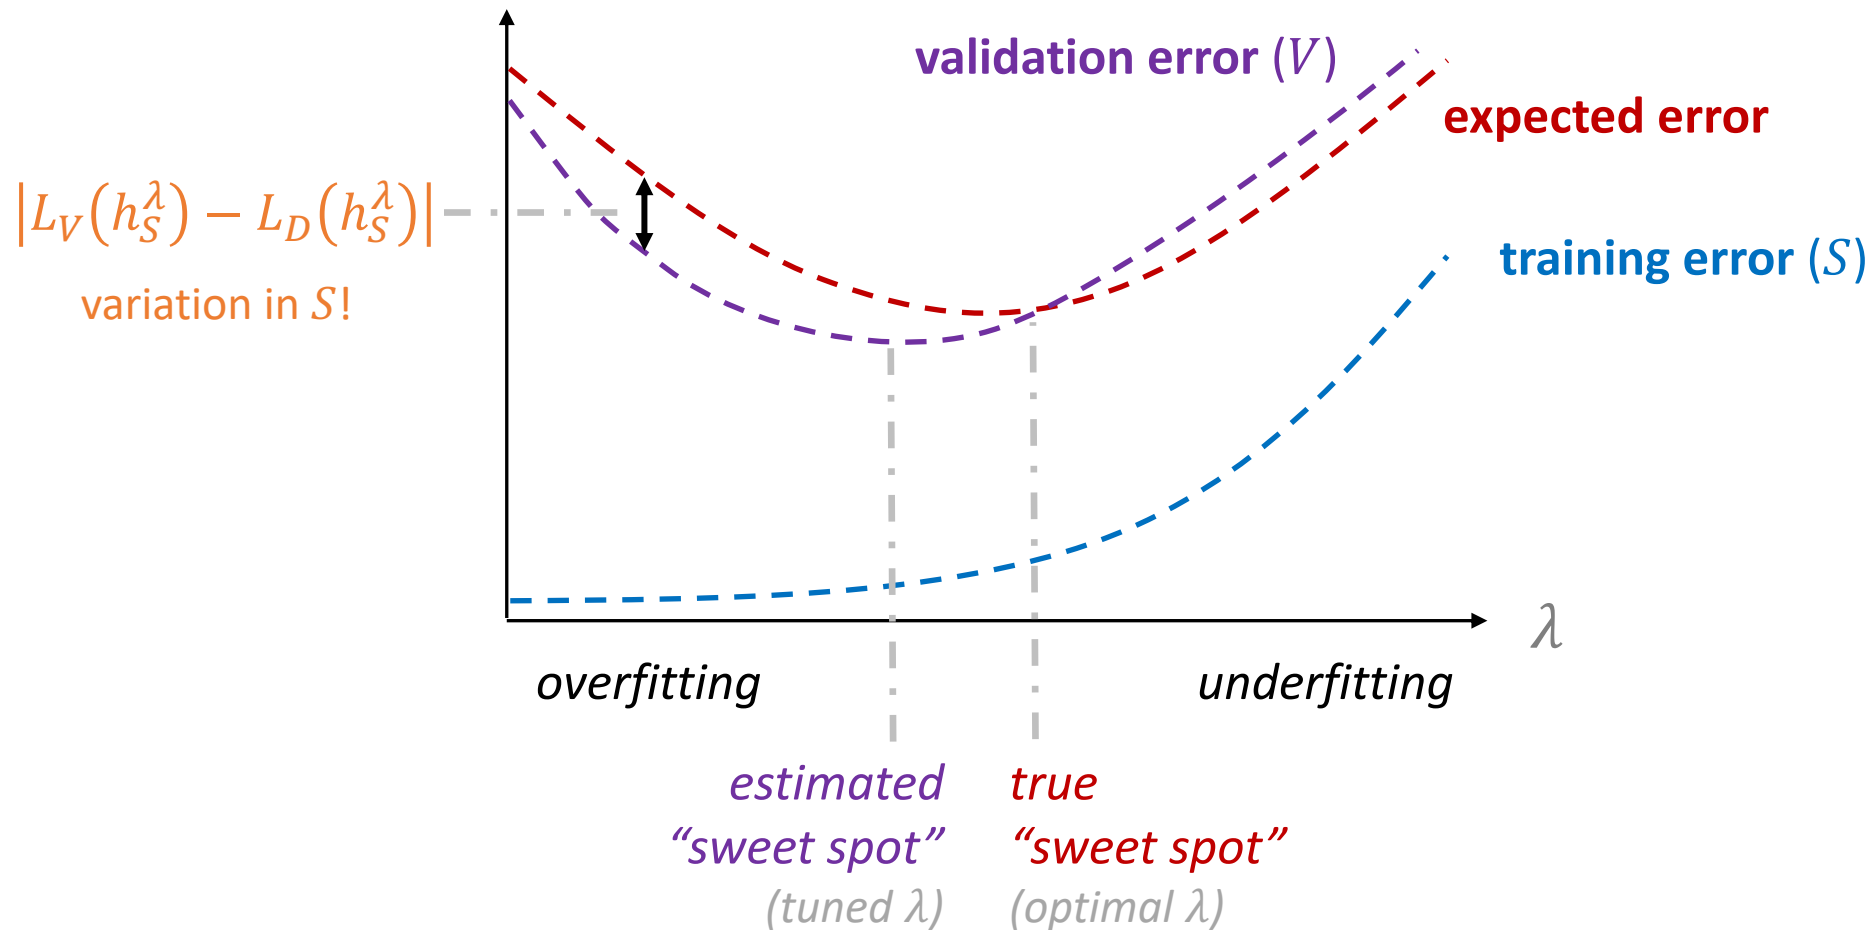
- V is called the *validation set* or the *held-out set*

- **Theorem:** Denote h_S^λ the model learned using λ . Let $H_\Lambda = \{h_S^\lambda : \lambda \in \Lambda\} \subset H$. Then for any $\delta \in [0,1]$, with probability $\geq 1 - \delta$, it holds that

$$\forall h \in H_\Lambda, \quad |L_V(h) - L_D(h)| \leq \sqrt{\frac{\log 2|\Lambda| + \log 1/\delta}{2m_v}}$$

- **A:** Think of H_Λ as a finite class, which we know is learnable – and apply PAC bound
- Works as long as V is sampled independently (of S , H , etc.)

Overfitting, revisited



Cross validation (CV)

- In reality, we usually don't have an “additional” sample set V
- In practice, must make use of S for both training (h) and tuning (λ)
- **Q:** How can we ensure training and tuning are independent?
- **Common solution:** k-fold cross-validation

Cross validation (CV)

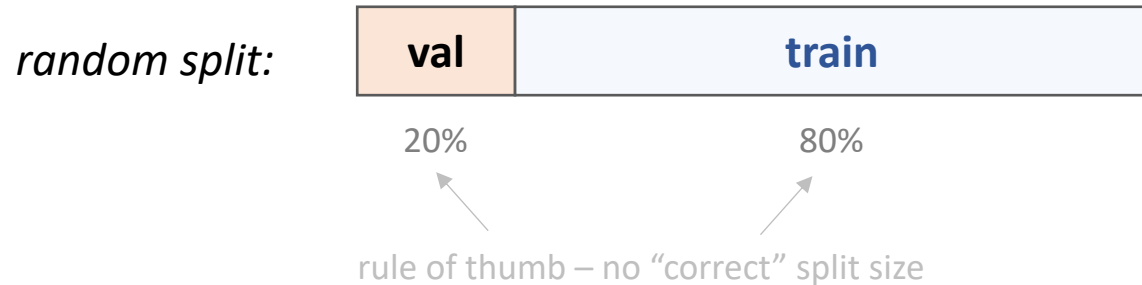
- In reality, we usually don't have an “additional” sample set V
- In practice, must make use of S for both training (h) and tuning (λ)
- **Q:** How can we ensure training and tuning are independent?
- **Common solution:** k-fold cross-validation

sampled data:

train

Cross validation (CV)

- In reality, we usually don't have an “additional” sample set V
- In practice, must make use of S for both training (h) and tuning (λ)
- **Q:** How can we ensure training and tuning are independent?
- **Common solution:** k-fold cross-validation



Cross validation (CV)

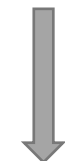
- In reality, we usually don't have an “additional” sample set V
- In practice, must make use of S for both training (h) and tuning (λ)
- **Q:** How can we ensure training and tuning are independent?
- **Common solution: k-fold cross-validation**

<i>fold 1</i>	val	train	train	train	train
<i>fold 2</i>	train	val	train	train	train
<i>fold 3</i>	train	train	val	train	train
<i>fold 4</i>	train	train	train	val	train
<i>fold 5</i>	train	train	train	train	val

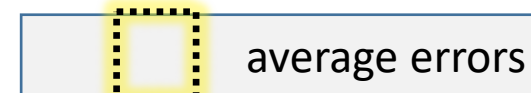
$k = 5$



$\lambda_1 \quad \lambda_2 \quad \lambda_3 \quad \dots \quad \lambda_{|\Lambda|}$



average



train on
all data

$\arg\max_{\lambda} \min$

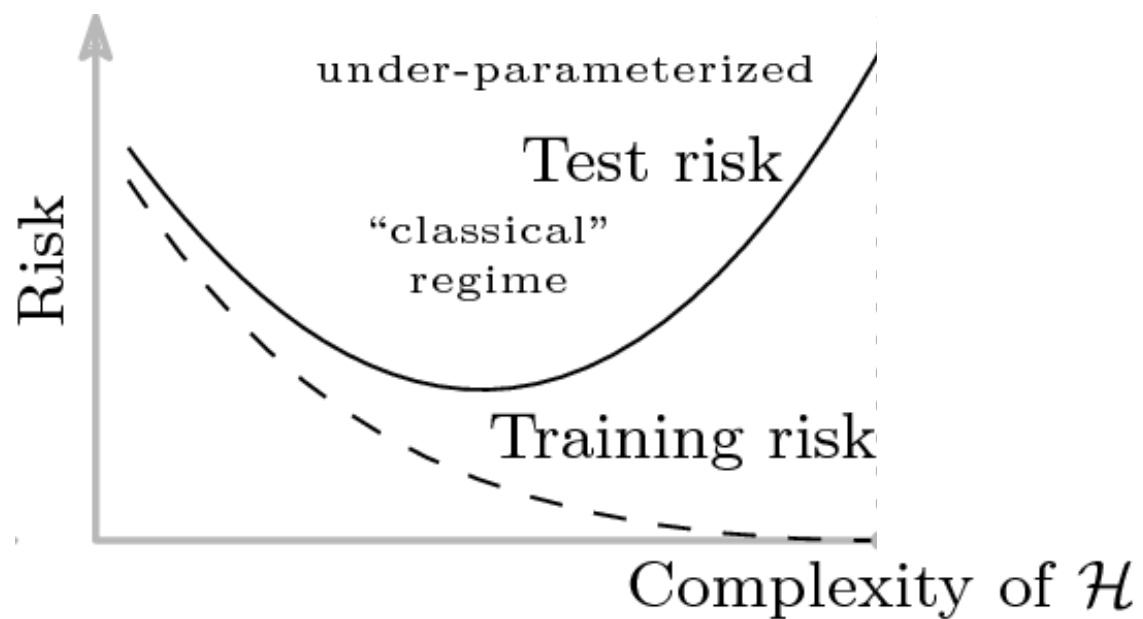
Discussion

- Three sources of error: **bias**, **variance**, and **noise**
- Smart modeling = carefully control each source
- **Different means to control:**
 - Feature collection/selection
 - Model class complexity (e.g., VC)
 - Hard-code into objective (e.g., RLM)
- **Thought experiment:** what is the difference between varying d and varying $\|w\|$?
- These inevitably introduce another layer of modeling, requiring either:
 - Prior knowledge (“decision trees should work well for this problem”)
 - Tuning by (cross-)validation (use data to determine optimal complexity)

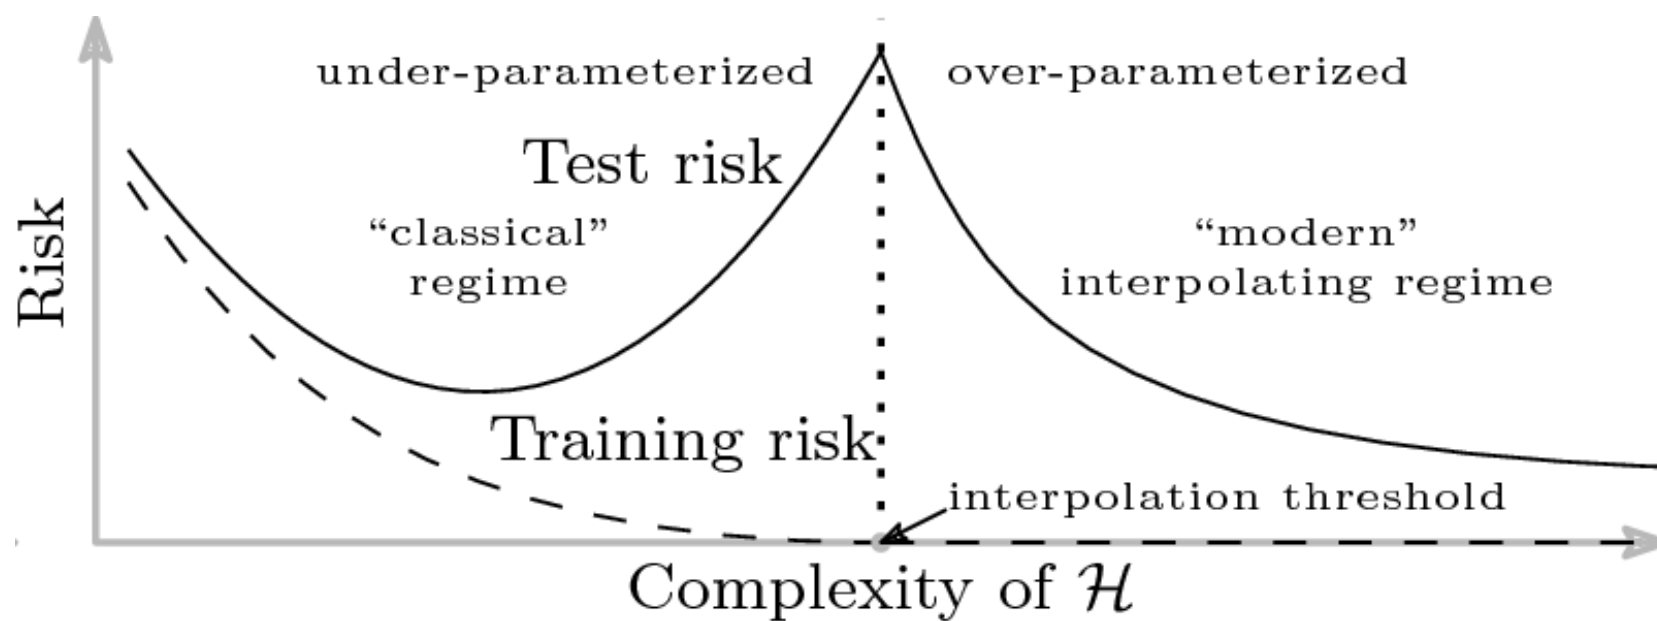
Discussion

- Many ways to tune hyper-params (e.g., λ): binary search, grid, random
 - **Remember:** tuning is **costly** (in samples, runtime, etc.)
 - **Important:** tuning is part of “learning”! (just not by ERM – still uses sample set)
 - **Bonus:** neural nets have tons of hyper-params, making tuning critical
-
- **In practice:** train-val-test split
 - Beware leakage! (this is the source of most “innocent” errors)
 - Beware “global” overfitting (e.g., repeated usage of same public dataset)

The limits of theory



The limits of theory



Next week(s)

- **Part II:** *the different aspects of learning*
 - ~~1. Statistics: generalization and PAC theory~~
 - ~~2. Modeling: model selection and evaluation~~
 3. Optimization: convexity, gradient descent
 4. Practical aspects and potential pitfalls

