

Introduction to Machine Learning (IML)

# LECTURE #9: REGRESSION

---

236756 – 2024 SPRING – TECHNION

LECTURER: NIR ROSENFELD



# Today

- **Part III:** *more supervised learning*
  1. Regression (today)
  2. Bagging and boosting
  3. Deep learning

# Learning setup

## Classification:

- **Features:**  $x \in \mathcal{X} = \mathbb{R}^d$
- **Labels:**  $y \in \mathcal{Y} = \{\pm 1\}$
- **Sample set:**  $S = \{(x_i, y_i)\}_{i=1}^m \stackrel{iid}{\sim} D^m$
- **Model class:**  $H = \{h \mid h: \mathcal{X} \rightarrow \mathcal{Y}\}$

- **Expected error:**

$$L_D(h) = \mathbb{P}_D[y \neq h(x)]$$

- **Empirical error:**

$$L_S(h) = \mathbb{P}_S[y \neq h(x)]$$

# Learning setup

## Classification:

- **Features:**  $x \in \mathcal{X} = \mathbb{R}^d$
- **Labels:**  $y \in \mathcal{Y} = \{\pm 1\}$
- **Sample set:**  $S = \{(x_i, y_i)\}_{i=1}^m \stackrel{iid}{\sim} D^m$
- **Model class:**  $H = \{h \mid h: \mathcal{X} \rightarrow \mathcal{Y}\}$
- **Loss function:**  $\ell^{0/1}(y, \hat{y}) = \mathbb{1}\{y \neq \hat{y}\}$
- **Expected error:**
$$L_D(h) = \mathbb{E}_D[\ell^{0/1}(y, h(x))]$$
- **Empirical error:**
$$L_S(h) = \frac{1}{m} \sum_{i=1}^m \ell^{0/1}(y_i, h(x_i))$$

# Learning setup

Classification: **Regression:**

- **Features:**  $x \in \mathcal{X} = \mathbb{R}^d$
- **Labels:**  $y \in \mathcal{Y} = \mathbb{R}$
- **Sample set:**  $S = \{(x_i, y_i)\}_{i=1}^m \stackrel{iid}{\sim} D^m$
- **Model class:**  $H = \{h \mid h: \mathcal{X} \rightarrow \mathcal{Y}\}$

- **Loss function:**  $\ell^{0/1}(y, \hat{y}) = \mathbb{1}\{y \neq \hat{y}\}$

- **Expected error:**

$$L_D(h) = \mathbb{E}_D[\ell^{0/1}(y, h(x))]$$

- **Empirical error:**

$$L_S(h) = \frac{1}{m} \sum_{i=1}^m \ell^{0/1}(y_i, h(x_i))$$

# Learning setup

~~Classification:~~ **Regression:**

- Features:  $x \in \mathcal{X} = \mathbb{R}^d$
- **Labels:**  $y \in \mathcal{Y} = \mathbb{R}$
- Sample set:  $S = \{(x_i, y_i)\}_{i=1}^m \stackrel{iid}{\sim} D^m$
- Model class:  $H = \{h \mid h: \mathcal{X} \rightarrow \mathcal{Y}\}$

- **Loss function:**  $\ell^{0/1}(y, \hat{y}) = \mathbb{1}\{y \neq \hat{y}\}$

- Expected error:

$$L_D(h) = \mathbb{E}_D[\ell^{0/1}(y, h(x))]$$

- Empirical error:


$$L_S(h) = \frac{1}{m} \sum_{i=1}^m \ell^{0/1}(y_i, h(x_i))$$

- Possible, but somewhat silly:
  - penalizes harshly even tiny errors
  - same penalty regardless of error size
- Solution is fairly straightforward, but we'll take the long road there!

Regression via  
statistical modeling

# History

- As old as time (Legendre 1805, Gauss 1809)
- **Why is it called “regression”?**
  - Galton (1889) and later Pearson (his student) observed that:
    - children of *tall* parents are *shorter* than their parents (on average)
    - children of *short* parents are *taller* than their parents (on average)
  - They concluded that observations have a statistical tendency to “regress” (i.e., “return”) to the mean (following extreme observations)
- Regression is deeply rooted in statistics; basis for many statistical tasks (beyond prediction)
- Way before ML came to be (and was popularized through classification)
- Today we’ll think about regression like *statisticians* would





# Statistical modeling

- **Recall:** in prediction, we care about  $p(y|x)$
- In discriminative learning, we've assumed  $y \stackrel{iid}{\sim} D_{Y|X=x}$
- Alternatively, **statistical modeling** assumes:
  - This is a direct assumption on the data generating process (more concretely, on  $p(y|x)$ , with arbitrary  $p(x)$ ; still assume  $x$  sampled iid)
  - Seems strange from a purely discriminative perspective (we've worked hard to assume as little as possible!)
  - **But remember:** regression emerged in statistics (where explicit assumptions are routine)
  - Like everything, such assumptions have pros and cons (**example:** easier to reason about non-iid data by modeling correlations in noise)

*random noise: unobserved, but sampled from assumed "error" distribution*

$$y = f^*(x) + \epsilon, \quad f^* \in F, \quad \epsilon \sim D_{\text{ERR}}$$

*true model: unknown, but from assumed model class*

# Linear regression

- In **linear regression**, we assume:

1. True model is **linear**:  $f^* \in \{f(x) = w^\top x : w \in \mathbb{R}^d\} = F$
2. Error distribution is **normal**:  $\epsilon \stackrel{iid}{\sim} N(0, \sigma^2)$ 
  - Variance  $\sigma^2$  is **unknown** but **uniform** across  $x$  (“homoscedastic”)
3. Features sampled from some (unknown)  $p(x)$

- **Conditional distributions become:**

$$y = w^\top x + \epsilon, \quad \epsilon \stackrel{iid}{\sim} N(0, \sigma^2) \quad \Rightarrow \quad P(y|x; w) = N(w^\top x, \sigma^2) \quad (\text{from linearity})$$

- **Interpretation:**

- Relationship between  $x$  and  $y$  is, in principle, linear
- Observations corrupted by many additive noise (e.g., a sum of many small errors + CLT = Normal!)

- **Take away:** structural assumptions are ok when we know something about the problem (conversely, with knowledge, not assuming anything (except iid) can be suboptimal)

# Learning

density  
↓

- Just made a *parametric* distributional assumption:  $P(y|x; w) = N(w^\top x, \sigma^2)$  for some  $w$
- **Q:** How can we use this to derive an appropriate learning objective?
- **A:** Maximize data *likelihood*
- **Observation:** two ways to interpret  $P(y|x; w)$ 
  - as function of **data**: given  $w$ , what is the probability of observing  $y$  given  $x$ ?
  - as function of **parameters**: given  $(x, y)$ , what it is the likelihood it was generated by  $w$ ?
- More generally, given sample set  $S = \{(x_i, y_i)\}_{i=1}^m$ , can define (as a function of  $w$ ):

$$\textbf{Likelihood: } L(w; S) := P(S; w) = P((x_1, y_1), \dots, (x_m, y_m); w)$$

- **Idea:** learn  $w$  that “best explains” data:

$$\textbf{Maximum-likelihood estimation (MLE): } \hat{w} = w_{\text{MLE}} = \operatorname{argmax}_{w \in \mathbb{R}^d} L(w; S)$$

# Maximum likelihood estimation

$$w_{\text{MLE}} = \operatorname{argmax}_w L(w; S)$$

$$= \operatorname{argmax}_w P(S; w)$$

$$= \operatorname{argmax}_w \prod_i P(y_i, x_i; w)$$

$$= \operatorname{argmax}_w \prod_i P(y_i | x_i; w) P(x_i | w)$$

$$= \operatorname{argmax}_w \prod_i P(y_i | x_i; w) P(x_i)$$

$$= \operatorname{argmax}_w \prod_i P(y_i | x_i; w)$$

$$= \operatorname{argmax}_w \log \prod_i P(y_i | x_i; w)$$

$$= \operatorname{argmax}_w \sum_i \log P(y_i | x_i; w)$$

*definition*

*iid*

*chain rule*

*x indep. of w*

$\prod_i P(x_i)$  is scalar

*log preserves argmax*

$\log \Pi = \Sigma \log$

$$= \operatorname{argmax}_w \sum_i \log \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(w^\top x_i - y_i)^2}{2\sigma^2}} \quad y|x \sim N(w^\top x, \sigma^2)$$

$$= \operatorname{argmax}_w \sum_i \left( \log \frac{1}{\sqrt{2\pi\sigma^2}} + \log e^{-\frac{(w^\top x_i - y_i)^2}{2\sigma^2}} \right)$$

$$= \operatorname{argmax}_w -\frac{1}{2\sigma^2} \sum_i (w^\top x_i - y_i)^2 \quad \text{indep. of } w; \log e^z = z$$

$$= \operatorname{argmin}_w \frac{1}{m} \sum_i (w^\top x_i - y_i)^2 \quad \text{negate; scalar prod.}$$

$$= \operatorname{argmin}_w \frac{1}{m} \sum_i \ell^{\text{sqr}}(y_i, w^\top x_i)$$

$$w_{\text{MLE}} = \operatorname{argmax}_{w \in \mathbb{R}^d} L(w; S) \quad \text{likelihood} = \operatorname{argmax}_{w \in \mathbb{R}^d} \log L(w; S) \quad \text{log-likelihood} = \operatorname{argmin}_{w \in \mathbb{R}^d} -\log L(w; S) \quad \text{negative log-likelihood (NLL)}$$

# Optimization

- MLE objective for linear regression is **Ordinary Least Squares**:

$$\operatorname{argmin}_{w \in \mathbb{R}} \frac{1}{m} \sum_{i=1}^m (y_i - w^\top x_i)^2$$


- **Possible ways to optimize:** (will not go in depth here)
  1. Gradient descent!  
(objective is convex – think why)
  2. Closed form  
(when  $X^\top X$  is invertible; as seen in numerical algorithms)
  3. Using SVD  
(when  $X^\top X$  is non-invertible ; as seen in numerical algorithms)

The story behind  
least squares

# Interpretation


- Consider some  $x_i$
- What would be a good prediction  $\hat{y}_i$ ?

●  $\hat{y}_i$ ?



# Interpretation

- Consider some  $x_i$
- What would be a good prediction  $\hat{y}_i$ ?
- To simplify, first assume we've observed multiple  $y_{ij} \sim p(y|x_i)$  for this particular  $x_i$
- What would be a good prediction now?





# Interpretation

- **Puzzle:**

Let  $y_1, \dots, y_n \in \mathbb{R}$ .

What is  $a \in \mathbb{R}$  which minimizes mean squared distances,  $\frac{1}{n} \sum_i (y_i - a)^2$ ?


$$\begin{aligned} \frac{1}{n} \sum_i (y_i - a)^2 &= \frac{1}{n} \sum_i y_i^2 - 2a \frac{1}{n} \sum_i y_i + a^2 & \dots &= a^2 - 2a\bar{y} + \bar{y}^2 - \bar{y}^2 + \bar{y} \\ &= \bar{y} - 2a\bar{y} + a^2 = \dots & &= (a - \bar{y})^2 - \bar{y}^2 + \bar{y} \end{aligned}$$

- **Solution:** mean squared distances are minimized by **the average,  $\bar{y}$**
- The average  $\bar{y}$  is a **statistic**: a useful way to summarize an entire sample as a single scalar

# Interpretation


- Back to our  $x_i$
- **Q:** What would be a good prediction  $\hat{y}_i$ ?

●  $\hat{y}_i$ ?




# Interpretation

- Back to our  $x_i$
- **Q:** What would be a good prediction  $\hat{y}_i$ ?
- **A:** If we observed multiple  $y_{ij} \sim p(y|x_i)$ , then  $\bar{y}_i$  would be a good prediction
- Unfortunately:
  - At *train time*, we only see one  $y_i$  per  $x_i$
  - At *test time*, we don't see any  $y_i$ -s!




# Interpretation

- Fortunately, training data includes *multiple*  $x_i$ , each with its own (single)  $y_i \sim P(y|x_i)$



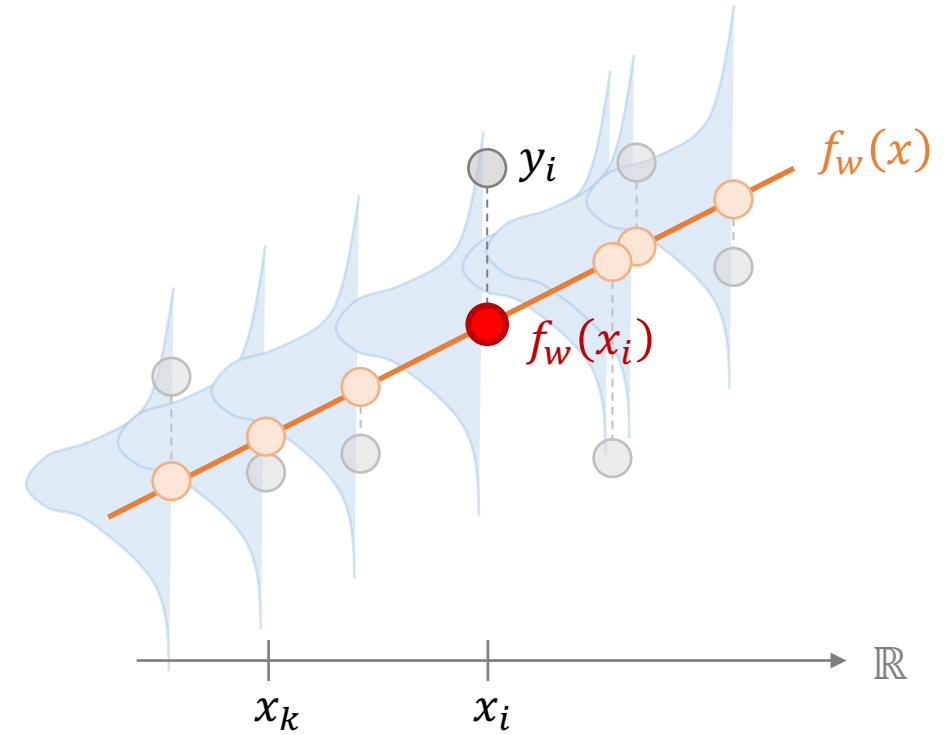
# Interpretation

- Fortunately, training data includes *multiple*  $x_i$ , each with its own (single)  $y_i \sim P(y|x_i)$
- Linear regression allows us to **share label information across examples**
- By assuming a parametric model  $f_w(x) = w^\top x$ , we can estimate  $\bar{y}_i$  for  $x_i$  using all other  $(x_j, y_j)$
- Minimizing mean squared errors  $(w^\top x - y)^2$  means we aim for the “line”  $f_w(x) = w^\top x$  to pass through all true averages  $\bar{y}_i$




# Interpretation

- Fortunately, training data includes *multiple*  $x_i$  each with its own (single)  $y_i \sim P(y|x_i)$
- Linear regression allows us to **share label information across examples**
- By assuming a parametric model  $f_w(x) = w^\top x$ , we can estimate  $\bar{y}_i$  for  $x_i$  using all other  $(x_j, y_j)$
- Minimizing mean squared errors  $(w^\top x - y)^2$  means we aim for the “line”  $f_w(x) = w^\top x$  to pass through all true averages  $\bar{y}_i$
- This comes from our assumption:  $y \sim N(\mu = w^\top x, \sigma^2)$




# Interpretation

- Fortunately, training data includes *multiple*  $x_i$  each with its own (single)  $y_i \sim P(y|x_i)$
- Linear regression allows us to **share label information across examples**
- By assuming a parametric model  $f_w(x) = w^\top x$ , we can estimate  $\bar{y}_i$  for  $x_i$  using other  $(x_j, y_j)$
- Minimizing mean squared errors  $(w^\top x - y)^2$  means we aim for the “line”  $f_w(x) = w^\top x$  to pass through all true averages  $\bar{y}_i$
- This comes from our assumption:  $y \sim N(\mu = w^\top x, \sigma^2)$
- $f_w(x) \approx \bar{y}$  only *estimates*  $\mu$  since we learn  $w$  from finite data



# Interpretation


- Estimates become **predictions**:  $\hat{y}_i = f_w(x_i) \approx \bar{y}_i$
- Ideally, we'd like  $\mathbb{E}[\hat{y}] = \mathbb{E}[\bar{y}] = \mu$  for new  $(x, y)$
- Hope is to approach this as  $m$  increases  
(i.e., more training examples  $(x_i, y_i)$ )
- **Won't work if we made the wrong assumptions!**





# Interpretation

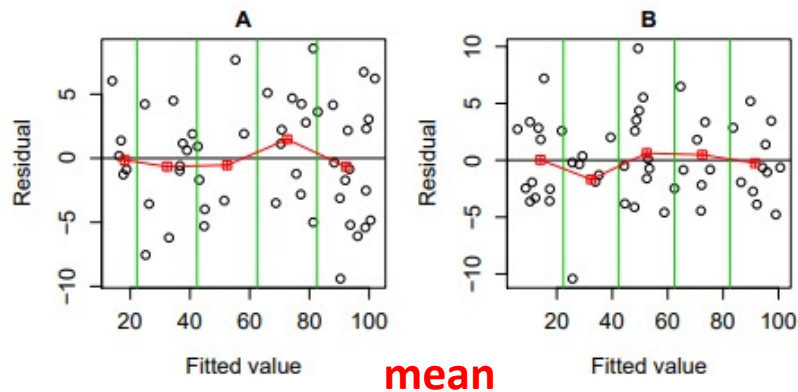
- Estimates become **predictions**:  $\hat{y}_i = f_w(x_i) \approx \bar{y}_i$
  - Ideally, we'd like  $\mathbb{E}[\hat{y}] = \mathbb{E}[\bar{y}] = \mu$  for new  $(x, y)$
  - Hope is to approach this as  $m$  increases  
(i.e., more training examples  $(x_i, y_i)$ )
  - **Won't work if we made the wrong assumptions!**
- 
- **Residual analysis** can help us understand  
(in hindsight) if our assumptions were sensible
  - Prediction “errors” are called **residuals**:  $r_i = \hat{y}_i - y_i$
  - Note “ $\mathbb{E}[\hat{y}] = \mathbb{E}[\bar{y}]$ ?” is like asking “ $\mathbb{E}[r] = 0$ ?”



# Residuals

- Recall our assumptions:  $y = w^\top x + \epsilon$ ,  $\epsilon \stackrel{iid}{\sim} N(0, \sigma^2)$


linear



# Residuals


- Recall our assumptions:  $y = w^T x + \epsilon$ ,  $\epsilon \stackrel{iid}{\sim} N(0, \sigma^2)$

linear



mean

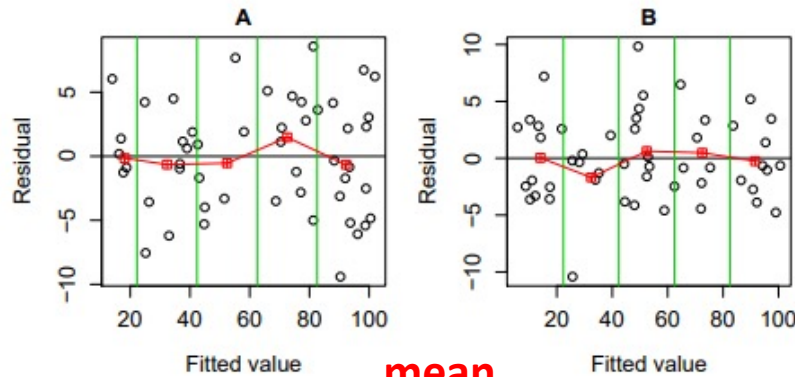
non-linear



# Residuals

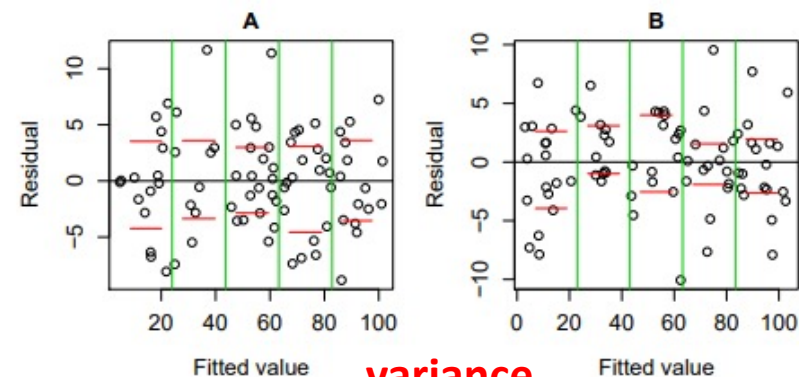
- Recall our assumptions:  $y = w^T x + \epsilon$ ,  $\epsilon \stackrel{iid}{\sim} N(0, \sigma^2)$

linear



mean

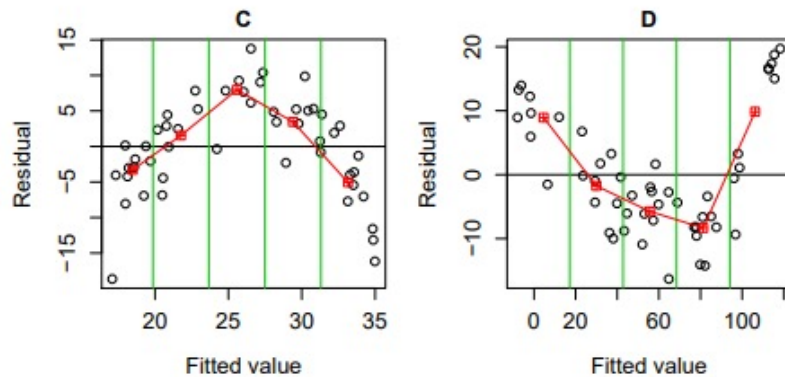
Fitted value



variance

homoscedastic

non-linear

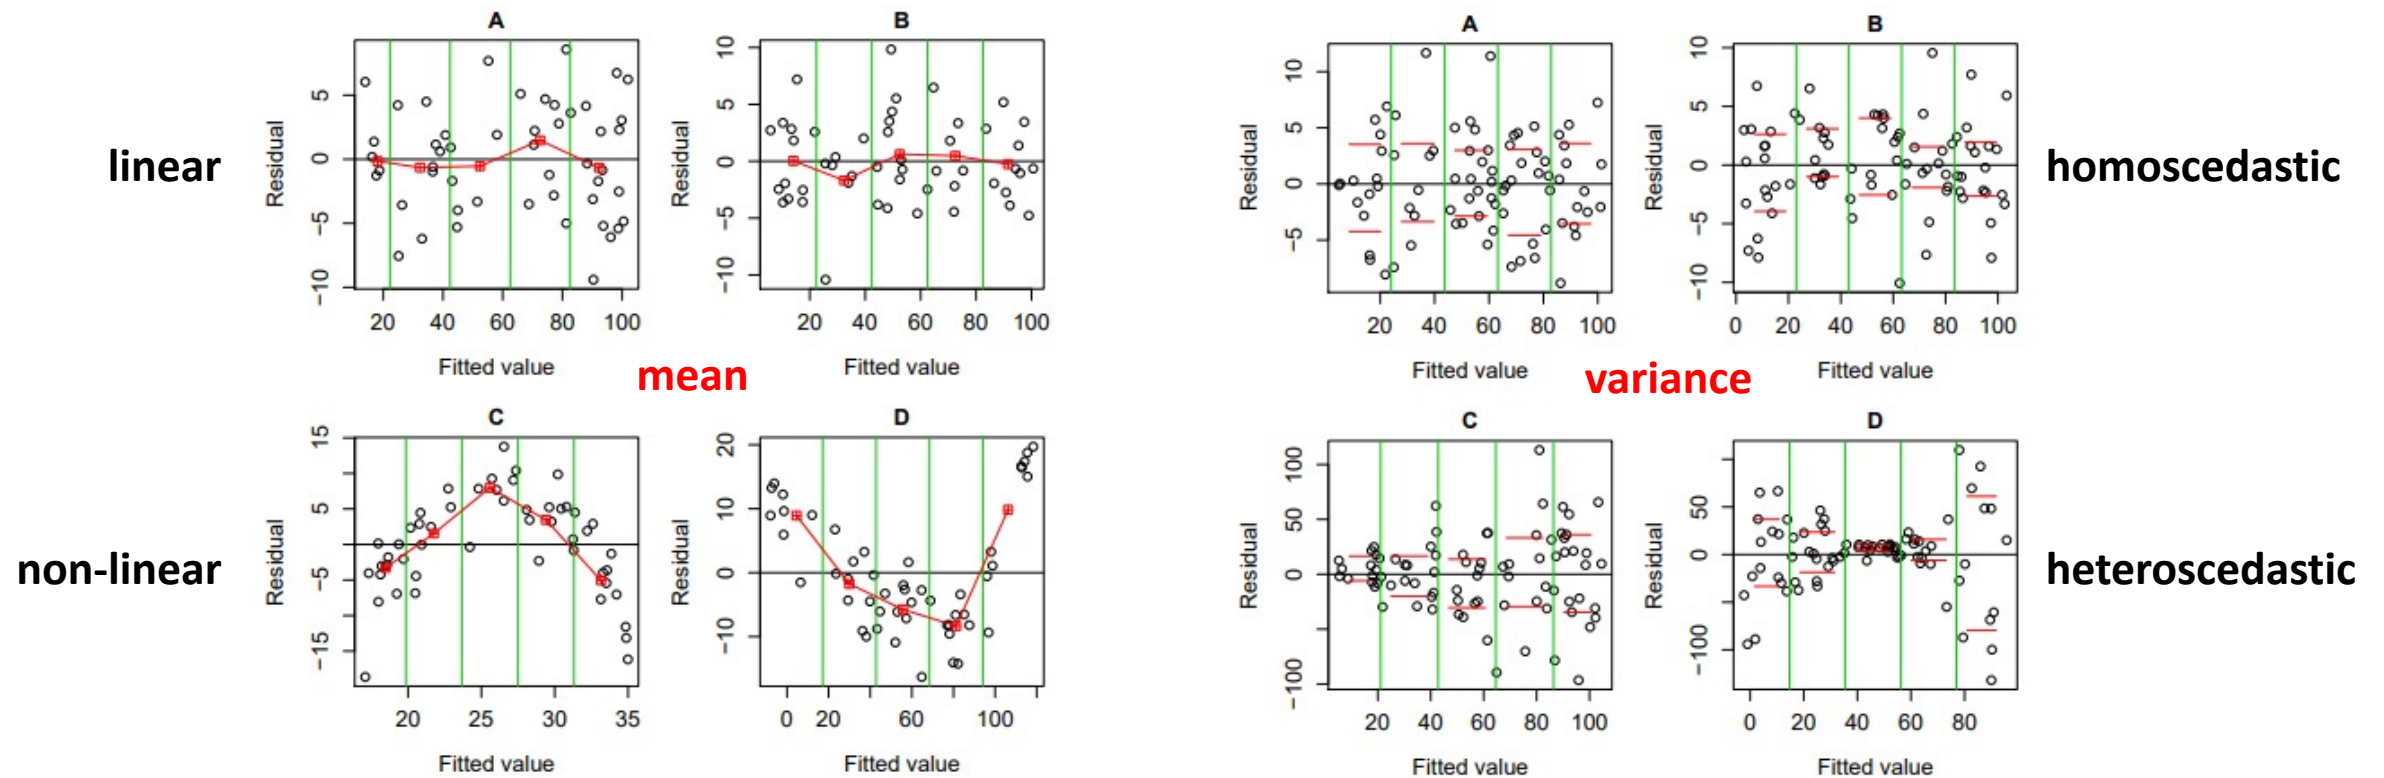


Fitted value

Fitted value

# Residuals

- Recall our assumptions:  $y = w^T x + \epsilon$ ,  $\epsilon \stackrel{iid}{\sim} N(0, \sigma^2)$



- Residuals will play big role next week!

Regression as loss minimization

**Regression** ( $y \in \mathbb{R}$ )

MLE objective:

$$\operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m \underbrace{(y_i - f(x_i))^2}_{= \ell^{\text{sqr}}(y, f(x))}$$

**Classification** ( $y \in \{\pm 1\}$ )

ERM objective:

$$\operatorname{argmin}_{h \in H} \frac{1}{m} \sum_{i=1}^m \underbrace{\mathbb{1}\{y_i \neq h(x_i)\}}_{= \ell^{0/1}(y, h(x))}$$

**Regression** ( $y \in \mathbb{R}$ )

~~MLE~~ **ERM objective!**

$$\operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m \underbrace{(y_i - f(x_i))^2}_{= \ell^{\text{sqr}}(y, f(x))}$$

**Classification** ( $y \in \{\pm 1\}$ )

ERM objective:

$$\operatorname{argmin}_{h \in H} \frac{1}{m} \sum_{i=1}^m \underbrace{\mathbb{1}\{y_i \neq h(x_i)\}}_{= \ell^{0/1}(y, h(x))}$$

- Squared error makes sense because:
  - correct prediction  $\Rightarrow$  loss=0
  - loss gradually increases with distance  $y - \hat{y}$
  - symmetric:  $\ell(y, y + a) = \ell(y, y - a)$
- Regression vs. classification – **just different losses!**
- (Or is it?)



**Regression** ( $y \in \mathbb{R}$ )

~~MLE~~ ERM objective!

$$\operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m \underbrace{(y_i - f(x_i))^2}_{= \ell^{\text{sqr}}(y, f(x))}$$

**Classification** ( $y \in \{\pm 1\}$ )

ERM objective:

$$\operatorname{argmin}_{h \in H} \frac{1}{m} \sum_{i=1}^m \underbrace{\mathbb{1}\{y_i \neq h(x_i)\}}_{= \ell^{0/1}(y, h(x))}$$

- Fundamental difference:  $f$  is scalar,  $h$  is binary

**Regression** ( $y \in \mathbb{R}$ )

~~MLE~~ ERM objective!

$$\operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m \underbrace{(y_i - f(x_i))^2}_{= \ell^{\text{sqr}}(y, f(x))}$$

**Classification** ( $y \in \{\pm 1\}$ )

ERM objective:

$$\operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m \underbrace{\mathbb{1}\{y_i \neq \text{sign}(f(x_i))\}}_{= \ell^{0/1}(y, f(x))}$$

- Fundamental difference:  $f$  is scalar,  $h$  is binary
- But...
  - most classifiers we've considered are based on scalar functions

### Regression ( $y \in \mathbb{R}$ )

MLE ERM objective!

$$\operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m \underbrace{(y_i - f(x_i))^2}_{= \ell^{\text{sqr}}(y, f(x))}$$

### Classification ( $y \in \{\pm 1\}$ )

ERM objective:

$$\operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m \underbrace{\max\{0, 1 - y_i f(x_i)\}}_{= \ell^{\text{hinge}}(y, f(x))}$$

- Fundamental difference:  $f$  is scalar,  $h$  is binary
- But...
  - most classifiers we've considered are based on **scalar functions**
  - and what we actually optimize is a **continuous proxy loss**
- **Surprise:** we've been doing regression all along!

# Classification vs. regression

- On its face, difference appears minor:

- **Classification:**  $y \in \{\pm 1\} \Rightarrow \ell^{0/1}$

- **Regression:**  $y \in \mathbb{R} \Rightarrow \ell^{\text{sqr}}$

(and most learning textbooks end the discussion here)


- But digging deeper, it turns out that:
  - some things remain exactly the same
  - some things differ but are easy to adapt
  - some things are fundamentally different
- Let's see *what* changes, *how*, and *why* (and what to do about it!)

# 1) Losses and proxies

- **Classification:** optimize hinge loss, want 0/1 loss
- **Regression:** optimize squared loss, want... squared loss?
- In regression we typically don't need a proxy loss
- Problem is natively continuous! (this is good news)
- Crux is that there is no “natural” performance measure  
(i.e., why  $(y - \hat{y})^2$ , and not  $(y - \hat{y})^4$ ? or  $|y - \hat{y}|$ ? or  $|y - \hat{y}|^{1.5}$ ? or  $|y - \hat{y}|^{0.5}$ ? or ...)
- Statistical modeling gives partial answer  
(that is not entirely satisfactory from a discriminative perspective)
- But still raises the issue – *how should we interpret performance results?*

## 2) The meaning of error

- **Classification:** 0/1-error = 0.08  $\Rightarrow$  92% future predictions are correct
- **Regression:** sqr-error = 0.08  $\Rightarrow$  ???
- **Q:** How can we
  - *interpret evaluated performance?*
  - *meaningfully compare performance?*
- **A:**  $R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2} = 1 - \frac{SS_{\text{RES}}}{SS_{\text{TOT}}}$   
 $= \frac{1}{m} \sum_i y_i$
- Measures how much of the variance the model can explain by using the features  $x_i$   
(think of  $\bar{y}$  as “best guess” if you observed only  $y$ -s but no  $x$ -s)
- (Recall that even 0-1 error is not always meaningful  
(e.g., imbalanced data) and other measures are needed)



## 2) The meaning of error

- **Classification:** 0/1-error = 0.08  $\Rightarrow$  92% future predictions are correct

- **Regression:** sqr-error = 0.08  $\Rightarrow$  ???

- **Q:** How can we





- *interpret evaluated performance?*
- *meaningfully compare performance?*

- **A:**  $R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2} = 1 - \frac{SS_{\text{RES}}}{SS_{\text{TOT}}}$   
 $= \frac{1}{m} \sum_i y_i$

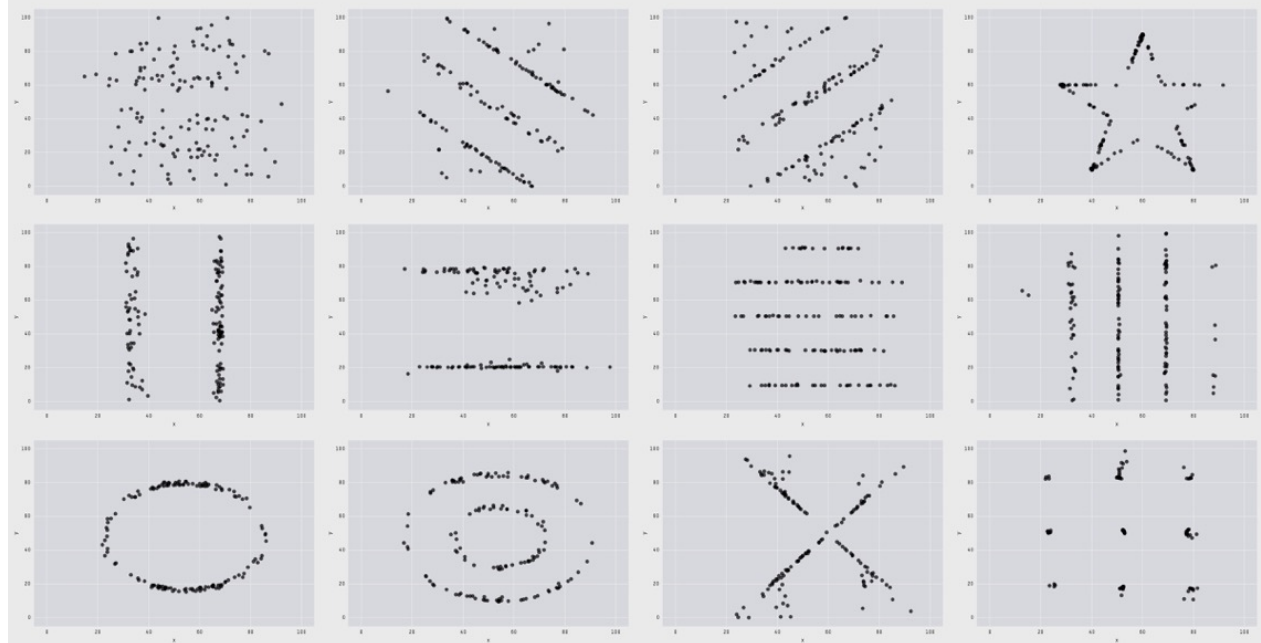
- Measures how much of the variance the model can explain by using the features  $x_i$   
(think of  $\bar{y}$  as “best guess” if you observed only  $y$ -s b

- (Recall that even 0-1 error is not always meaningful  
(e.g., imbalanced data) and other measures are needed)

But even  $R^2$  has its limits...



X Mean: 54.26  
Y Mean: 47.83  
X SD : 16.76  
Y SD : 26.93  
Corr. : -0.06






### 3) Feature transformations

- **Classification:** in theory, unaffected by monotone feature transforms  
(in practice, this does matter, e.g. running GD on proxy loss)
- **Regression:** not true!
- **Example:** threshold classifiers –  $\text{sign}(f(x)) = \text{sign}(\alpha f(x)) \quad \forall \alpha \in \mathbb{R}$
- In regression, direct evaluation of  $f(x)$  (i.e., no sign) means:
  - loss is sensitive to transforms
  - statistical modeling assumptions mean different things
- **Example:**
  - Consider  $x \mapsto \log x$  (assume  $x > 0$ )
  - Compare:  $y = a + bx \Rightarrow$  one unit change in  $x$  results in  $b$  units change in  $y$   
 $y = a + b \log x \Rightarrow$  one % change in  $x$  results in  $\frac{b}{100}$  units change in  $y$

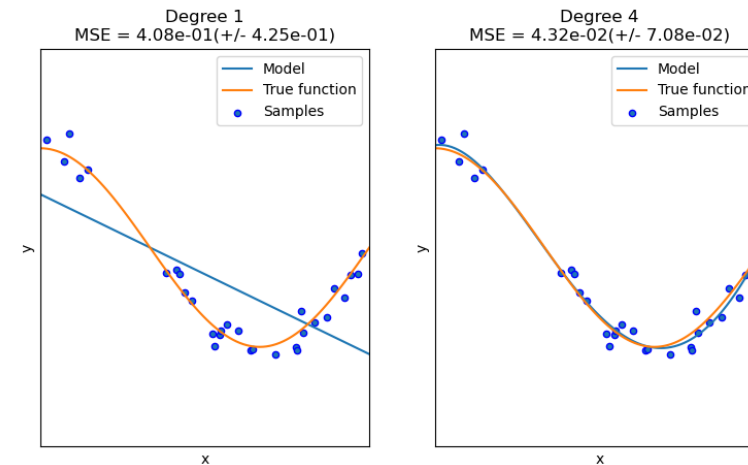
## 4) Generalization

- **Classification:** overfitting: (i) quantified by VC, and (ii) reduced by norm regularizers
- **Regression:**
  - VC theory does not apply (specific to binary classification – think what shattering does!)
  - Min-norm as max-margin does not apply (there is no longer even a notion of “margin”)
- **Q:** Does overfitting even happen?
- **A:** Most certainly yes! (and easier to draw)



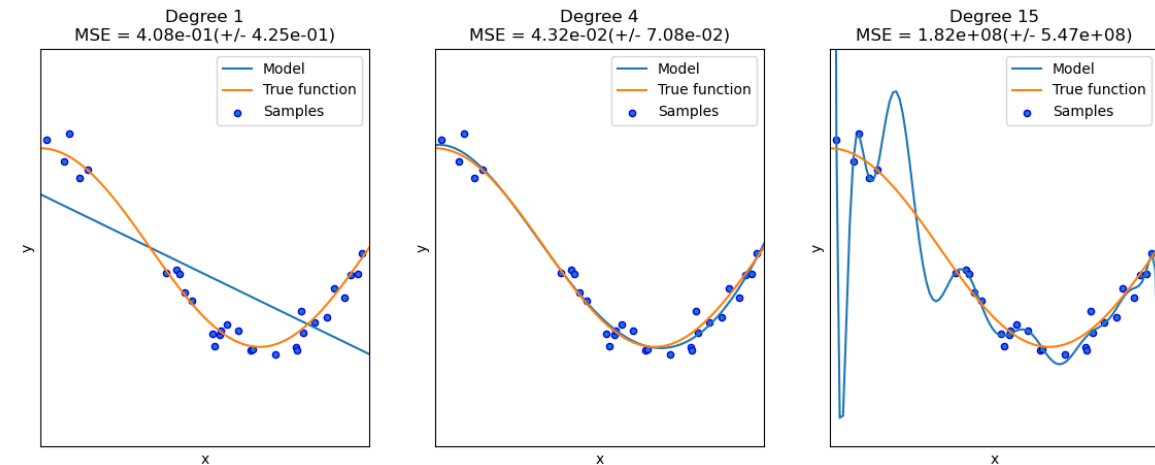
## 4) Generalization

- **Classification:** overfitting: (i) quantified by VC, and (ii) reduced by norm regularizers
- **Regression:**
  - VC theory does not apply (specific to binary classification – think what shattering does!)
  - Min-norm as max-margin does not apply (there is no longer even a notion of “margin”)
- **Q:** Does overfitting even happen?
- **A:** Most certainly yes! (and easier to draw)



# 4) Generalization

- **Classification:** overfitting: (i) quantified by VC, and (ii) reduced by norm regularizers
- **Regression:**
  - VC theory does not apply (specific to binary classification – think what shattering does!)
  - Min-norm as max-margin does not apply (there is no longer even a notion of “margin”)
- **Q:** Does overfitting even happen?
- **A:** Most certainly yes! (and easier to draw)
- **Q:** Do norm regularizers help with overfitting?
- **A:** Yes! More on this in tirgul.



# 5) Least Squares vs Least Squares

## Empirical Risk Minimization (ERM):

- Learning objective – **empirical loss**:

$$f_{\text{ERM}} = \underset{f \in F}{\operatorname{argmin}} \frac{1}{m} \sum_{i=1}^m (y_i - f(x_i))^2$$

## Maximum Likelihood Estimation (MLE):

- Learning objective – **likelihood**:

$$f_{\text{MLE}} = \underset{f \in F}{\operatorname{argmin}} \frac{1}{m} \sum_{i=1}^m (y_i - f(x_i))^2$$

# 5) Least Squares vs Least Squares

## Empirical Risk Minimization (ERM):

- Learning objective – **empirical loss**:

$$f_{\text{ERM}} = \operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m (y_i - f(x_i))^2$$

- Loss function:  $\ell^{\text{sqr}}(y, \hat{y}) = (y - \hat{y})^2$
- Distribution-independent
- Arbitrary  $H$
- Care about:
  - minimizing expected loss
  - generalization:  $\mathbb{E}[(y - f_{\text{ERM}}(x))^2]$
  - finite sample performance

## Maximum Likelihood Estimation (MLE):

- Learning objective – **likelihood**:

$$f_{\text{MLE}} = \operatorname{argmin}_{f \in F} \frac{1}{m} \sum_{i=1}^m (y_i - f(x_i))^2$$

- Statistical model:  $y = f^*(x) + \epsilon$ 
  - Assumed “true” model  $f^* \in F$
  - Assumed error distribution  $\epsilon \sim D_{\text{ERR}}$
- Care about: (for example; out of our scope)
  - consistency (asymptotic property)
  - identifiability (asymptotic property)
  - confidence intervals


**take away:** same same (formula) but different (story; and language)

Statistical modeling, revisited

# Other noise models

- We saw Normal noise  $\Rightarrow$  squared loss
- **Q:** What happens when assuming other forms of noise?
- **Example:**

$$\epsilon \sim \text{Laplace}(0, \eta)$$






# Other noise models

- We saw Normal noise  $\Rightarrow$  squared loss
- **Q:** What happens when assuming other forms of noise?
- **Example:**

$$\epsilon \sim \text{Laplace}(0, \eta) = \frac{1}{2\eta} e^{-\frac{|x|}{\eta}}$$




# Other noise models

- We saw Normal noise  $\Rightarrow$  squared loss
- **Q:** What happens when assuming other forms of noise?

- **Example:**


$$\epsilon \sim \text{Laplace}(0, \eta) = \frac{1}{2\eta} e^{-\frac{|x|}{\eta}}$$

- **A:** Gives absolute loss  $\ell^{\text{abs}}(y, \hat{y}) = |y - \hat{y}|$
- (To get this, just follow MLE derivation for Normal)
- Called ***least absolute deviation***
- Solution estimate **median** of  $p(y|x)$  (vs. average in squared loss)
  - **pros:** more robust to outliers
  - **cons:** not smooth



# What about classification?

- Can statistical modelling handle classification? ( $y \in \{0,1\}$ )
- In principle, can just use linear regression... (it's well defined)
- ...but it's a little silly!
- E.g., we know  $y$  is bounded in  $[0,1]$ , but  $f(x)$  is not
- **Fix:** change assumption on  $p(y|x)$



# Logistic regression

- In **linear regression**, we:

1. Assumed  $y|x \sim N(\mu_x, \sigma^2)$ , where  $\mu_x = \mathbb{E}[Y|X = x] \in \mathbb{R}$
2. Modeled  $\hat{y} = \mathbb{E}[Y|X = x] = w^\top x$

- In **logistic regression**:

1. Assume  $y|x \sim \text{Bernoulli}(p_x)$ , for some  $p_x \in [0,1]$

- This gives:

$$P(y|x) = \begin{cases} p_x & \text{if } y = 1 \\ 1 - p_x & \text{if } y = 0 \end{cases}$$


- Recall Bernoulli  $\Rightarrow \mathbb{E}[Y|X = x] = p_x$

2. Can now model  $f(x) = \mathbb{E}[Y|X = x] = p_x$

# Logistic regression

- We would like to model  $f(x) = \mathbb{E}[Y|X = x] = p_x$
- Need  $p_x$  to be parametric:  $p_x = f(x; w)$
- This allows to *learn*  $w$  from data with MLE
- Keep it simple: use *linear*  $w^\top x$ !
- **Q:** Would  $p_x = w^\top x$  work?
- **A:** No!  $p_x \in [0,1]$ , but  $w^\top x$  isn't
- **Idea:** model probabilities by passing  $w^\top x$  through “sigmoid”

$$p_x(w) = P(Y = 1|x; w) = \frac{1}{1 + e^{-w^\top x}} := \sigma(w^\top x) \in [0,1]$$




# Logistic regression

- We would like to model  $f(x) = \mathbb{E}[Y|X = x] = p_x$
- Need  $p_x$  to be parametric:  $p_x = f(x; w)$
- This allows to *learn*  $w$  from data with MLE
- Keep it simple: use *linear*  $w^\top x$ !
- **Q:** Would  $p_x = w^\top x$  work?
- **A:** No!  $p_x \in [0,1]$ , but  $w^\top x$  isn't

- **Idea:** model probabilities by passing  $w^\top x$  through “sigmoid”

$$p_x(w) = P(Y = 1|x; w) = \frac{1}{1 + e^{-w^\top x}} := \sigma(w^\top x) \in [0,1]$$

- Relies on common empirical observation:  
as  $p$  reaches extremes, changes in  $x$  matter less



# Logistic regression


- Learning  $w$  gives us a **predictor**:

$$\hat{p} = \sigma(w^\top x + b) = \frac{1}{1 + e^{-w^\top x + b}} \in [0,1]$$


- $\sigma$  is called a **logistic function**
- Note that labels  $y \in \{\pm 1\}$  are *binary*,  
but **predictions**  $\hat{p} \in [0,1]$  are *probabilities*
- To classify, can threshold:  $\hat{y} = \mathbb{1}\{\hat{p} > 0.5\}$
- But what does  $w^\top x$  "mean"?
- Interpretation: invert to get *linear log-odds*:

$$w^\top x = \log \frac{p(x)}{1-p(x)} \quad (\text{a.k.a. "logit"})$$

$$\text{logistic: } \frac{1}{1+e^{-z}}$$



$$\text{logit/log-odds: } \log \frac{p}{1-p}$$



# Learning objective

$$NLL(w; S) = -\log P(S; w)$$

= ...

$$= -\sum_i \log P(y_i | x_i; w) = \begin{cases} \hat{p}_i & \text{if } y_i = 1 \\ 1 - \hat{p}_i & \text{if } y_i = 0 \end{cases}$$

(remember  $\hat{p}_i$  depends on  $w$ )

$$= -\sum_i \log \hat{p}_i^{y_i} (1 - \hat{p}_i)^{1-y_i}$$

$$= -\sum_i y_i \log \hat{p}_i + (1 - y_i) \log(1 - \hat{p}_i)$$

“soft” prediction  
 $\hat{p} \in [0,1]$

- **Cross-entropy loss:**

$$\ell^{\text{CE}}(y, \hat{p}) = -y \log \hat{p} - (1 - y) \log(1 - \hat{p})$$

- **Interpretation:** logistic regression =

`cross_entropy(sigmoid(linear))`

- Nonetheless – **convex** in  $w$ !

(so can optimize with gradient descent)

- Naturally extends to multiclass

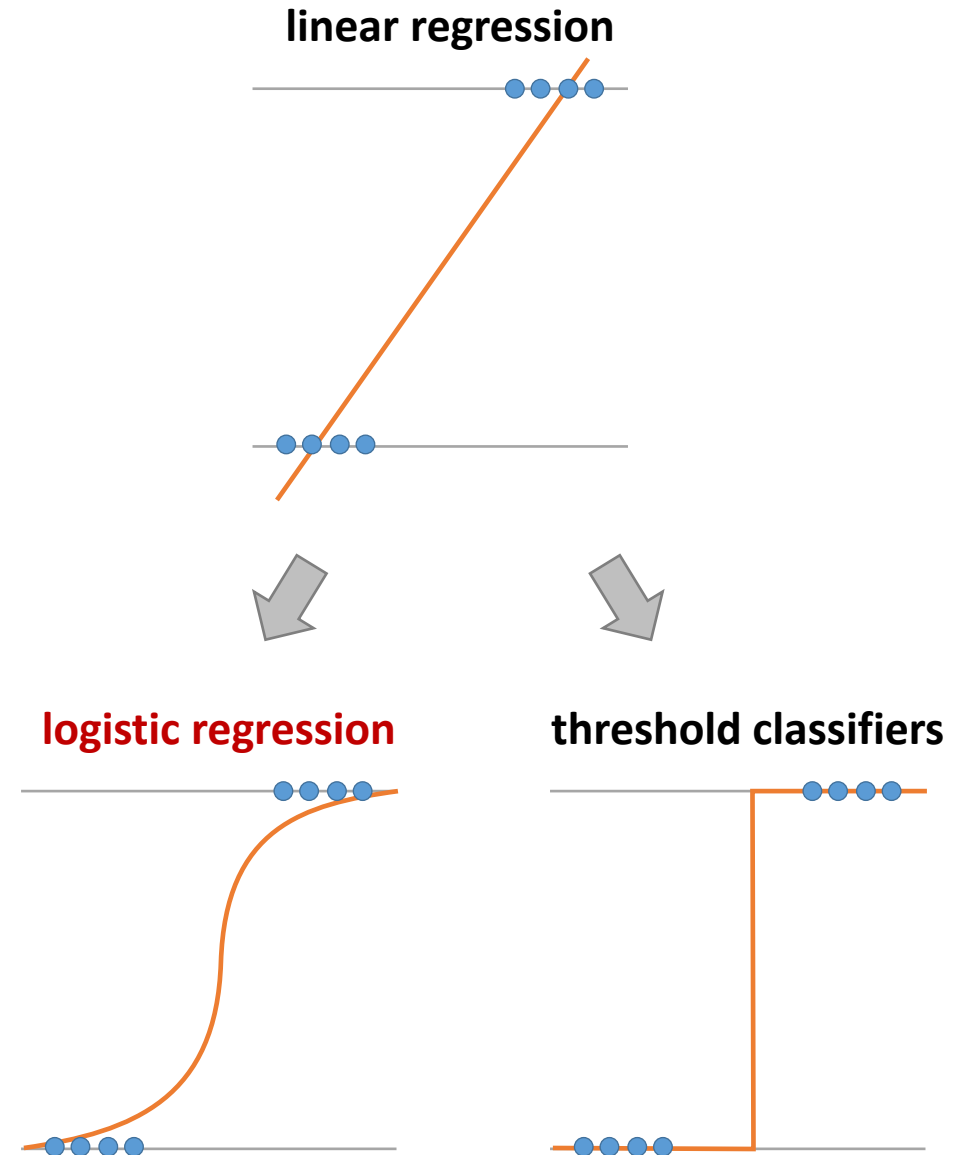
(think multinomial distribution)

- Very popular in deep learning



# Vs. discriminative

- **Statistical modeling** handles classification by wrapping linear model with *logistic* “link function”
- Compare this to **discriminative ML**, which wraps a linear model with a *step function* (e.g., sign, or  $\mathbb{1}$ )
- Not so different after all!
- In fact, sigmoid  $\rightarrow$  step function in the limit:
$$\sigma_{\alpha}(z) = \frac{1}{1 + e^{-\alpha z}} \xrightarrow[\alpha \rightarrow \infty]{} \mathbb{1}\{z > 0\}$$
- Since  $\ell^{\text{CE}}$  is continuous, can use as proxy for 0-1 loss (just like hinge is!)



# Next week

- **Part III:** *more supervised learning*
  1. Regression (today)
  2. Bagging and boosting ← *residuals!*
  3. Deep learning

