

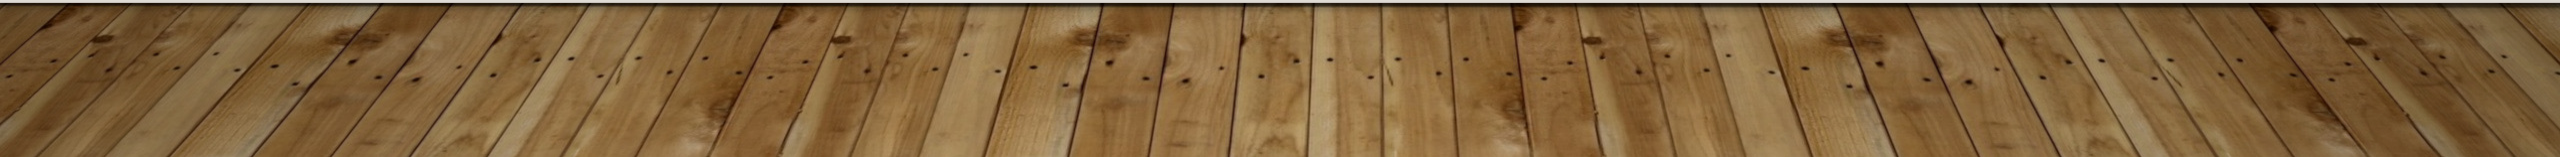
Introduction to Machine Learning (IML)

# LECTURE #10: BAGGING AND BOOSTING

---

236756 – 2023-2024 WINTER – TECHNION

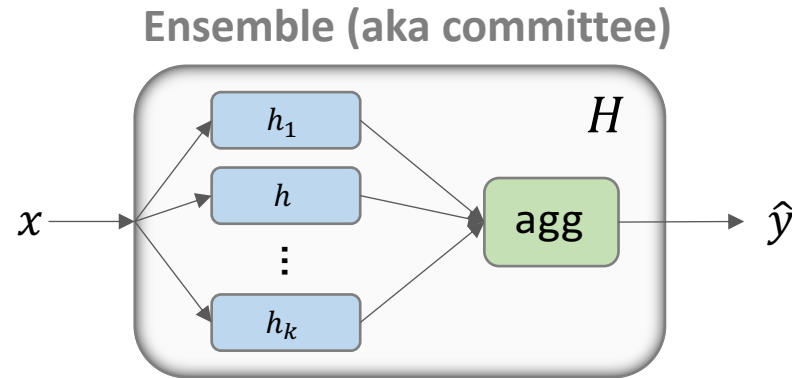

LECTURER: YONATAN BELINKOV



# Today

- **Part III:** *more supervised learning*
  1. Regression
  2. Bagging and boosting (today)
  3. Generative models
  4. Deep learning
- Back to a discriminative frame-of-thought
- Will consider classification and regression interchangeably (i.e., use a different loss)

# Ensemble methods



- **Ensemble methods:**

aggregate multiple models into a single, (hopefully) more powerful model

- **Q:** are ensembles more powerful than any individual models?
- **A:** it's complicated
- Today's goal is to shed light on this question
- **Intuitively, ensembling works well when:**
  - Models differ in strengths and weaknesses (e.g., accurate on different parts of  $D$ )
  - Aggregation amplifies strengths and mitigates weaknesses
  - Optimization is likely to work reasonably well

# Ensemble methods

- Template (and **today**):

## 1. Base-class learner

*ID3 (decision trees)*

## 2. Aggregation rule

*linear:  $H = \sum_i \alpha_i h_i$*

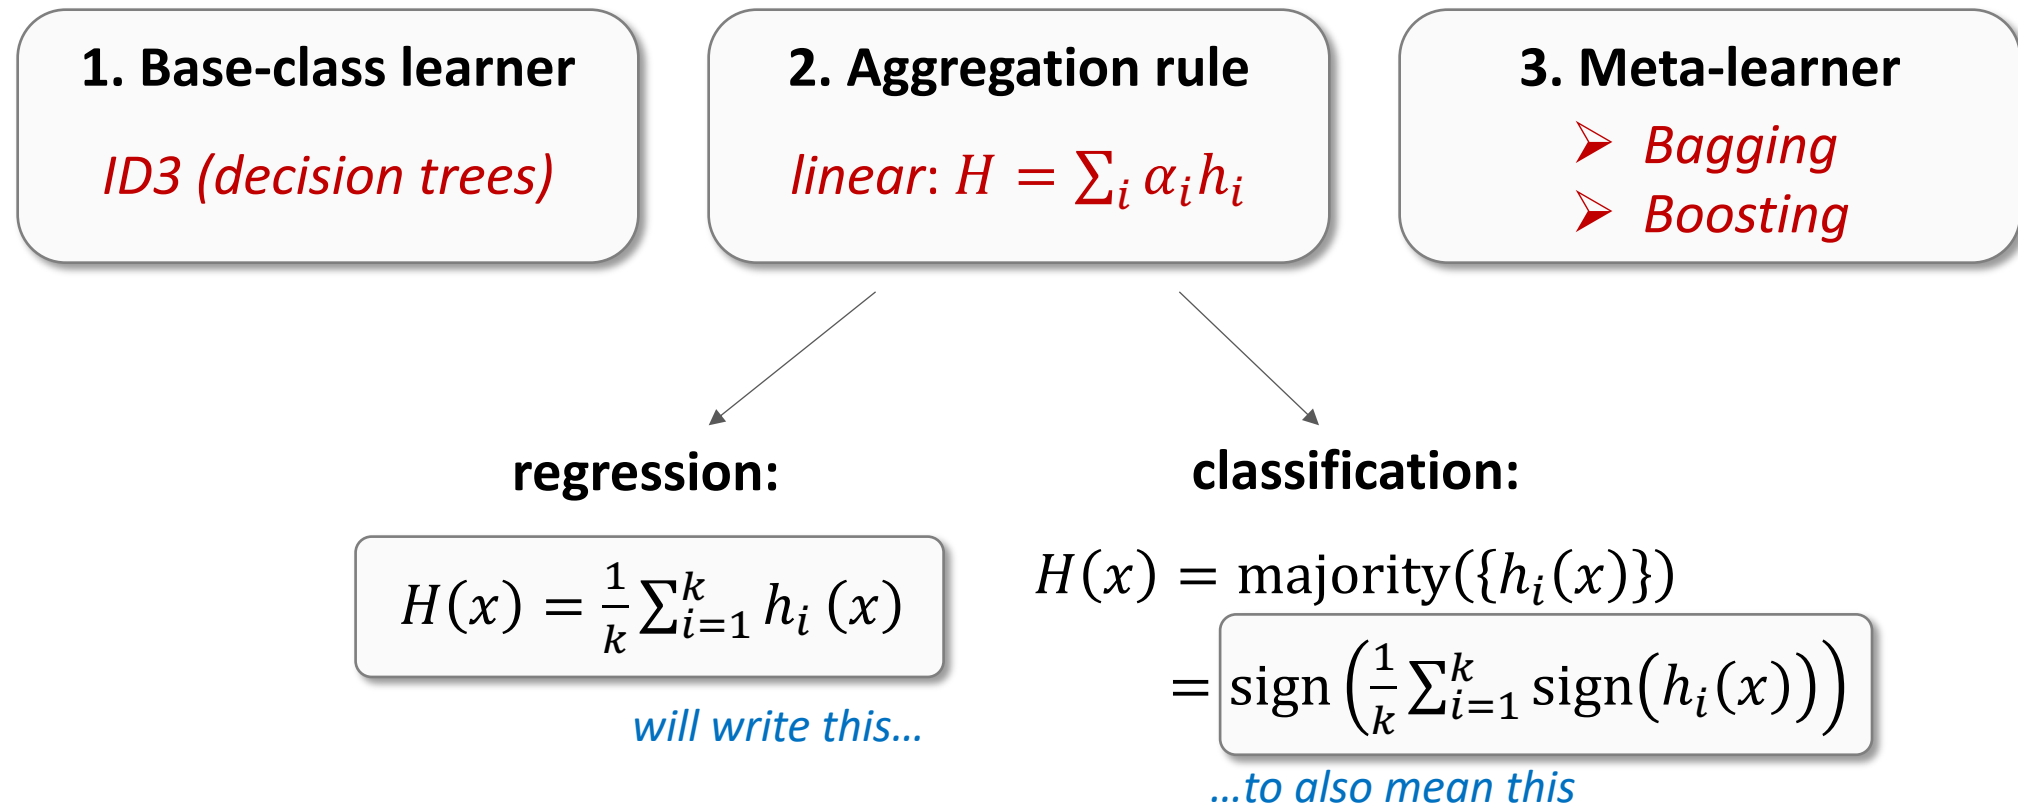
## 3. Meta-learner

- *Bagging*
- *Boosting*

- **Base model class:**  $h \in \mathcal{H}$  (classification, regression, ...)
- **Base-class learner** (=algorithm):  $h_S = A(S)$
- **Linear ensemble model:**  $H(x) = \sum_i \alpha_i h_i(x)$ ,  $\alpha_i \in \mathbb{R}, h_i \in \mathcal{H}$
- **Meta-learner** (=algorithm):  $H_S = M(S)$  (learns  $\alpha_i, h_i$ )

# Ensemble methods

- Template (and **today**):



# Ensemble methods

- Template (and **today**):

## 1. Base-class learner

*ID3 (decision trees)*

## 2. Aggregation rule

*linear:  $H = \sum_i \alpha_i h_i$*

## 3. Meta-learner

- *Bagging*
- *Boosting*

- **Q:** Does it make sense to aggregate linear models?
- **A:** Question is ill-posed
  - Need to also say what base learner is (e.g., ERM, RLM, SVM)
  - For linear classes, ensembles  $H$  are also linear! so member of the base class
  - But this does not imply that learning  $H$  is the same as learning  $h$
- We will return to this!

# Bagging

# Bagging

- **Recall:** bias-variance tradeoff

$$\mathbb{E}_{S \sim D^m} [L_D^{sqr}(h_S)] = \underbrace{\mathbb{E}_{x,y} [(\bar{y}(x) - y)^2]}_{\text{expected error}} + \underbrace{\mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2]}_{\text{noise}} + \underbrace{\mathbb{E}_{S,x} [(h_S(x) - \bar{h}(x))^2]}_{\text{bias}^2} + \underbrace{\mathbb{E}_{S,x} [(h_S(x) - \bar{h}(x))^2]}_{\text{variance}}$$

- where:
  - Expected label:  $\bar{y}(x) = \mathbb{E}_{y \sim D_{Y|X=x}}[y] \in [0,1]$
  - Expected classifier:  $\bar{h} = \mathbb{E}_{S \sim D^m}[h_S]$




# Bagging

- **Recall:** bias-variance tradeoff

$$\mathbb{E}_{S \sim D^m} [L_D^{sqr}(h_S)] = \underbrace{\mathbb{E}_{x,y} [(\bar{y}(x) - y)^2]}_{\text{noise}} + \underbrace{\mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2]}_{\text{bias}^2} + \underbrace{\mathbb{E}_{S,x} [(h_S(x) - \bar{h}(x))^2]}_{\text{variance}}$$

*expected error*



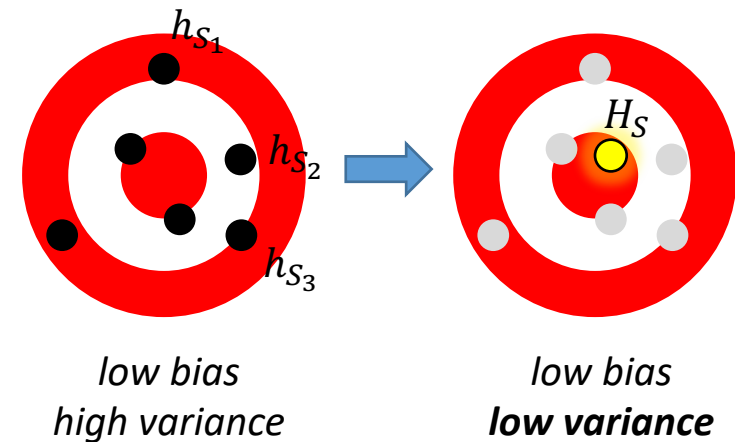
# Bagging

- **Recall:** bias-variance tradeoff

$$\mathbb{E}_{S \sim D^m} [L_D^{sqr}(h_S)] = \underbrace{\mathbb{E}_{x,y} [(\bar{y}(x) - y)^2]}_{\text{expected error}} + \underbrace{\mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2]}_{\text{noise}} + \underbrace{\mathbb{E}_{S,x} \left[ \left( \cancel{h_S(x)} - \bar{h}(x) \right)^2 \right]}_{\text{bias}^2 + \text{variance}}$$



The equation shows the decomposition of the expected squared loss into three components: noise (irreducible error), bias squared, and variance. The variance component is highlighted in a box, and the term  $h_S(x)$  is crossed out with a red slash, with a red arrow pointing to it from the label  $H_S$  above.

- Bagging is an ensembling approach aimed at **reducing variance** *without decreasing complexity*
- **Idea:** learn an  $H_S$  that tries to “approximate”  $\bar{h}$
- Bagging works well for *high-variance, low-bias* base model classes
- For example, **deep decision trees**



# Decision trees

- **Recall:** *decision trees*
- **ID3 learning algorithm:**
  1. Start with root as leaf.
  2. For each feature, compute best possible split (using entropy).
  3. Create decision node (and split data) using best feature.
  4. Recurse.
- **Trees can be highly non-linear**
- Without restrictions, always exists a tree  $T$  having zero empirical error,  $L_S(H_T) = 0$
- Generally, large depth implies **high variance**



# Variance reduction via averaging

- **Bagging** reduces variance by **averaging**.
- **Recall:** law of large numbers roughly states that:

If  $z_i$  are i.i.d. with mean  $\mu$  and variance  $\sigma^2$ , then

1.  $\bar{z}_{(k)} = \frac{1}{k} \sum_{i=1}^k z_i \xrightarrow[k \rightarrow \infty]{} \mu$
2.  $\text{Var}(\bar{z}_{(k)}) = \frac{\sigma^2}{k}$

- **Idea:** average multiple models

# Variance reduction via averaging

- **Bagging** reduces variance by **averaging**.
- **Recall:** law of large numbers roughly states that:

If  $h_i$  are i.i.d. with mean  $\mu$  and variance  $\sigma^2$ , then\*

1.  $\bar{h}_{(k)} = \frac{1}{k} \sum_{i=1}^k h_i \xrightarrow[k \rightarrow \infty]{} \mu \rightarrow \text{bias remains roughly the same}$

2.  $\text{Var}(\bar{h}_{(k)}) = \frac{\sigma^2}{k} \rightarrow \text{variance diminishes quickly! (compared to base models)}$

- **Idea:** average multiple models
- Can hope for error (or at least its **variance** term) to decrease at rate  $\approx \frac{1}{k}$

$$\text{bias}^2: \mathbb{E}_x \left[ (\bar{h}(x) - \bar{y}(x))^2 \right]$$

$$\text{variance: } \mathbb{E}_{S,x} \left[ (h_S(x) - \bar{h}(x))^2 \right]$$

\* model mean and variance are not well-defined here; a proper treatment requires more elaborate tools

# Bootstrapping

- **Q:** How can we obtain  $k$  different models  $h_1, \dots, h_k$ ?
- **A:** (obvious) train on  $k$  different (iid) data sets  $S_1, \dots, S_k$ :  
$$h_1 = A(S_1), \dots, h_k = A(S_k)$$
- (recall  $A$  is our *base learner*)
- **Problem:** we don't have multiple (iid) data sets  $S_1, \dots, S_k$  – just one!
- **Naïve solution:** partition  $S$  into  $k$  sets of size  $m/k$
- **Think:** can/should/will this lead to improved performance?
- **Better solution:** use single  $S$  to “simulate” multiple sets by repeated sub-sampling
  - Many ways to subsample
  - Bagging uses ***bootstrapping***

# Bootstrapping

- **Bootstrap** = sample uniformly with replacement: (=“take out and put back in”)

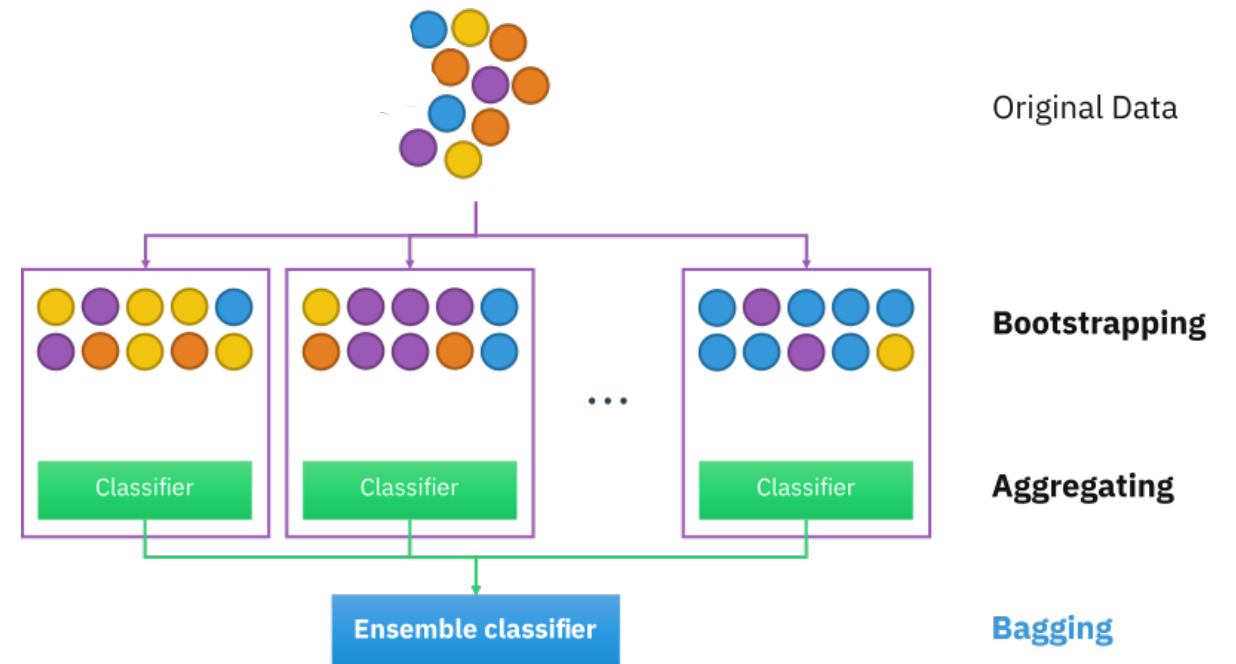
- Define distribution  $Q((x, y)|S) = \begin{cases} \frac{1}{m} & (x, y) \in S \\ 0 & o.w. \end{cases}$

- Sample  $S_i \sim Q^m$  iid

- Then:

- Train  $h_i = A(S_i)$
- Aggregate  $H = \text{agg}(h_1, \dots, h_k)$   
(by averaging)

- Bagging = **B**ootstrap **A**ggregating





# Example

single sample set  $S$





single tree trained on  $S$

other trees trained on  
other sample sets  $S'$



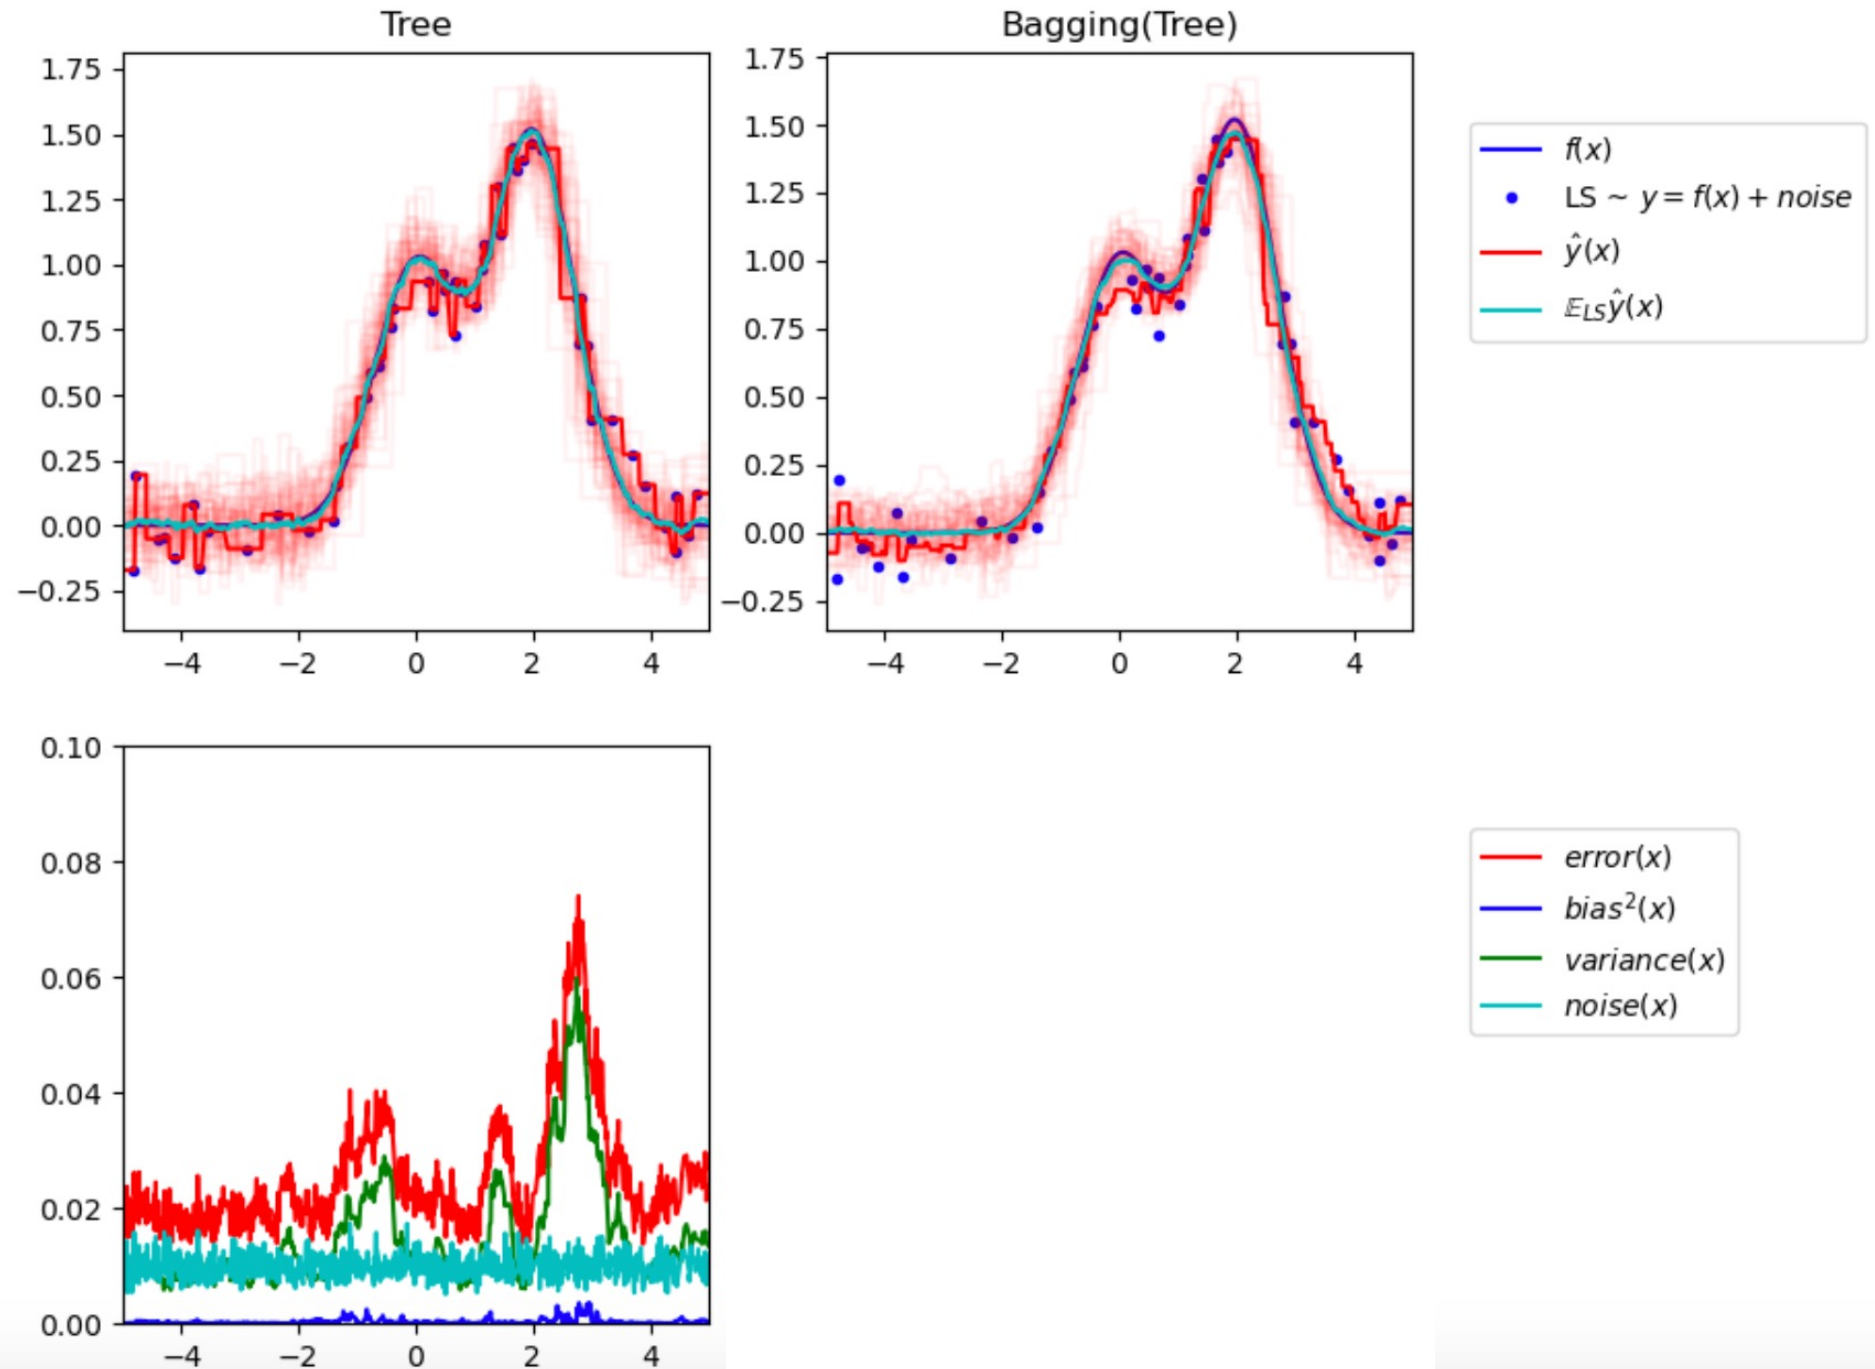


# Example

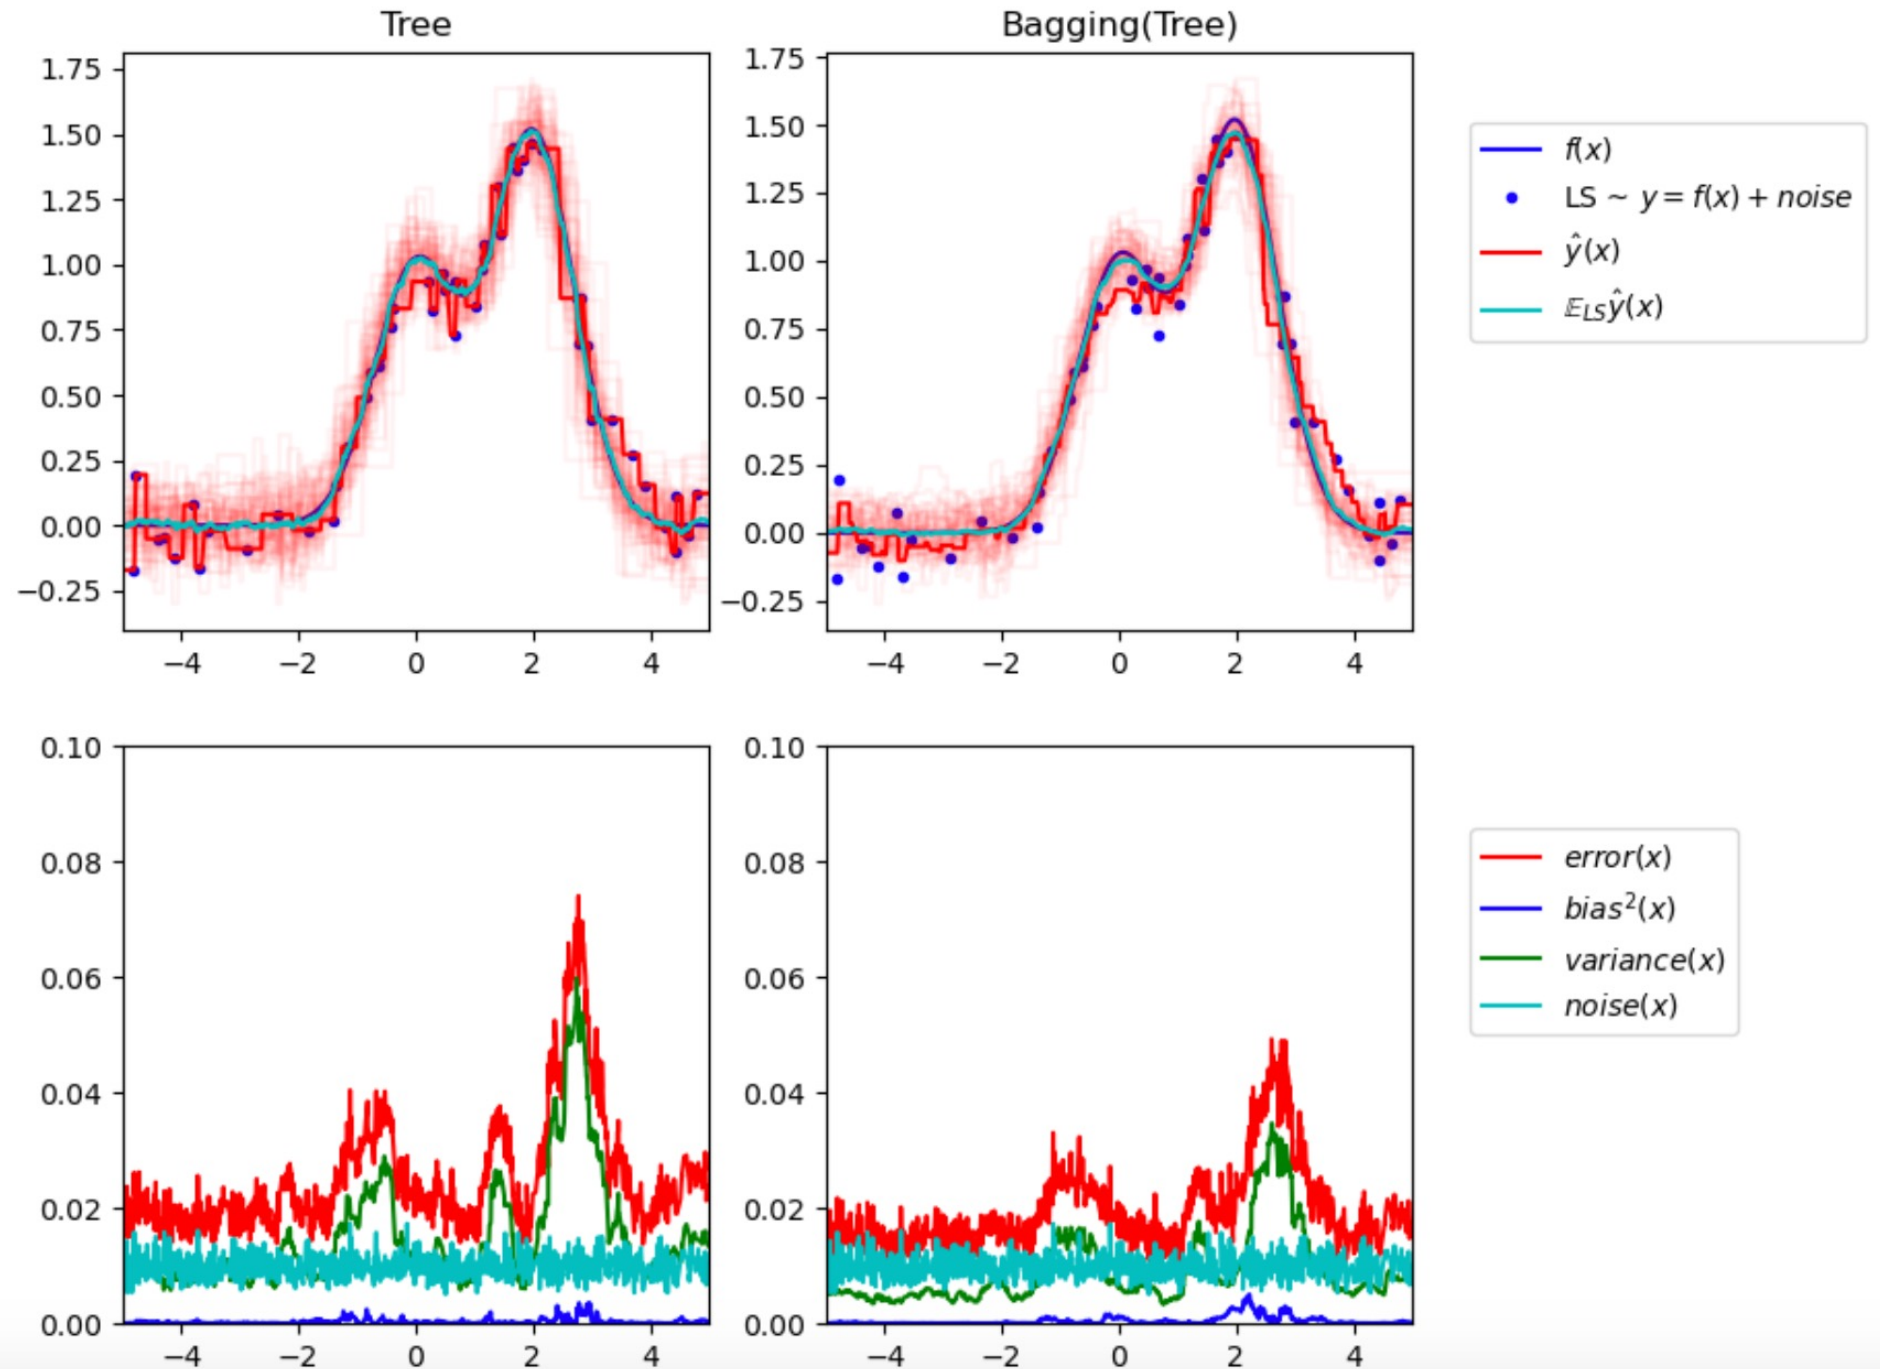


[https://scikit-learn.org/stable/auto\\_examples/ensemble/plot\\_bias\\_variance.html](https://scikit-learn.org/stable/auto_examples/ensemble/plot_bias_variance.html)

# Example



# Example



# Why it works #1

- Assume for simplicity  $y = h^*(x)$  (i.e., no noise)
- Each base learner can be written as:  $h_i(x) = h^*(x) + \epsilon_i(x)$

← Error specific to each base learner

- Expected (squared-loss) error of each **base learner**:

$$\text{err}_i = \mathbb{E}_x \left[ (h_i(x) - h^*(x))^2 \right] = \mathbb{E}_x \left[ (h^*(x) + \epsilon_i(x) - h^*(x))^2 \right] = \mathbb{E}_x [\epsilon_i(x)^2]$$

- Denote “average error” of base learners, *had they been acting individually*:

$$\text{err}_{avg} = \frac{1}{k} \sum_{i=1}^k \mathbb{E}_x [\epsilon_i(x)^2]$$

- Decomposition of error of the **ensemble**:

$$\begin{aligned} \text{err}_H &= \mathbb{E}_x \left[ (H(x) - h^*(x))^2 \right] = \mathbb{E}_x \left[ \left( \frac{1}{k} \sum_{i=1}^k h_i(x) - h^*(x) \right)^2 \right] = \mathbb{E}_x \left[ \left( \frac{1}{k} \sum_{i=1}^k \epsilon_i(x) \right)^2 \right] \\ &= \mathbb{E}_x \left[ \frac{1}{k^2} \left( \sum_{i=1}^k \epsilon_i(x)^2 + \sum_{i \neq j} \epsilon_i(x) \epsilon_j(x) \right) \right] = \frac{1}{k} \text{err}_{avg} + \frac{1}{k^2} \mathbb{E}_x \left[ \sum_{i \neq j} \epsilon_i(x) \epsilon_j(x) \right] \end{aligned}$$

# Why it works #1

- Decomposition of error of the **ensemble**:

$$\text{err}_H = \frac{1}{k} \text{err}_{avg} + \frac{1}{k^2} \mathbb{E}_x \left[ \sum_{i \neq j} \epsilon_i(x) \epsilon_j(x) \right]$$

- If errors  $\epsilon_i$ :
  - have zero mean:  $\mathbb{E}_x[\epsilon_i(x)] = 0$
  - are uncorrelated:  $\mathbb{E}_x[\epsilon_i(x) \epsilon_j(x)] = \mathbb{E}_x[\epsilon_i(x)] \mathbb{E}_x[\epsilon_j(x)]$ ,

then:  $\text{err}_H = \frac{1}{k} \text{err}_{avg}$

- **Dramatic reduction in error just by averaging!**
- In practice, errors are usually correlated
- Suffer additional “correlation errors”  $\frac{1}{k^2} \mathbb{E}_x \left[ \sum_{i \neq j} \epsilon_i(x) \epsilon_j(x) \right]$

# Why it works #2

- **Intuition:** “more” variation in input  $\Rightarrow$  less variation in output (through averaging)
- Note that examples are no longer independent (they are linked by  $S$ )
- This means  $H = \frac{1}{k} \sum_{i=1}^k h_i \xrightarrow{m \rightarrow \infty} \mu$  ( $S_i$  are not iid and so LLN does not kick in)
- Nonetheless, their probability under  $D$  is preserved
- **Claim:**  $P_Q(x, y) = P_D(x, y)$  (that is, sampling from  $D$  is like sampling  $S$  and then from  $Q$ )
- To simplify, consider finite  $x \in \{x_1, \dots, x_n\}$  with  $P(x_i) = p_i$ , and discard  $y$
- **Proof:**

# Why it works #2

- **Intuition:** “more” variation in input  $\Rightarrow$  less variation in output (through averaging)
- Note that examples are no longer independent (they are linked by  $S$ )
- This means  $H = \frac{1}{k} \sum_{i=1}^k h_i \xrightarrow{m \rightarrow \infty} \mu$  ( $S_i$  are not iid and so LLN does not kick in)
- Nonetheless, their probability under  $D$  is preserved
- **Claim:**  $P_Q(x, y) = P_D(x, y)$  (that is, sampling from  $D$  is like sampling  $S$  and then from  $Q$ )
- To simplify, consider finite  $x \in \{x_1, \dots, x_n\}$  with  $P(x_i) = p_i$ , and discard  $y$
- **Proof:**  
$$P_Q(x_i) = P(x_i \in S \text{ and was picked})$$

# Why it works #2

- **Intuition:** “more” variation in input  $\Rightarrow$  less variation in output (through averaging)
- Note that examples are no longer independent (they are linked by  $S$ )
- This means  $H = \frac{1}{k} \sum_{i=1}^k h_i \xrightarrow{m \rightarrow \infty} \mu$  ( $S_i$  are not iid and so LLN does not kick in)
- Nonetheless, their probability under  $D$  is preserved

- **Claim:**  $P_Q(x, y) = P_D(x, y)$  (that is, sampling from  $D$  is like sampling  $S$  and then from  $Q$ )
- To simplify, consider finite  $x \in \{x_1, \dots, x_n\}$  with  $P(x_i) = p_i$ , and discard  $y$

- **Proof:**
$$P_Q(x_i) = P(x_i \in S \text{ and was picked}) = \sum_{n=1}^m \binom{m}{n} p_i^n (1 - p_i)^{m-n} \cdot \frac{n}{m}$$

prob. of having  $n$  copies of  $x_i$  in  $S$

probability of picking one of those copies



# Why it works #2

- **Intuition:** “more” variation in input  $\Rightarrow$  less variation in output (through averaging)
- Note that examples are no longer independent (they are linked by  $S$ )
- This means  $H = \frac{1}{k} \sum_{i=1}^k h_i \not\rightarrow_{m \rightarrow \infty} \mu$  ( $S_i$  are not iid and so LLN does not kick in)
- Nonetheless, their probability under  $D$  is preserved

- **Claim:**  $P_Q(x, y) = P_D(x, y)$  (that is, sampling from  $D$  is like sampling  $S$  and then from  $Q$ )
- To simplify, consider finite  $x \in \{x_1, \dots, x_n\}$  with  $P(x_i) = p_i$ , and discard  $y$

- **Proof:**  
$$P_Q(x_i) = P(x_i \in S \text{ and was picked}) = \sum_{n=1}^m \underbrace{\binom{m}{n} p_i^n (1 - p_i)^{m-n}}_{\text{prob. of having } n \text{ copies of } x_i \text{ in } S} \cdot \underbrace{\frac{n}{m}}_{\text{probability of picking one of those copies}}$$
$$= \frac{1}{m} \sum_{n=1}^m \binom{m}{n} p_i^n (1 - p_i)^{m-n} n$$

expected value of binomial distribution

$$\mathbb{E}[\text{Binomial}(p_i, m)] = mp_i$$

# Why it works #2

- **Intuition:** “more” variation in input  $\Rightarrow$  less variation in output (through averaging)
- Note that examples are no longer independent (they are linked by  $S$ )
- This means  $H = \frac{1}{k} \sum_{i=1}^k h_i \xrightarrow{m \rightarrow \infty} \mu$  ( $S_i$  are not iid and so LLN does not kick in)
- Nonetheless, their probability under  $D$  is preserved

- **Claim:**  $P_Q(x, y) = P_D(x, y)$  (that is, sampling from  $D$  is like sampling  $S$  and then from  $Q$ )
- To simplify, consider finite  $x \in \{x_1, \dots, x_n\}$  with  $P(x_i) = p_i$ , and discard  $y$

- **Proof:**

$$\begin{aligned} P_Q(x_i) &= P(x_i \in S \text{ and was picked}) = \sum_{n=1}^m \underbrace{\binom{m}{n} p_i^n (1 - p_i)^{m-n}}_{\text{prob. of having } n \text{ copies of } x_i \text{ in } S} \cdot \underbrace{\frac{n}{m}}_{\text{probability of picking one of those copies}} \\ &= \frac{1}{m} \sum_{n=1}^m \binom{m}{n} p_i^n (1 - p_i)^{m-n} n = \frac{1}{m} m p_i = p_i = P_D(x_i) \end{aligned}$$

$$\mathbb{E}[\text{Binomial}(p_i, m)] = m p_i$$

# 0.632 bootstrapping

- For any  $x_i$  and  $S_j$ :
  - At each step,  $P(\text{choose } x_i) = \frac{1}{m}$
  - Overall,  $P(x_i \notin S_j) = \left(1 - \frac{1}{m}\right)^m \xrightarrow{m \rightarrow \infty} e^{-1} = 0.368$
- This means each  $S_j$  includes roughly 2/3 of the data, *regardless of m*
- Hence, each  $h_j$  is effectively trained on a different, random 1/3-2/3 “split” of the data
- Performance on non-selected points provides unbiased estimate of test error
- Sounds like cross-validation, but free! (i.e., didn't “throw away” validation examples)
- Gives insight into how variance reduction works


# Random Forests

- An extremely popular and practically useful instantiation of Bagging
- Uses decision trees, but not using ID3 – rather, uses a **different base learner** (having additional randomization)

## Random Forest( $S, k, d'$ )

1. Sub-sample  $k$  sample sets  $S_1, \dots, S_k$  of size  $m = |S|$  with replacement from  $S$
2. For each  $S_i$ , learn an unbounded-depth decision tree  $h_i$ , but with a twist:  
**at each decision node, use only a randomly sampled subset of  $d' < d$  features**  
(without replacement! i.e., each split criterion is based on different features)
3. Return  $H(x) = \frac{1}{k} \sum_{i=1}^k h_i(x)$

- Usually fairly insensitive to choices of hyper-parameters  $k, d'$  (often set  $d' = \sqrt{d}$ )



# Boosting

### Bagging:

- mostly aimed at reducing variance
- targets high-variance models  
e.g., deep trees
- linear aggregation, uniform weights:

$$H = \sum_{i=1}^k \frac{1}{k} h_i$$

- parallel training  
( $h_i$  independent given  $S$ )

### Boosting:

- mostly aimed at reducing bias
- targets high-bias models  
e.g., stumps (depth-one trees)
- linear aggregation, varying weights:

$$H = \sum_{i=1}^k \alpha_i h_i$$

- incremental (greedy) training  
( $h_i$  depends on  $h_1, \dots, h_{i-1}$ )

# Weak learners

- Boosting works well for **weak base learners**  
(this is how we can operationalize the notion of “high-bias”)

- **Definition:** (for binary classification)

The base class  $\mathcal{H} = \{h: \mathcal{X} \rightarrow \{\pm 1\}\}$  is a  $\gamma$ -weakly learnable if exists  $A$  and  $m(\delta)$  s.t. for all  $\delta \in [0,1]$  and all  $D$ , for  $S \sim D^m$ ,

$$P_D(\text{err}(A(S)) \leq 0.5 - \gamma) \geq 1 - \delta$$

- Intuitively, means learned models perform at least better than random
- A fairly minimal requirement...  
notice that if both  $h, -h \in \mathcal{H}$ , then at least one has  $\text{err} < 0.5$



# Weak learners

- Boosting works well for **weak base learners**  
(this is how we can operationalize the notion of “high-bias”)

- **Definition:** (for binary classification)

The base class  $\mathcal{H} = \{h: \mathcal{X} \rightarrow \{\pm 1\}\}$  is a  $\gamma$ -weakly learnable if exists  $A$  and  $m(\delta)$  s.t. for all  $\delta \in [0,1]$  and all  $D$ , for  $S \sim D^m$ ,

$$P_D(\text{err}(A(S)) \leq 0.5 - \gamma) \geq 1 - \delta$$

## Vs. PAC-learnable:

... and  $m(\delta, \epsilon)$  s.t.  $\forall \delta, \epsilon \in [0,1]$

$$P_D(\text{err}(A(S)) \leq \epsilon) \geq 1 - \delta$$

- Intuitively, means learned models perform at least better than random
- A fairly minimal requirement...  
notice that if both  $h, -h \in \mathcal{H}$ , then at least one has  $\text{err} < 0.5$
- **Compare to PAC;** in weak learning,  $m$  does not depending on error ( $\epsilon$ )

# Weak learners

- Boosting works well for **weak base learners**  
(this is how we can operationalize the notion of “high-bias”)

- **Definition:** (for binary classification)

The base class  $\mathcal{H} = \{h: \mathcal{X} \rightarrow \{\pm 1\}\}$  is a  $\gamma$ -weakly learnable if exists  $A$  and  $m(\delta)$  s.t. for all  $\delta \in [0,1]$  and all  $D$ , for  $S \sim D^m$ ,

$$P_D(\text{err}(A(S)) \leq 0.5 - \gamma) \geq 1 - \delta$$

- We call such  $A$  a **weak learner**,  
and the learned model  $h_S$  a **weak model**
- Simply averaging weakly-learned models (i.e., bagging) isn't going to work
- Enter **Boosting**

# Boosting

- **Boosting** uses a weak learner to form a strong learner (i.e., that PAC-learns with any  $\epsilon$ )
- **Main idea** – iteratively build aggregate model:
  - for  $t = 1 \dots T$ 
    - learn  $h_t$  with weak learner  $A$
    - set  $\alpha_t$
    - update  $H_t = H_{t-1} + \alpha_t h_t$
  - return  $H = H_T = \sum_{t=1}^T \alpha_t h_t$
- $\alpha_t$  can be either pre-determined (constant, diminishing, etc) or optimized (also greedily) (this depends on the particular learning setting)
- **Intuition:** each  $h_t$  “compensates” for errors made by its predecessors in  $H_{t-1}$
- Let’s start with an example, and then proceed to a more fundamental understanding

# AdaBoost

- **AdaBoost = Adaptive Boosting**
- Requires weak learner  $A(S; D)$  that takes as input:
  - data  $S = \{(x_i, y_i)\}_{i=1}^m$
  - vector  $D \in [0,1]^m$  having  $\sum_i D_i = 1$  that encodes per-example importance weights
- Each weak model  $h_t$  minimizes  $\sum_{i=1}^m D_i \mathbb{1}\{y_i \neq h_t(x_i)\}$

- initialize  $D^{(1)} = (1/m, \dots, 1/m)$ 
  - for  $t = 1 \dots T$ 
    - Learn weak model  $h_t = A(S; D^{(t)})$
    - set  $\alpha_t$
    - update  $D^{(t+1)}$  based on  $D^{(t)}$
  - return  $H = H_T = \sum_{t=1}^T \alpha_t h_t$

## Intuition:

- $\alpha_t$  should be high if  $h_t$  is “good”
- $D^{(t+1)}$  should up-weight examples that  $h_t$  misclassified

## For example (not real derivation):

- $\alpha_t \approx 1/\text{err}(h_t)$
- $D_i^{(t+1)} \approx 1/y_i h_t(x_i)$

## AdaBoost( $S, A, T$ )

- initialize  $D^{(1)} = (1/m, \dots, 1/m)$

- for  $t = 1 \dots T$ :

1.  $h_t = A(S; D^{(t)})$

*0-1 loss*

2.  $\varepsilon_t = \sum_{i=1}^m D_i^{(t)} \mathbb{1}\{y_i \neq h_t(x_i)\}$

3.  $\alpha_t = 1/2 \log(1/\varepsilon_t - 1)$

*guess?*

4.  $\forall i, D_i^{(t+1)} \propto D_i^{(t)} \exp\{-\alpha_t y_i h_t(x_i)\}$

# train with weak learner

# compute weighted 0/1 error

# set model coefficient

# update normalized weights

- return  $H = \text{sign}(\sum_{t=1}^T \alpha_t h_t)$


*learned (=optimized) weights*

*"more" wrong =>  
more importance*



- More to follow!

most popular base-class for boosting:



**stumps**

*= trees of depth 0*

(vs. deep trees in bagging)

# Boosting vs. Bagging



- Boosting trains each  $h_t$  on a “reweighted” sample, with importance scores  $D^{(t)}$
- Each  $D^{(t)}$  is updated according to  $D^{(t-1)}$
- Can think of Bagging as special case where  $\alpha_t = \frac{1}{T}$  and each  $D^{(t)}$ :
  - is set independently
  - and at random
  - to include “count” entries  $D_i \in \{0,1,2, \dots\}$  (and then normalize)
  - (by sampling indices with replacement)
- **Conclusion:** boosting generalizes bagging
- **Q:** What generalizes boosting?

Gradient boosting





# A gentle start

- Consider regression with least squares as a loss
- Ideally, we want  $\hat{y} = y$ .
- In practice, get residual errors:  $r_i = y_i - \hat{y}_i$
- **Idea:**
  - Train sequence of models  $h_1, h_2, \dots$
  - Train each  $h_t$  to fit the **residuals** of  $h_{t-1}$
  - Hope that each  $h_t$  “compensates” for previous errors
  - Output  $H = \sum_t h_t$



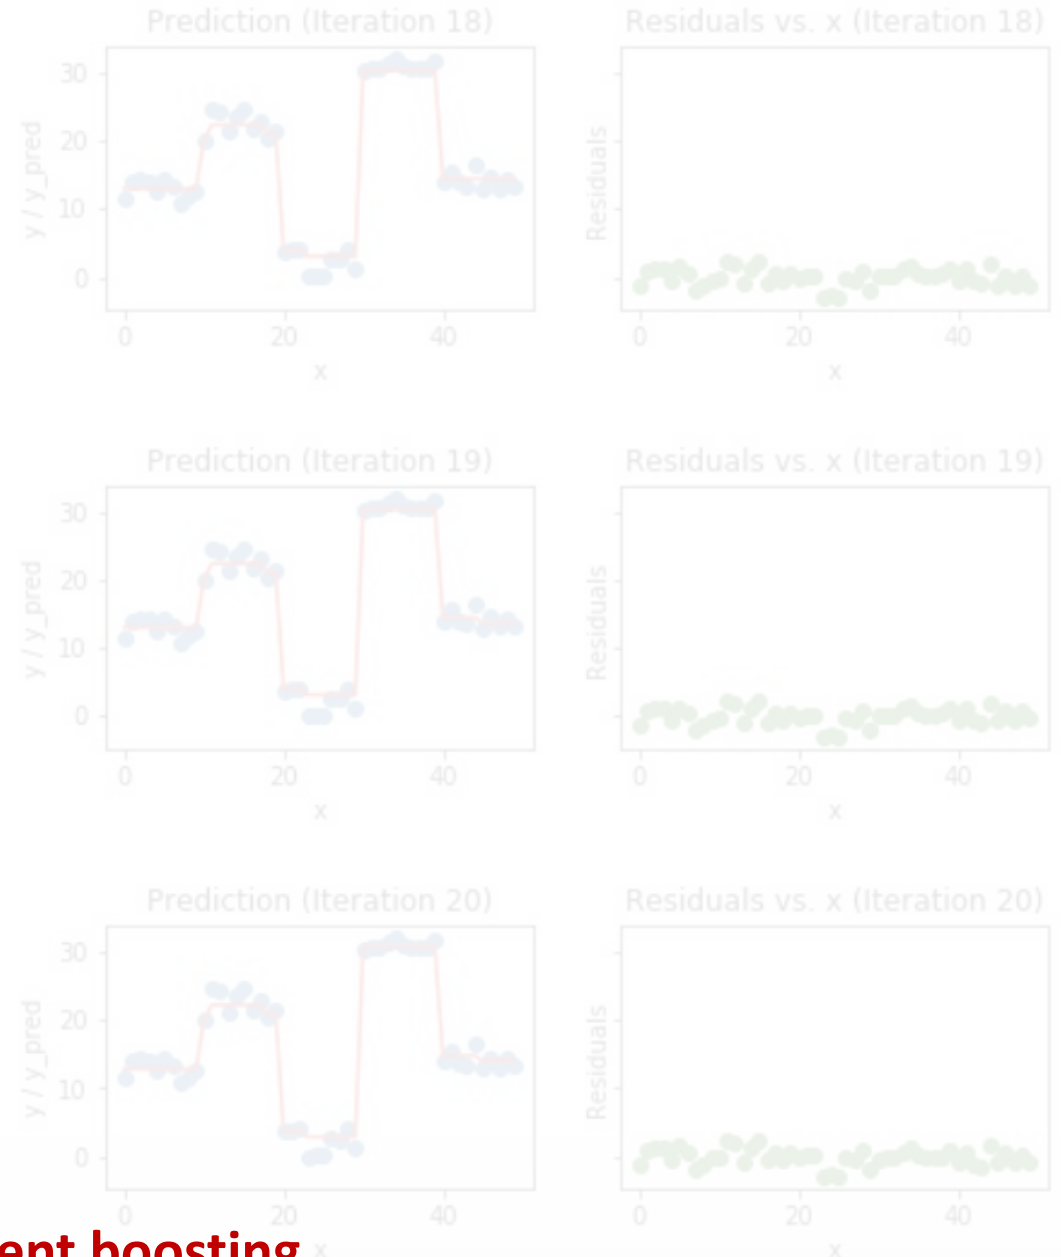
# A gentle start

- Consider regression with least squares as a loss
- Ideally, we want  $\hat{y} = y$ .
- In practice, get residual errors:  $r_i = y_i - \hat{y}_i$
- **Idea:**
  - Train sequence of models  $h_1, h_2, \dots$
  - Train each  $h_t$  to fit the **residuals** of  $h_{t-1}$
  - Hope that each  $h_t$  “compensates” for previous errors
  - Output  $H = \sum_t h_t$



# A gentle start

- Consider regression with least squares as a loss
- Ideally, we want  $\hat{y} = y$ .
- In practice, get residual errors:  $r_i = y_i - \hat{y}_i$
- **Idea:**
  - Train sequence of models  $h_1, h_2, \dots$
  - Train each  $h_t$  to fit the **residuals** of  $h_{t-1}$
  - Hope that each  $h_t$  “compensates” for previous errors
  - Output  $H = \sum_t h_t$
- But we’d like a more general solution; enter **gradient boosting**



# Boosted least squares

- At step  $t$ , say we want to minimize  $L_S(H_{t-1}) = \sum_i \frac{1}{2} (H_{t-1}(x_i) - y_i)^2$ , by adjusting prior predictions  $H(x_i)$

- Think of  $H_{t-1}(x_i)$  as parameters and take the gradient:

$$g_i = [\nabla L_S(H_{t-1})]_{x_i} = \frac{\partial L_S(H_{t-1})}{\partial H_{t-1}(x_i)} = \frac{\partial \frac{1}{2} (H_{t-1}(x_i) - y_i)^2}{\partial H_{t-1}(x_i)} = H_{t-1}(x_i) - y_i$$

- Residuals are **negative gradients(!)**:  $y_i - H_{t-1}(x_i) = -g_i$   


- $h$  fits the residuals and thus approximates the negative gradient

$$h(x_i) \approx -g_i$$

# Boosted least squares

- **Gradient-Boosted Least Squares:**

Start with initial model, say  $H_0 = \frac{1}{m} \sum_i y_i$

at each step  $t$ ,

1. Compute negative gradient  $r_i = y_i - H_{t-1}(x_i)$


2. Learn  $h_t \in \operatorname{argmin}_{h \in \mathcal{H}} \sum_{i=1}^m (h(x_i) - r_i)^2$

3. Take optimal local step:  $H_t = H_{t-1} + h_t$


- Notice the update rule


$$H_t(x_i) = H_{t-1}(x_i) + h_t(x_i) \approx H_{t-1}(x_i) - g_i = H_{t-1}(x_i) - 1[\nabla L_S(H_{t-1})]_{x_i}$$

**residuals / negative gradient (Round 0)**



**Observed Data vs. Fit Vector (Round 0)**





# Functional gradient descent

- Boosted least squares approximates the following update:

$$H_t = H_{t-1} + h_t \approx H_{t-1} - 1[\nabla L_S(H_{t-1})]$$

- Generally, we *want*:  $H_t = H_{t-1} - \eta \nabla L_S(H_{t-1})$
- What does this remind you of?
- Recall gradient descent in *parameter space* (e.g., for linear  $h_w(x) = w^\top x$ ):

$$w_t = w_{t-1} - \eta \nabla L_S(w_t)$$

- We can now think of GD as learning an “ensemble”:

$$\hat{h} = h_{\hat{w}}, \quad \hat{w} = \sum_t \alpha_t h_t, \quad h_t = \nabla L_S(w_t), \quad \alpha_t = -\eta$$

- Gradient boosting performs **gradient descent in function space**



# Functional gradient descent

- Gradient descent updates:  $H_t = H_{t-1} - \eta \nabla L_S(H_{t-1})$
- Great! But – two problems:
  1. no one said  $\nabla L_S(H_{t-1}) \in \mathcal{H}$  (recall we want to learn an ensemble)
  2. we can't compute  $\nabla L_S(H_{t-1})$  for all  $x$  – we only have  $S$ !
- Fortunately, one solution solves both problems at once

- **Idea:**

1. Compute “empirical” gradient, with entries only for observed data points:

$$g_i = [\nabla L_S(H_{t-1})]_{x_i} = \frac{\partial \ell(y_i, H_{t-1}(x_i))}{\partial H_{t-1}(x_i)}$$

2. Find  $h_t \in \mathcal{H}$  that best fits empirical gradient, i.e.,  $h_t(x_i) \approx g_i \quad \forall i \in [m]$
3. Update with inferred model:  $H_t = H_{t-1} - \eta h_t$

# Functional gradient descent

- **Gradient Boosting:** at each step  $t$ ,
  1. Compute “empirical” gradient  $g_i = [\nabla L_S(H_{t-1})]_{x_i}$
  2. Find  $h_t \in \mathcal{H}$  that best fits empirical gradient, i.e.,  $h_t(x_i) \approx g_i \quad \forall i \in [m]$
  3. Update with inferred model:  $H_t = H_{t-1} - \eta h_t$

# Functional gradient descent

- **Gradient Boosting:** at each step  $t$ ,

1. Compute “empirical” gradient  $g_i = [\nabla L_S(H_{t-1})]_{x_i}$

2. Learn  $h_t \in \operatorname{argmin}_{h \in \mathcal{H}} \sum_{i=1}^m \ell(g_i, h(x_i))$

(can use *any* loss function; should choose according to what  $g_i$  are)

3. Update with inferred model:  $H_t = H_{t-1} - \eta h_t$

# Functional gradient descent

- **Gradient Boosting:** at each step  $t$ ,

1. Compute “empirical” gradient  $g_i = [\nabla L_S(H_{t-1})]_{x_i}$

2. Learn  $h_t \in \operatorname{argmin}_{h \in \mathcal{H}} \sum_{i=1}^m \ell(g_i, h(x_i))$

(can use *any* loss function; should choose according to what  $g_i$  are)

3. Take optimal local step:  $H_t = \operatorname{argmin}_{\alpha \in \mathbb{R}} H_{t-1} + \alpha h_t$

(aka “line search”; when possible; sometimes has closed-form solution)

- Together, can think of as approximating  $\operatorname{argmin}_{h \in \mathcal{H}, \alpha \in \mathbb{R}} L_S(H_{t-1} + \alpha h)$
- This approach is called **forward stage-wise optimization**
- Gradient-Boosted least squares is one example

# AdaBoost, revisited

- **Claim:**

AdaBoost is an instance of Gradient Boosting

- **Loss:** exponential,  $\ell^{\text{exp}}(y, \hat{y}) = e^{-y\hat{y}}$

- **Gradient:**

$$g_i = \frac{\partial \ell(y_i, H(x_i))}{\partial H(x_i)} = -y_i e^{-y_i H(x_i)}$$


- **Will see in tirgul:**

1. updating  $D$  approximates gradient step
2.  $\alpha_t$  is the optimal step size

- AdaBoost converges *exponentially fast* (won't show)

## AdaBoost( $S, A, T$ )

- initialize  $D^{(1)} = (1/m, \dots, 1/m)$
- for  $t = 1 \dots T$ :
  - # train with weak learner
  - 1.  $h_t = A(S; D^{(t)})$
  - # compute weighted 0/1 error
  - 2.  $\varepsilon_t = \sum_{i=1}^m D_i^{(t)} \mathbb{1}\{y_i \neq h_t(x_i)\}$
  - # set model coefficient
  - 3.  $\alpha_t = 1/2 \log(1/\varepsilon_t - 1)$
  - # update normalized weights
  - 4.  $\forall i, D_i^{(t+1)} \propto D_i^{(t)} \exp\{-\alpha_t y_i h_t(x_i)\}$
- return  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t)$

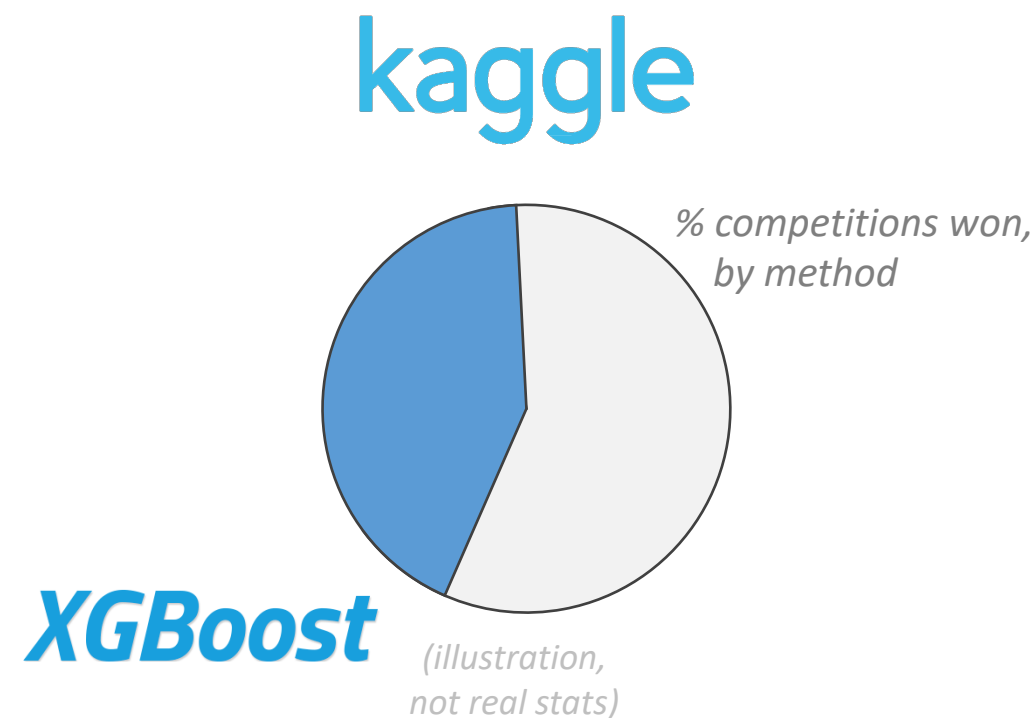



# Discussion

- Ensembling is an approach to combine multiple models into a single, better model
  - **Bagging** works well when base learner has **high variance**
  - **Boosting** works well when base learner has **high bias**
- Both propose ways to optimize over a class of linear ensemble models,  $H = \sum_i \alpha_i h_i$   
(helpful when direct optimization is hard)
- **Cons:** runtime (training, prediction)
- **Pros:**
  - expressivity
  - works very, very well

# Discussion


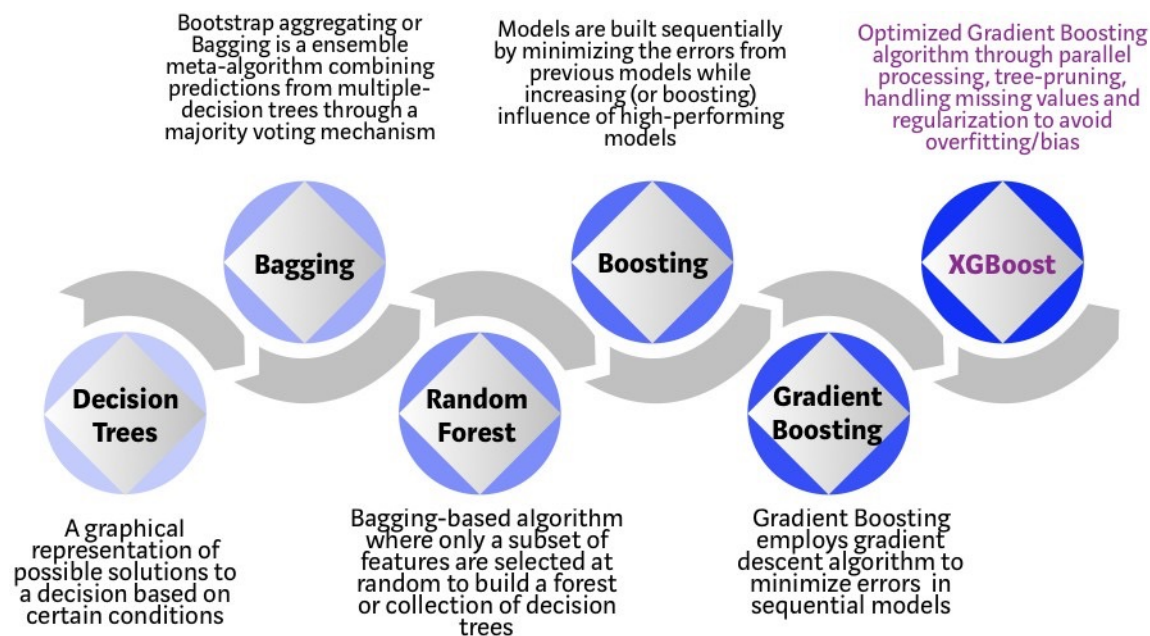
**XGBoost:** highly-optimized flexible gradient boosting tool





# Discussion

**XGBoost:** highly-optimized flexible gradient boosting tool



# Discussion

- **Interpretation:**

ensembles are linear models over a representation consisting of weak model predictions:

$$H(x) = \sum_i \alpha_i h_i(x) \equiv H(x) = w^\top \phi(x), \quad (\phi(x))_i = h_i(x)$$

- Can think of learning as done in two steps:
  1. learn representation mapping  $\phi$
  2. learn linear model  $w$  over representation  $\phi$
- We will see that deep learning can be thought of as jointly learning  $\phi$  and  $w$   
(and so entries in  $\phi$  will now longer correspond to predictions of some  $h$ )

# Next week

- **Part III:** *more supervised learning*
  1. Regression
  2. Bagging and boosting (today)
  3. **Generative models**
  4. Deep learning

